

CS 2200

LAB 2

Announcements

- **Homework 1** released after lab hopefully, due **Sep 5th @ 6:30 PM**
- **Project 1** released, is due on **Sep 17th @ 11:59 PM**
- **Lab attendance** will be taken starting today

Endianness & Padding

Coming Up!

We need a test C program that tells you whether your platform is big or little endian.


What does that mean?

How should we go about it?

Endianness

Endianness defines how the system will organize the bytes that make up numbers. All data is comprised into chunks of bytes, with the largest byte in the number called the Most Significant Byte (MSB), and the smallest byte called the Least Significant Byte (LSB).

- Big Endian = The most significant byte is at the lowest memory address
- Little Endian = The least significant byte is at the lowest memory address



Endianness Example

```
int i = 0x12345678 /* 4 bytes */
```

Little endian

	+0	+1	+2	+3
0x00				

Big endian

	+0	+1	+2	+3
0x00				

Endianness Example

```
int i = 0x12345678 /* 4 bytes */
```

Little endian

	+0	+1	+2	+3
0x00	0x78	0x56	0x34	0x12

Big endian

	+0	+1	+2	+3
0x00	0x12	0x34	0x56	0x78

A little DIY!

In groups, try to write and test a C program that tells you whether your platform is big or little endian. Take 10 mins and feel free to collaborate!

HINT: Try using a union with int and char members

A little DIY!

In groups, try to write and test a C program that tells you whether your platform is big or little endian. Take 10 mins and feel free to collaborate!

HINT: Try using a union with int and char members

The solution C file is in the attendance quiz endianness question

Endianness In C

If the computer that is running the program is little endian, what would the expected output be?

Similarly what would the output be for a big endian system?

```
void show_mem_rep(char *start, int n) {  
    for (int i = 0; i < n; i++)  
        printf(" %.2x", start[i]);  
    printf("\n");  
}  
  
int main()  
{  
    int i = 0x01234567;  
    show_mem_rep((char *)&i, sizeof(i));  
    return 0;  
}
```

Endianness In C

If the computer that is running the program is little endian, what would the expected output be?

67 45 23 01

Similarly what would the output be for a big endian system?

01 23 45 67

```
void show_mem_rep(char *start, int n) {  
    for (int i = 0; i < n; i++)  
        printf(" %.2x", start[i]);  
    printf("\n");  
}  
  
int main()  
{  
    int i = 0x01234567;  
    show_mem_rep((char *)&i, sizeof(i));  
    return 0;  
}
```

Padding in C

How do we predict what the sizes of these structures will be?

Presume we're using an ARM system that requires natural alignment.

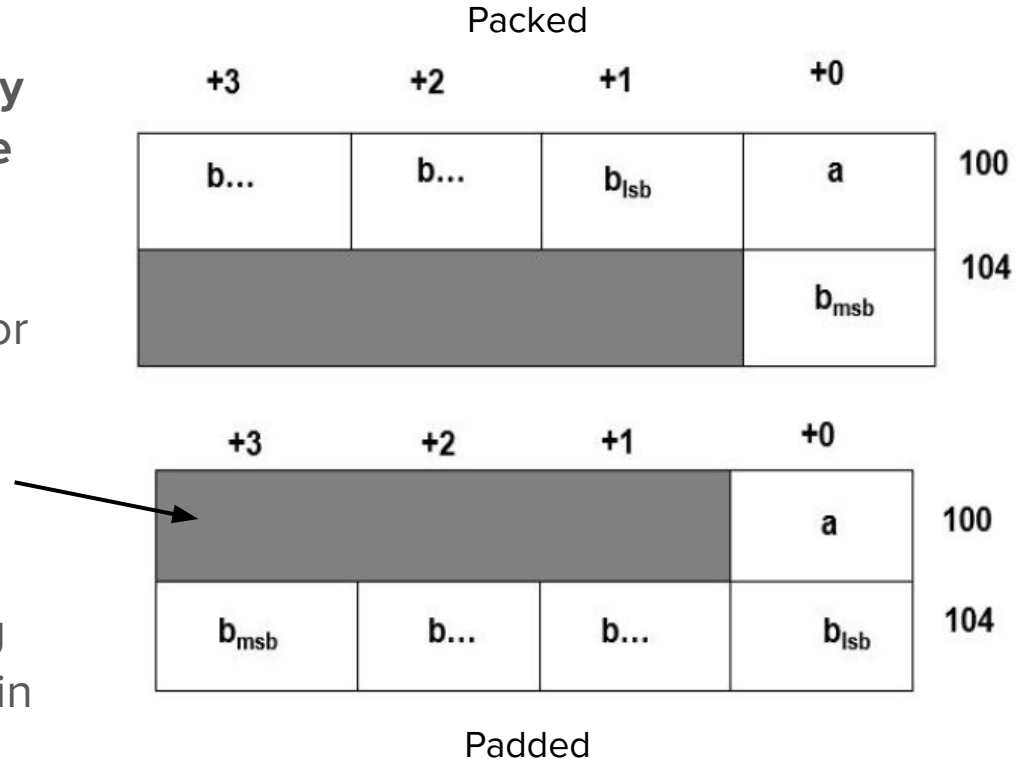
Natural alignment means that a data type will be stored at a memory address that is a **multiple** of its size!

And structures will be padded so their size will be a **multiple** of the **largest** data type they contain. (Why?)

```
struct A {  
    char *ptr; // 8 bytes  
    char c;    // 1 byte  
    long num;  // 8 bytes  
};  
  
struct B {  
    char c;    // 1 byte  
    short s[2]; // 2 x 2 bytes  
    int i;     // 4 bytes  
};  
  
struct C {  
    int i;     // 4 bytes  
    short s[2]; // 2 x 2 bytes  
    char c;    // 1 byte  
};
```

Padding Structs

- **Unaligned access to memory is often expensive and some architectures don't allow it.**
- For example, ARM CPUs require “natural alignment” for their data types. MIPS and Intel CPUs implement unaligned access, but do it significantly more slowly.
- Compilers will apply **padding** to small data types to maintain alignment in memory.



Padding Example

```
struct Q {  
    char c;      /* 1 bytes */  
    short s;     /* 2 bytes */  
    int i;       /* 4 bytes */  
};
```

	+0	+1	+2	+3
0x00				
0x04				

Padding Example

```
struct Q {  
    char c;      /* 1 bytes */  
    short s;     /* 2 bytes */  
    int i;       /* 4 bytes */  
};
```

	+0	+1	+2	+3
0x00	c		s0	s1
0x04	i0	i1	i2	i3

Padding in C

What do you think the following output will be?

****Word is 8 Bytes****

What are the sizes of each struct?

- A
- B
- C

```
struct A {
    char *ptr; // 8 bytes
    char c;    // 1 byte
    long num;  // 8 bytes
};

struct B {
    short s;    // 2 bytes
    int i;      // 4 bytes
    short s2;   // 2 bytes
};

struct C {
    int i;      // 4 bytes
    short s;    // 2 bytes
    short s2;   // 2 bytes
};

int main() {
    printf("A size: %lu\n", sizeof(struct A));
    printf("B size: %lu\n", sizeof(struct B));
    printf("C size: %lu\n", sizeof(struct C));
}
```


Padding in C

What do you think the following output will be?

****Word is 8 Bytes****

What are the sizes of each struct?

- A: **24** bytes
- B: **12** bytes
- C: **8** bytes

```
struct A {
    char *ptr; // 8 bytes
    char c;    // 1 byte
    long num;  // 8 bytes
};

struct B {
    short s;    // 2 bytes
    int i;      // 4 bytes
    short s2;   // 2 bytes
};

struct C {
    int i;      // 4 bytes
    short s;    // 2 bytes
    short s2;   // 2 bytes
};

int main() {
    printf("A size: %lu\n", sizeof(struct A));
    printf("B size: %lu\n", sizeof(struct B));
    printf("C size: %lu\n", sizeof(struct C));
}
```

Padding in C cont.

- Let's revisit struct B from last slide
- What goes where?
- What did you get from your calculation?

```
struct B {
    short s;    // 2 bytes
    int i;      // 4 bytes
    short s2;   // 2 bytes
};
```

```
struct C {
    int i;        // 4 bytes
    short s;      // 2 bytes
    short s2;     // 2 bytes
};
```

[illegible]


Padding in C cont.

```
struct B {  
    short s;    // 2 bytes  
    int i;      // 4 bytes  
    short s2;   // 2 bytes  
};
```

```
struct C {  
    int i;      // 4 bytes  
    short s;    // 2 bytes  
    short s2;   // 2 bytes  
};
```

- What did we get?

Padding!



- Struct B takes up **12 bytes** of space
 - Memory addresses 0x02, 0x03, 0x0a, 0x0b are used as padding

Padding in C cont.


```
struct B {  
    short s;    // 2 bytes  
    int i;      // 4 bytes  
    short s2;   // 2 bytes  
};
```

- Now let's do struct C

- Has the **same members** as B, just arranged differently

```
struct C {  
    int i;      // 4 bytes  
    short s;    // 2 bytes  
    short s2;   // 2 bytes  
};
```

Padding!



- Struct B takes up **12 bytes** of space

- Memory addresses 0x02, 0x03, 0x0a, 0x0b are used as padding

Padding in C cont.

```
struct B {  
    short s;    // 2 bytes  
    int i;      // 4 bytes  
    short s2;   // 2 bytes  
};
```

```
struct C {  
    int i;      // 4 bytes  
    short s;    // 2 bytes  
    short s2;   // 2 bytes  
};
```

- What did we get for struct C?

		+0	+1	+2	+3	+4	+5	+6	+7	+8	+9	+10	+11
B	0x00	s0	s1			i0	i1	i2	i3	s0	s1		
	0x0c	i0	i1	i2	i3	s0	s1	s0	s1				
C													

- Struct B takes up **12 bytes** of space in mem
- Struct C takes up **8 bytes** of space in mem

Padding in C cont.

```
struct B {  
    short s;    // 2 bytes  
    int i;      // 4 bytes  
    short s2;   // 2 bytes  
};
```

```
struct C {  
    int i;      // 4 bytes  
    short s;    // 2 bytes  
    short s2;   // 2 bytes  
};
```

- Rule of thumb: Arranging struct members in **decreasing** order of size will minimize padding!
- Will allow more efficient use of memory space while retaining memory access efficiency

		+0	+1	+2	+3	+4	+5	+6	+7	+8	+9	+10	+11
B	0x00	s0	s1			i0	i1	i2	i3	s0	s1		
	0x0c	i0	i1	i2	i3	s0	s1	s0	s1				
C													

- Struct B takes up **12 bytes** of space in mem
- Struct C takes up **8 bytes** of space in mem

Today's Attendance Number

67,890

LC-2200

Instruction Set Architecture

Problem


- We need a LC-2200 assembly language function, maxmindiff, that computes the difference between the largest and smallest elements in an integer array
- Arguments: the address of the array and the length of the array
- Return value: the difference of the max and min elements in the array

How do we go about this? What do we need to know?

In a bit, we'll use this pseudo-code

```
int maxmindiff(int a[], int alen) {  
  
    int min, max, i;  
    min = a[0];  
    max = a[0];  
    for (i = 1; i < alen; i++) {  
        if (a[i] < min)  
            min = a[i];  
        if (a[i] > max)  
            max = a[i];  
    }  
    return max - min;  
  
}
```

Software



Instruction Set


Hardware

Instruction Set Architecture

- Provides an abstraction for using an implementation of a data path (usually hardware, but it can be software)
- The implementation (hardware or software) interprets the machine instructions described by the ISA
- The ISA provides a target for a compiler's output
- Ex: x86, ARM, RISC-V, MIPS, POWERPC, etc ...

LC-2200 Architectural Details

- 32-bit Architecture
- 32-bit word addressable memory
- Fixed Length (32 bit) Instructions
- 16 General Purpose Registers
- Load-Store Architecture
 - All data operations occur on registers



LC 2200 Instruction Set

Mnemonic Example	Format	Opcode	Action Register Transfer Language
add add \$v0, \$a0, \$a1	R	0 0000 ₂	Add contents of reg Y with contents of reg Z, store results in reg X. RTL: $\$v0 \leftarrow \$a0 + \$a1$
nand nand \$v0, \$a0, \$a1	R	1 0001 ₂	Nand contents of reg Y with contents of reg Z, store results in reg X. RTL: $\$v0 \leftarrow \sim(\$a0 \&\& \$a1)$
addi addi \$v0, \$a0, 25	I	2 0010 ₂	Add Immediate value to the contents of reg Y and store the result in reg X. RTL: $\$v0 \leftarrow \$a0 + 25$
lw lw \$v0, 0x42(\$fp)	I	3 0011 ₂	Load reg X from memory. The memory address is formed by adding OFFSET to the contents of reg Y. RTL: $\$v0 \leftarrow \text{MEM}[\$fp + 0x42]$
sw sw \$a0, 0x42(\$fp)	I	4 0100 ₂	Store reg X into memory. The memory address is formed by adding OFFSET to the contents of reg Y. RTL: $\text{MEM}[\$fp + 0x42] \leftarrow \$a0$
beq beq \$a0, \$a1, done	I	5 0101 ₂	Compare the contents of reg X and reg Y. If they are the same, then branch to the address PC+1+OFFSET, where PC is the address of the beq instruction. RTL: if($\$a0 == \$a1$) PC \leftarrow PC+1+OFFSET

Note: For programmer convenience (and implementer confusion), the assembler computes the OFFSET value from the number or symbol given in the instruction and the assembler's idea of the PC. In the example, the assembler stores done-(PC+1) in OFFSET so that the machine will branch to label "done" at run time.

jalr jalr \$at, \$ra	J	6 0110 ₂	First store PC+1 into reg Y, where PC is the address of the jalr instruction. Then branch to the address now contained in reg X. Note that if reg X is the same as reg Y, the processor will first store PC+1 into that register, then end up branching to PC+1. RTL: $\$ra \leftarrow \text{PC}+1$; PC \leftarrow \$at Note that an unconditional jump can be realized using jalr \$ra, \$r0, and discarding the value stored in \$r0 by the instruction. This is why there is no separate jump instruction in LC-2200.
nop	n.a.	n.a.	Actually a pseudo instruction (i.e. the assembler will emit: add \$zero, \$zero, \$zero)
halt halt	O	7 0111 ₂	

Additions to the textbook instructions

<code>lea</code> <code>lea \$v0, 0x10</code>	I	9 1001_2	Load reg X from with PC-relative address. The effective address is formed by adding OFFSET to the contents of the incremented PC. RTL: $\$v0 \leftarrow PC + 1 + 0x10$
<code>blt, bgt</code> <code>blt \$a0, \$a1,</code> <code>done</code> <code>bgt \$a0, \$a1,</code> <code>done</code>	I	8 1000_2 10 1010_2	Compare the contents of reg X and reg Y. If X is less than Y (greater than), then branch to the address $PC+1+OFFSET$, where PC is the address of the blt(bgt) instruction. RTL: if ($\$a0 < \$a1$) $PC \leftarrow PC+1+OFFSET$

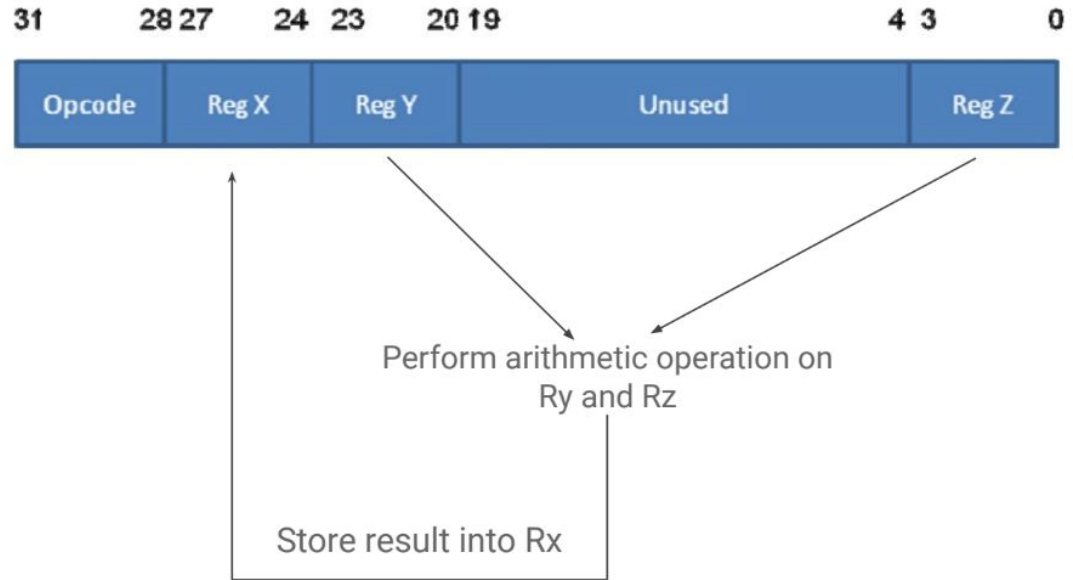
Addressing Mode

- The format of the instructions, and how the data within them is accessed.
- LC-2200 has 4 main types
 - R-type (Register)
 - I-type (Immediate)
 - J-type (Jump)
 - O-type also exists, but its sole use is NOP and HALT
 - Indicates no operands at all!

R-Type Instructions

Uses 3 operands contained in registers.

- **ADD** \$v0, \$a0, \$a1
 - Opcode 0b0000
- **NAND** \$v0, \$a0, \$a1
 - Opcode 0b0001



I-Type Instructions

Uses an immediate value in the instruction, not one pulled from a register or memory.

- **ADDI** \$v0, \$a1, 1
 - Opcode 0b0010
- **LW** \$v0, 0x2(\$a1)
 - Opcode 0b0011
- **SW** \$v0, 0x2(\$a1)
 - Opcode 0b0100
- **LEA** \$a0, LABEL
 - Opcode 0b1001
- **BEQ** \$a0, \$a1, LABEL
 - Opcode 0b0101
- **BLT** \$a0, \$a1, LABEL
 - Opcode 0b1000
- **BGT** \$a0, \$a1, LABEL
 - Opcode 0b1010




-
- lw/sw $\$data_register, offset(\$base_register)$
- Load/store word at
 - $MEM[\$base_register + offset]$
 - Into/from $\$data_register$


J-Type Instructions

Uses exactly two registers, the rest of the instruction space is unused.




- **JALR \$at, \$ra**
 - Stores the current PC in \$ra, then jumps to \$at



Unaligned



Aligned



Mnemonic Example	Format	Opcode	Action Register Transfer Language
add add \$v0, \$a0, \$a1	R	0 0000 ₂	Add contents of reg Y with contents of reg Z, store results in reg X. RTL: $\$v0 \leftarrow \$a0 + \$a1$
nand nand \$v0, \$a0, \$a1	R	1 0001 ₂	Nand contents of reg Y with contents of reg Z, store results in reg X. RTL: $\$v0 \leftarrow \sim(\$a0 \& \$a1)$
addi addi \$v0, \$a0, 25	I	2 0010 ₂	Add Immediate value to the contents of reg Y and store the result in reg X. RTL: $\$v0 \leftarrow \$a0 + 25$
lw lw \$v0, 0x42(\$fp)	I	3 0011 ₂	Load reg X from memory. The memory address is formed by adding OFFSET to the contents of reg Y. RTL: $\$v0 \leftarrow \text{MEM}[\$fp + 0x42]$
sw sw \$a0, 0x42(\$fp)	I	4 0100 ₂	Store reg X into memory. The memory address is formed by adding OFFSET to the contents of reg Y. RTL: $\text{MEM}[\$fp + 0x42] \leftarrow \$a0$
beq beq \$a0, \$a1, done	I	5 0101 ₂	Compare the contents of reg X and reg Y. If they are the same, then branch to the address PC+1+OFFSET, where PC is the address of the beq instruction. RTL: if(\$a0 == \$a1) PC ← PC+1+OFFSET
Note: For programmer convenience (and implementer confusion), the assembler computes the OFFSET value from the number or symbol given in the instruction and the assembler's idea of the PC. In the example, the assembler stores done-(PC+1) in OFFSET so that the machine will branch to label "done" at run time.			
jlr jlr \$a0, \$ra	J	6 0110 ₂	First store PC+1 into reg Y, where PC is the address of the jlr instruction. Then branch to the address now contained in reg X. Note that if reg X is the same as reg Y, the processor will first store PC+1 into that register, then end up branching to PC+1. RTL: $\$ra \leftarrow \text{PC}+1$; PC ← \$a0 Note that an unconditional jump can be realized using jlr \$ra, \$t0, and discarding the value stored in \$t0 by the instruction. This is why there is no separate jump instruction in LC-2200.
nop	n.a.	n.a.	Actually a pseudo instruction (i.e. the assembler will emit: add \$zero, \$zero, \$zero)
halt halt	O	7 0111 ₂	

Let's try to implement this now!

```
int maxmindiff(int a[], int alen) {  
  
    int min, max, i;  
    min = a[0];  
    max = a[0];  
    for (i = 1; i < alen; i++) {  
        if (a[i] < min)  
            min = a[i];  
        if (a[i] > max)  
            max = a[i];  
    }  
    return max - min;  
  
}
```

Main program stub for testing

```
lea $t0, stack          ! initialize the stack
pointer
lw  $t0, 0x0($t0)
lw  $sp, 0($t0)

lea  $a0, a             ! put address of a as argument 1
lw  $a0, 0x0($a0)
addi $a1, $zero, 12     ! put the value from 12 as
argument 2

lea $at, maxmindiff     ! call our function
jalr $ra, $at

halt

stack:.fill 0x8000
```

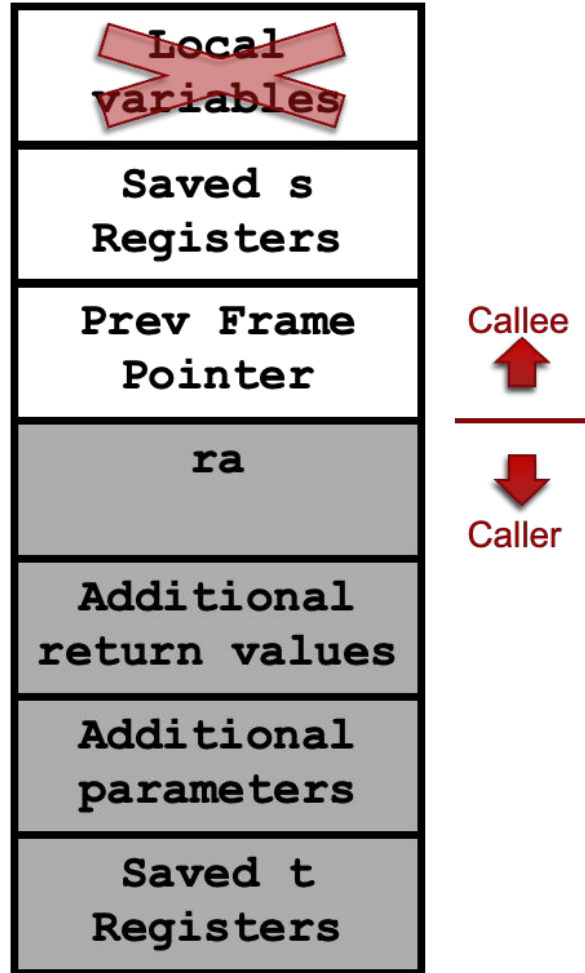
```
stack .fill 0x8000

a:
    .fill 7
    .fill 12
    .fill 8
    .fill 23
    .fill 2
    .fill 15
    .fill 9
    .fill 12
    .fill 10
    .fill 4
    .fill 6
    .fill 23
```

What do we need in the stack?

Since we're the callee, we just need to handle callee responsibilities.

So what does that look like?



Code to set up and tear down the stack frame

```
! int maxmindiff(int a[], int alen) {
```

maxmindiff:

```
addi $sp, $sp, -1    ! Push old $fp
```

```
sw    $fp, 0($sp)
```

```
addi  $fp, $sp, 0    ! Set new $fp
```

```
addi  $sp, $sp, -3    ! Push the 3 s  
registers
```

```
sw    $s0, -1($fp)
```

```
sw    $s1, -2($fp)
```

```
sw    $s2, -3($fp)
```

```
!int min, max, i;
```

```
!min = a[0];
```

```
!max = a[0];
```

```
!for (i = 1; i < alen; i++) {  
    !if (a[i] < min)  
        !min = a[i];  
    !if (a[i] > max)  
        !max = a[i];  
    !}  
!return max - min;
```

```
lw    $s2, -3($fp)    ! Restore the 3 s registers
```

```
lw    $s1, -2($fp)
```

```
lw    $s0, -1($fp)
```

```
lw    $fp, 0($fp) ! Restore the old $fp
```

```
addi  $sp, $sp, 4 ! Pop all 4 words off the stack
```

```
jalr  $at, $zero ! Return
```

```
!}
```


Implementing the pseudo-code

```
!int min, max, i;
```

!\$s0 is min, \$s1 is max, \$s2 is i

!\$a0 is a, \$a1 is alen

```
!min = a[0];
```

Implementing the pseudo-code

```
!int min, max, i;
```

```
!$s0 is min, $s1 is max, $s2 is i
```

```
!$a0 is a, $a1 is alen
```

```
!min = a[0];
```

```
lw $s0, 0($a0)
```

```
!max = a[0];
```

Implementing the pseudo-code

```
!int min, max, i;
```

```
!$s0 is min, $s1 is max, $s2 is i
```

```
!$a0 is a, $a1 is alen
```

```
!min = a[0];
```

```
lw $s0, 0($a0)
```

```
!max = a[0];
```

```
lw $s1, 0($a0)
```

```
!for (i = 1; i < alen; i++) {
```

Implementing the pseudo-code

```
!int min, max, i;
```

!\$s0 is min, \$s1 is max, \$s2 is i

!\$a0 is a, \$a1 is alen

```
!min = a[0];
```

```
lw $s0, 0($a0)
```

```
!max = a[0];
```

```
lw $s1, 0($a0)
```

```
!for (i = 1; i < alen; i++) {
```

```
addi $t2, $zero, 12
```

```
addi $s2, $zero, 1
```

```
loop:
```

```
bgt $s2, $t2, endloop
```

```
beq $s2, $t2, endloop
```

```
!if (a[i] < min)
```

Implementing the pseudo-code

```
!int min, max, i;
```

```
!$s0 is min, $s1 is max, $s2 is i
```

```
!$a0 is a, $a1 is alen
```

```
!min = a[0];
```

```
lw $s0, 0($a0)
```

```
!max = a[0];
```

```
lw $s1, 0($a0)
```

```
!for (i = 1; i < alen; i++) {
```

```
addi $t2, $zero, 12
```

```
addi $s2, $zero, 1
```

```
loop:
```

```
bgt $s2, $t2, endloop
```

```
beq $s2, $t2, endloop
```

```
!if (a[i] < min)
```

```
add $t0, $s2, $a0
```

```
lw $t0, 0($t0)
```

```
bgt $t0, $s0, nomin
```

```
beq $t0, $s0, nomin
```

```
!min = a[i];
```

Implementing the pseudo-code

```
!int min, max, i;  
!$s0 is min, $s1 is max, $s2 is i  
!$a0 is a, $a1 is alen  
!min = a[0];  
lw $s0, 0($a0)  
!max = a[0];  
lw $s1, 0($a0)  
!for (i = 1; i < alen; i++) {  
addi $t2, $zero, 12  
addi $s2, $zero, 1  
loop:  
bgt $s2, $t2, endloop  
beq $s2, $t2, endloop
```

```
!if (a[i] < min)  
add $t0, $s2, $a0  
lw $t0, 0($t0)  
bgt $t0, $s0, nomin  
beq $t0, $s0, nomin  
!min = a[i];  
addi $s0, $t0, 0  
nomin:  
!if (a[i] > max)
```

Implementing the pseudo-code

```
!int min, max, i;  
!$s0 is min, $s1 is max, $s2 is i  
!$a0 is a, $a1 is alen  
!min = a[0];  
lw $s0, 0($a0)  
!max = a[0];  
lw $s1, 0($a0)  
!for (i = 1; i < alen; i++) {  
addi $t2, $zero, 12  
addi $s2, $zero, 1  
loop:  
bgt $s2, $t2, endloop  
beq $s2, $t2, endloop
```

```
!if (a[i] < min)  
add $t0, $s2, $a0  
lw $t0, 0($t0)  
bgt $t0, $s0, nomin  
beq $t0, $s0, nomin  
!min = a[i];  
addi $s0, $t0, 0  
nomin:  
!if (a[i] > max)  
add $t0, $s2, $a0  
lw $t0, 0($t0)  
blt $t0, $s1, nomax  
beq $t0, $s1, nomax  
! max = a[i];
```

Implementing the pseudo-code

```
!int min, max, i;  
!$s0 is min, $s1 is max, $s2 is i  
!$a0 is a, $a1 is alen  
!min = a[0];  
lw $s0, 0($a0)  
!max = a[0];  
lw $s1, 0($a0)  
!for (i = 1; i < alen; i++) {  
addi $t2, $zero, 12  
addi $s2, $zero, 1  
loop:  
bgt $s2, $t2, endloop  
beq $s2, $t2, endloop
```

```
!if (a[i] < min)  
add $t0, $s2, $a0  
lw $t0, 0($t0)  
bgt $t0, $s0, nomin  
beq $t0, $s0, nomin  
!min = a[i];  
addi $s0, $t0, 0  
nomin:  
!if (a[i] > max)  
add $t0, $s2, $a0  
lw $t0, 0($t0)  
blt $t0, $s1, nomax  
beq $t0, $s1, nomax  
! max = a[i];  
addi $s1, $t0, 0  
nomax:  
!}
```


Implementing the pseudo-code

```
!int min, max, i;  
!$s0 is min, $s1 is max, $s2 is i  
!$a0 is a, $a1 is alen  
!min = a[0];  
lw $s0, 0($a0)  
!max = a[0];  
lw $s1, 0($a0)  
!for (i = 1; i < alen; i++) {  
addi $t2, $zero, 12  
addi $s2, $zero, 1  
loop:  
bgt $s2, $t2, endloop  
beq $s2, $t2, endloop  
!if (a[i] < min)  
    add $t0, $s2, $a0  
    lw $t0, 0($t0)  
    bgt $t0, $s0, nomin
```

```
    beq $t0, $s0, nomin  
    !min = a[i];  
    addi $s0, $t0, 0  
nomin:  
!if (a[i] > max)  
    add $t0, $s2, $a0  
    lw $t0, 0($t0)  
    blt $t0, $s1, nomax  
    beq $t0, $s1, nomax  
    ! max = a[i];  
    addi $s1, $t0, 0  
nomax:  
    !}  
    addi $s2, $s2, 1  
    beq $zero, $zero, loop  
endloop:  
!return max - min;
```

Implementing the pseudo-code

```
!int min, max, i;
```

!\$s0 is min, \$s1 is max, \$s2 is i

!\$a0 is a, \$a1 is alen

```
!min = a[0];
```

```
lw $s0, 0($a0)
```

```
!max = a[0];
```

```
lw $s1, 0($a0)
```

```
!for (i = 1; i < alen; i++) {
```

```
addi $t2, $zero, 12
```

```
addi $s2, $zero, 1
```

```
loop:
```

```
bgt $s2, $t2, endloop
```

```
beq $s2, $t2, endloop
```

```
!if (a[i] < min)
```

```
add $t0, $s2, $a0
```

```
lw $t0, 0($t0)
```

```
bgt $t0, $s0, nomin
```

```
beq $t0, $s0, nomin
```

```
!min = a[i];
```

```
addi $s0, $t0, 0
```

nomin:

```
!if (a[i] > max)
```

```
add $t0, $s2, $a0
```

```
lw $t0, 0($t0)
```

```
blt $t0, $s1, nomax
```

```
beq $t0, $s1, nomax
```

```
! max = a[i];
```

```
addi $s1, $t0, 0
```

nomax:

```
!}
```

```
addi $s2, $s2, 1
```

```
beq $zero, $zero, loop
```

endloop:

```
!return max - min;
```

```
nand $t0, $s0, $s0
```

```
addi $t0, $t0, 1
```

```
add $v0, $s1, $t0
```

Thanks for Attending :)