

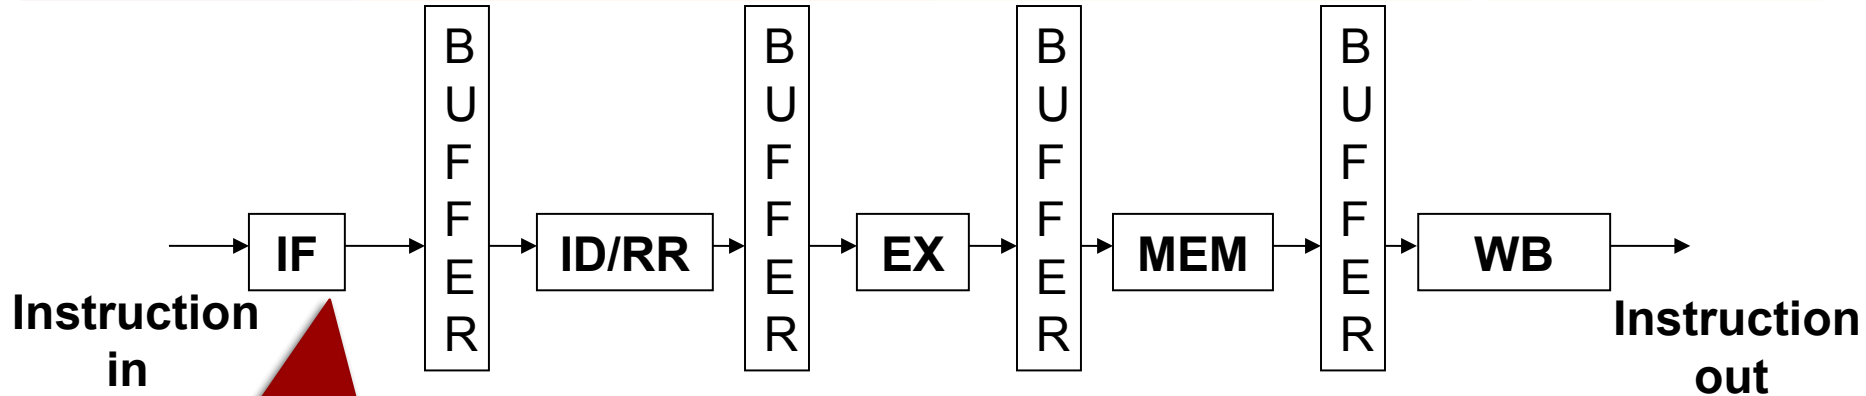


Memory Hierarchy

Chapter 9, Ramachandran and Leahy



Recall: In a paged memory system



Two memory accesses for every instruction:
- access PTE
- access instr



What can we do?



Increase cycle time

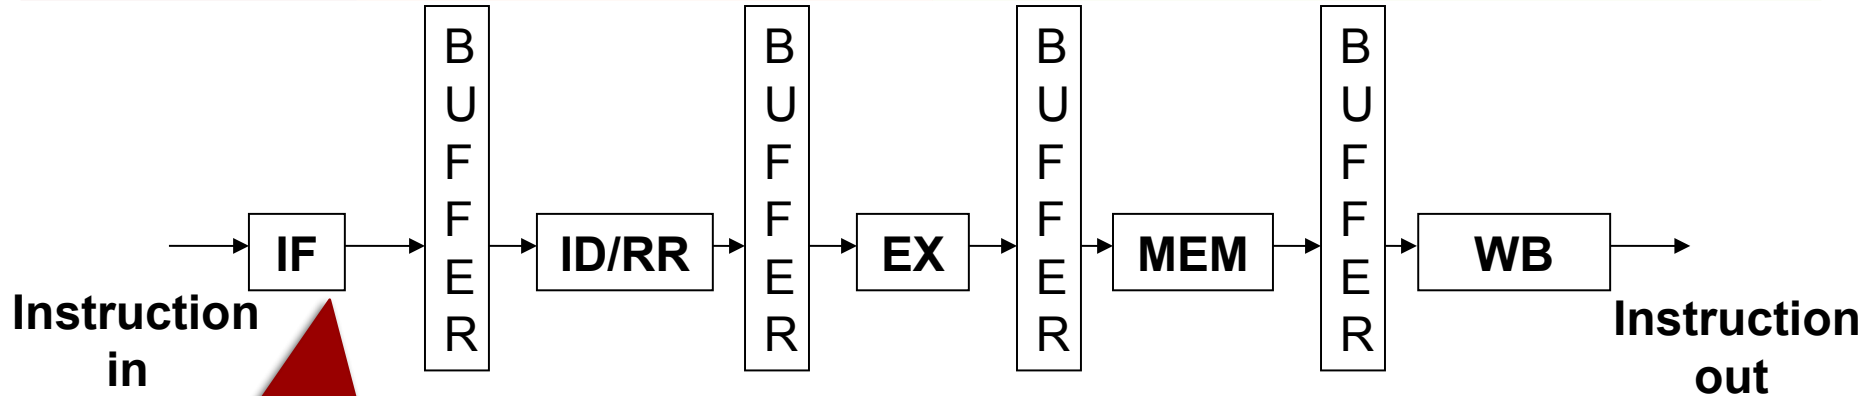


Take 2 cycles in IF



Both bad!

In a paged memory system



➤ TLB saves the day...

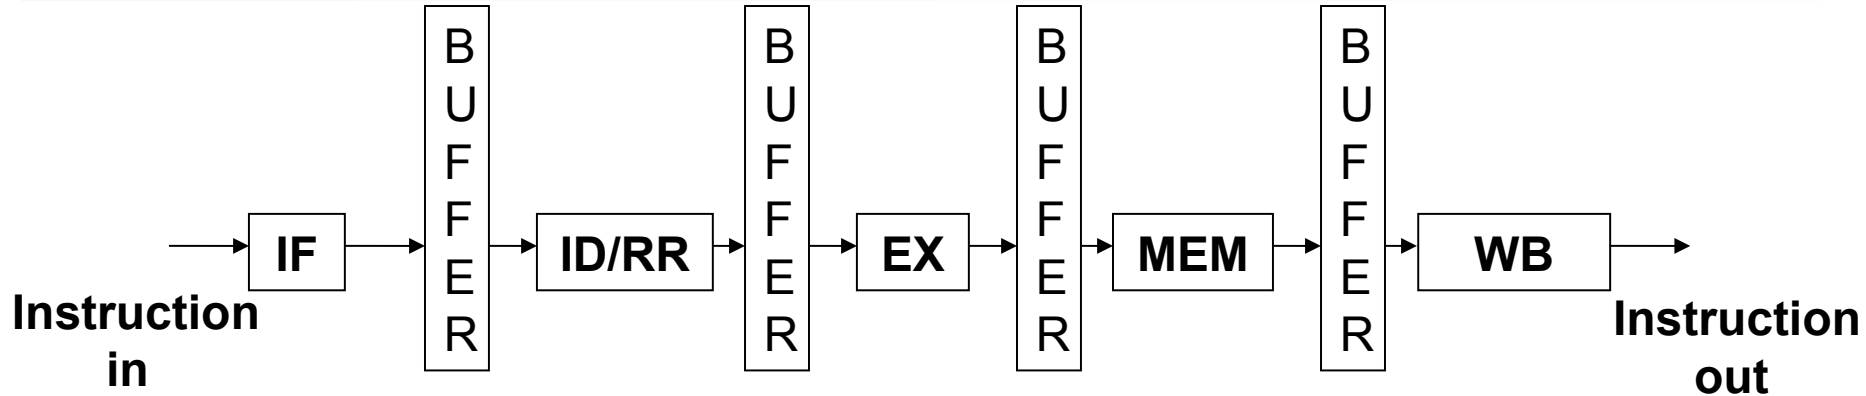
➤ PC @ TLB @ memory is only one memory access

➤ So we're back in business....

Are we really?

- What's the CPU clock cycle speed? ~1ns
- What's the memory access time? ~100ns
- It's roughly a 100:1 ratio!

What happens to the pipeline?



➤ With a 100:1 ratio, what happens in the IF stage on a memory fetch?

➤ 99 bubbles on every instruction

Not sustainable.
FAIL!

The register file and TLB...

- They're "kinda" memories, right?
- Why are they not a problem?
- TLB + Reg file ⊗ Static RAM (SRAM) technology
- Memory ⊗ Dynamic RAM (DRAM) technology
- SRAM ⊗ 6 transistors per cell ⊗ faster ⊗ bulkier
- DRAM ⊗ 1 transistor per cell ⊗ slower ⊗ denser

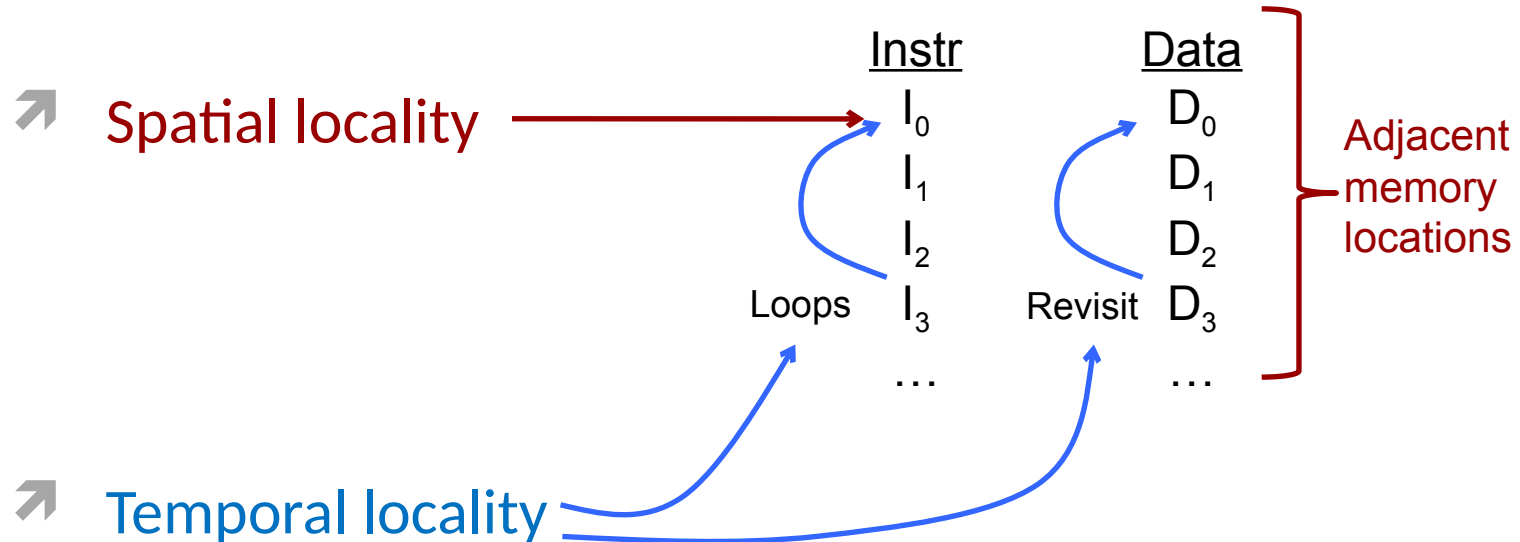
- Upshot: You can have **size** or **speed** but not both!
- This also means "memory access" cannot always be part of the CPU clock cycle time

What principle makes TLB work?

➤ Locality

- Sequentiality of instructions in the program
- Sequentiality of data structures like arrays and structs
- “Recent” translations are stashed away from in-memory page tables into high-speed hardware
- Subsequent translations become faster
 - Ⓟ no need to go to memory for most PTEs
- What if we apply the same principle to the data and instructions we need?


The principle is **locality**



How do we put these two ideas together?

- Locality
- Small & fast or large & slow storage
- © The concept of “cache”
- TLB is a special instance of caching “addresses”

The model



- How do you use
 - Spatial locality?
 - Temporal locality?

Terminologies

| | | |
|-------------------------------------|-------------------------------------|-------------|
| Hit | Hit ratio h | $h + m = 1$ |
| Miss | Miss ratio m | |
| Cycle time | T_c cache access time | |
| | T_m memory access time | |
| Memory access time | Same for SRAM Different for DRAM | |
| Memory cycle time | | |
| Miss penalty | T_m | |
| Effective memory access time (EMAT) | $EMAT = T_c + T_m * m$ | |

Current-day memory hierarchy

$$EMAT_i = T_i + M_i * EMAT_{i+1}$$

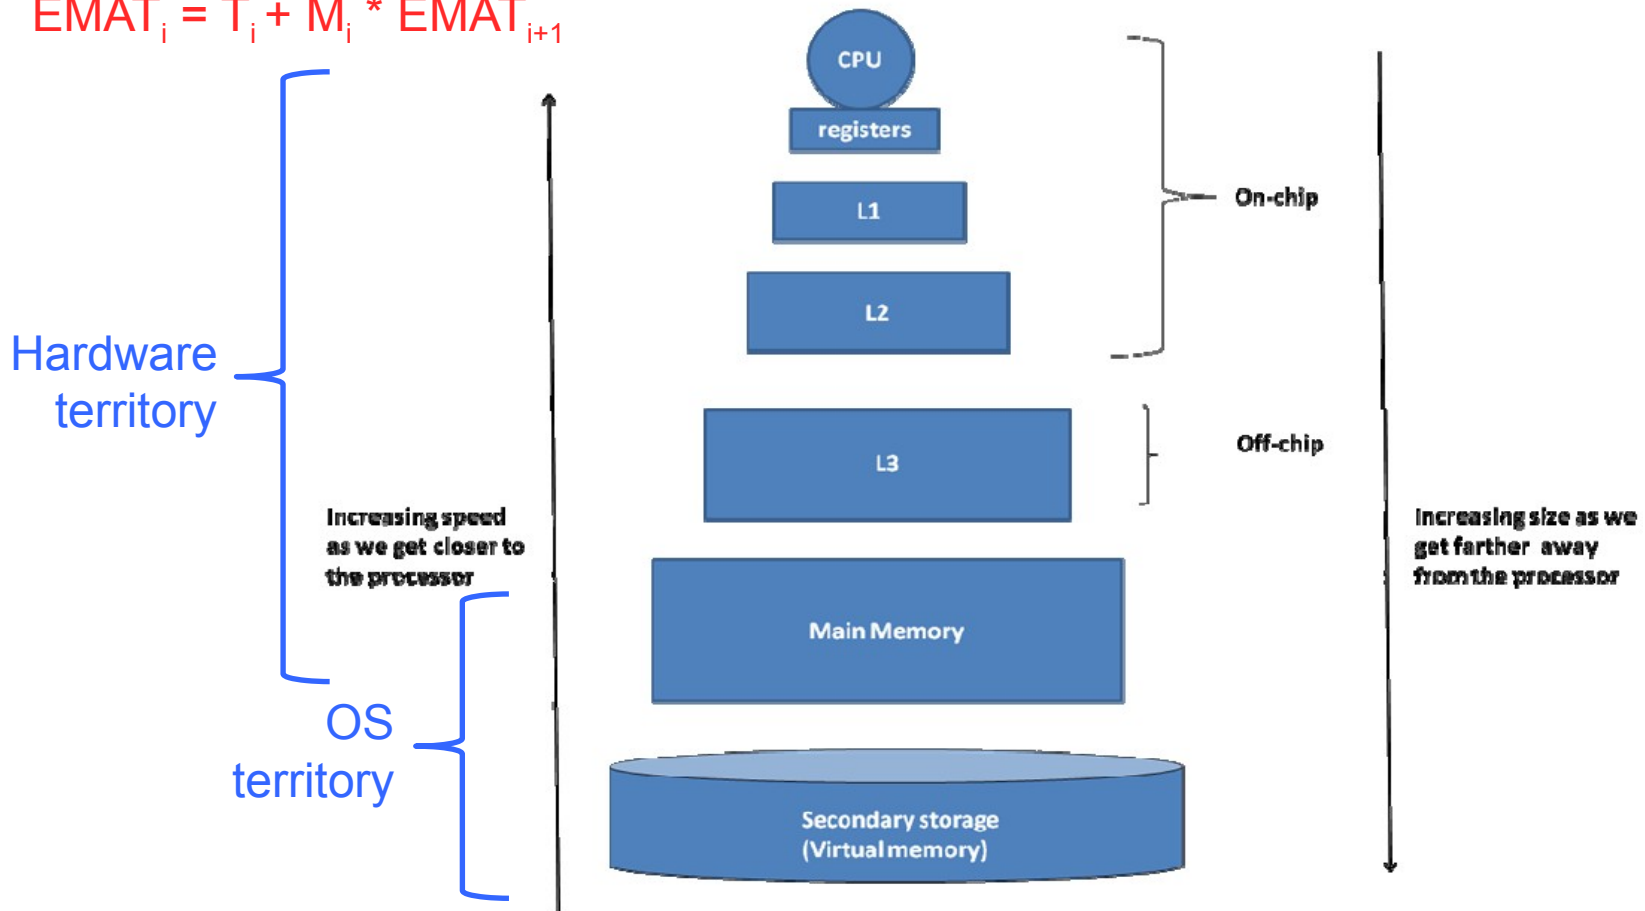


Figure 9.2: The entire memory hierarchy stretching from processor registers to the virtual memory.

Calculating EMAT

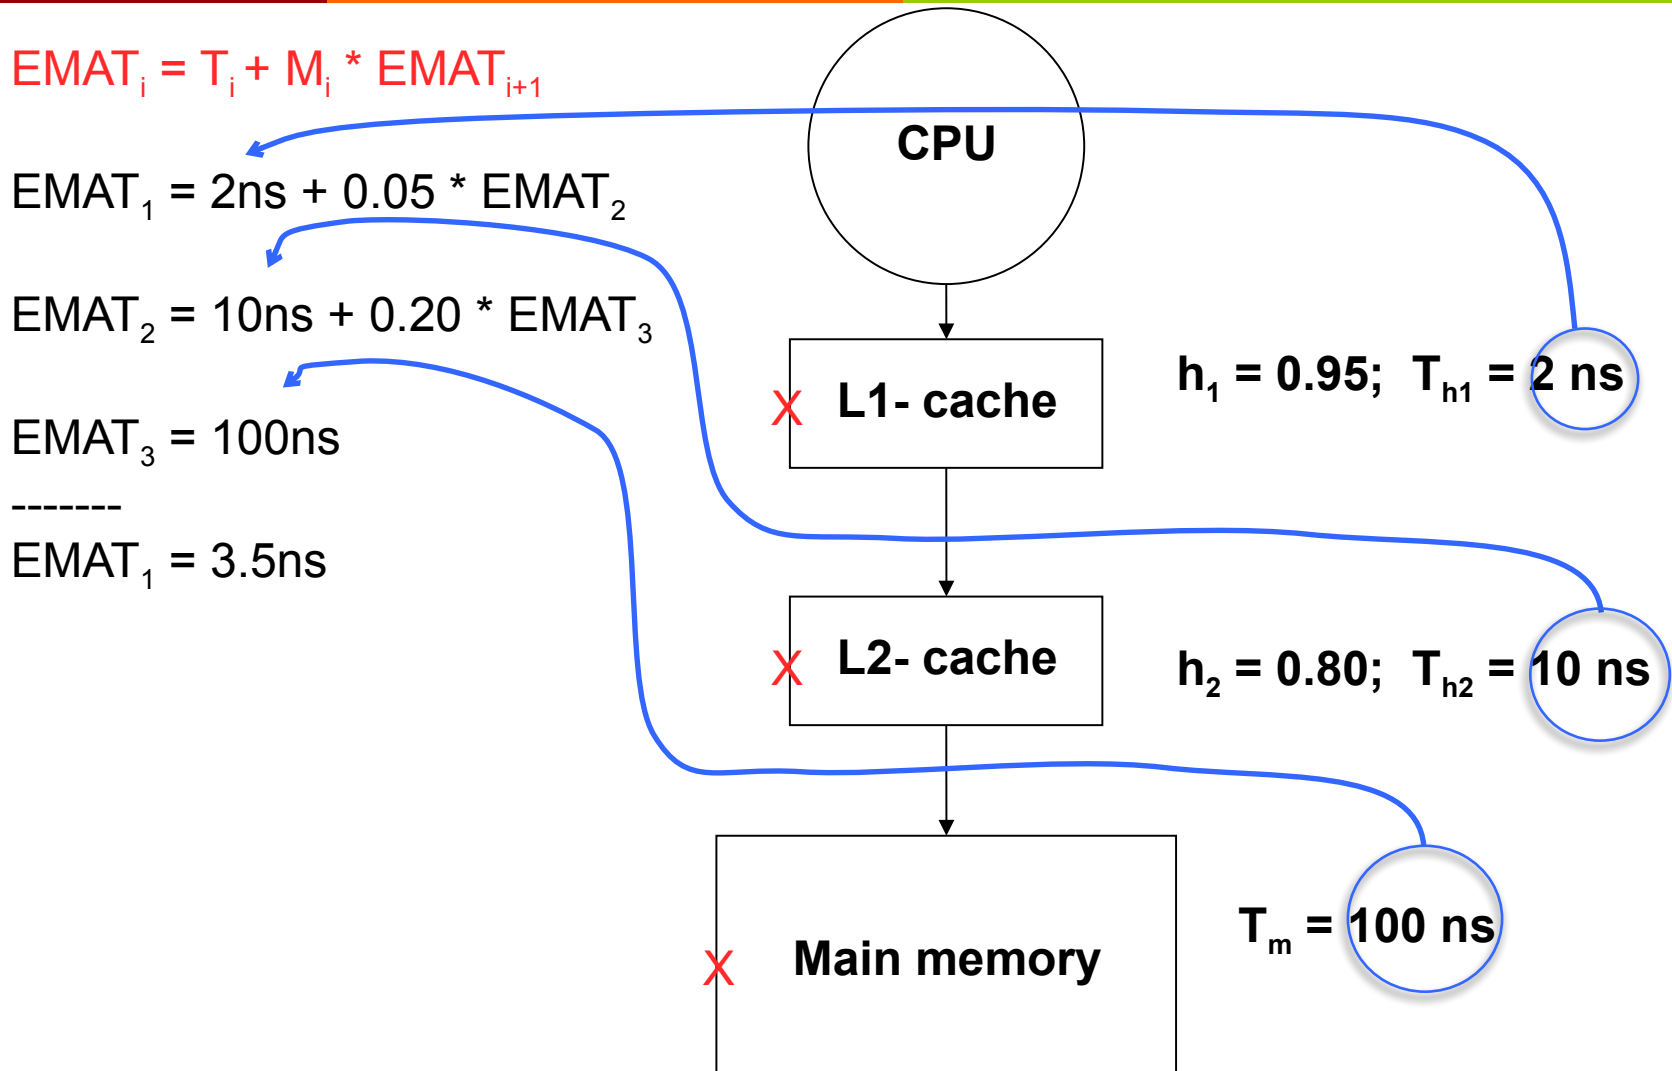
$$EMAT_i = T_i + M_i * EMAT_{i+1}$$

$$EMAT_1 = 2ns + 0.05 * EMAT_2$$

$$EMAT_2 = 10ns + 0.20 * EMAT_3$$

$$EMAT_3 = 100ns$$

$$EMAT_1 = 3.5ns$$



Calculate the EMAT?

You have a one-level cache that has a hit rate of 90% and an access time of 10ns; your main memory has an access time of 100 ns.

- A. 11ns
- B. 110ns
- C. 12ns
- D. 20ns



Today's number is 30,333

$$EMAT_i = T_i + M_i * EMAT_{i+1}$$

$$EMAT_1 = 10ns + 0.10 * EMAT_2$$

$$EMAT_2 = 100ns$$

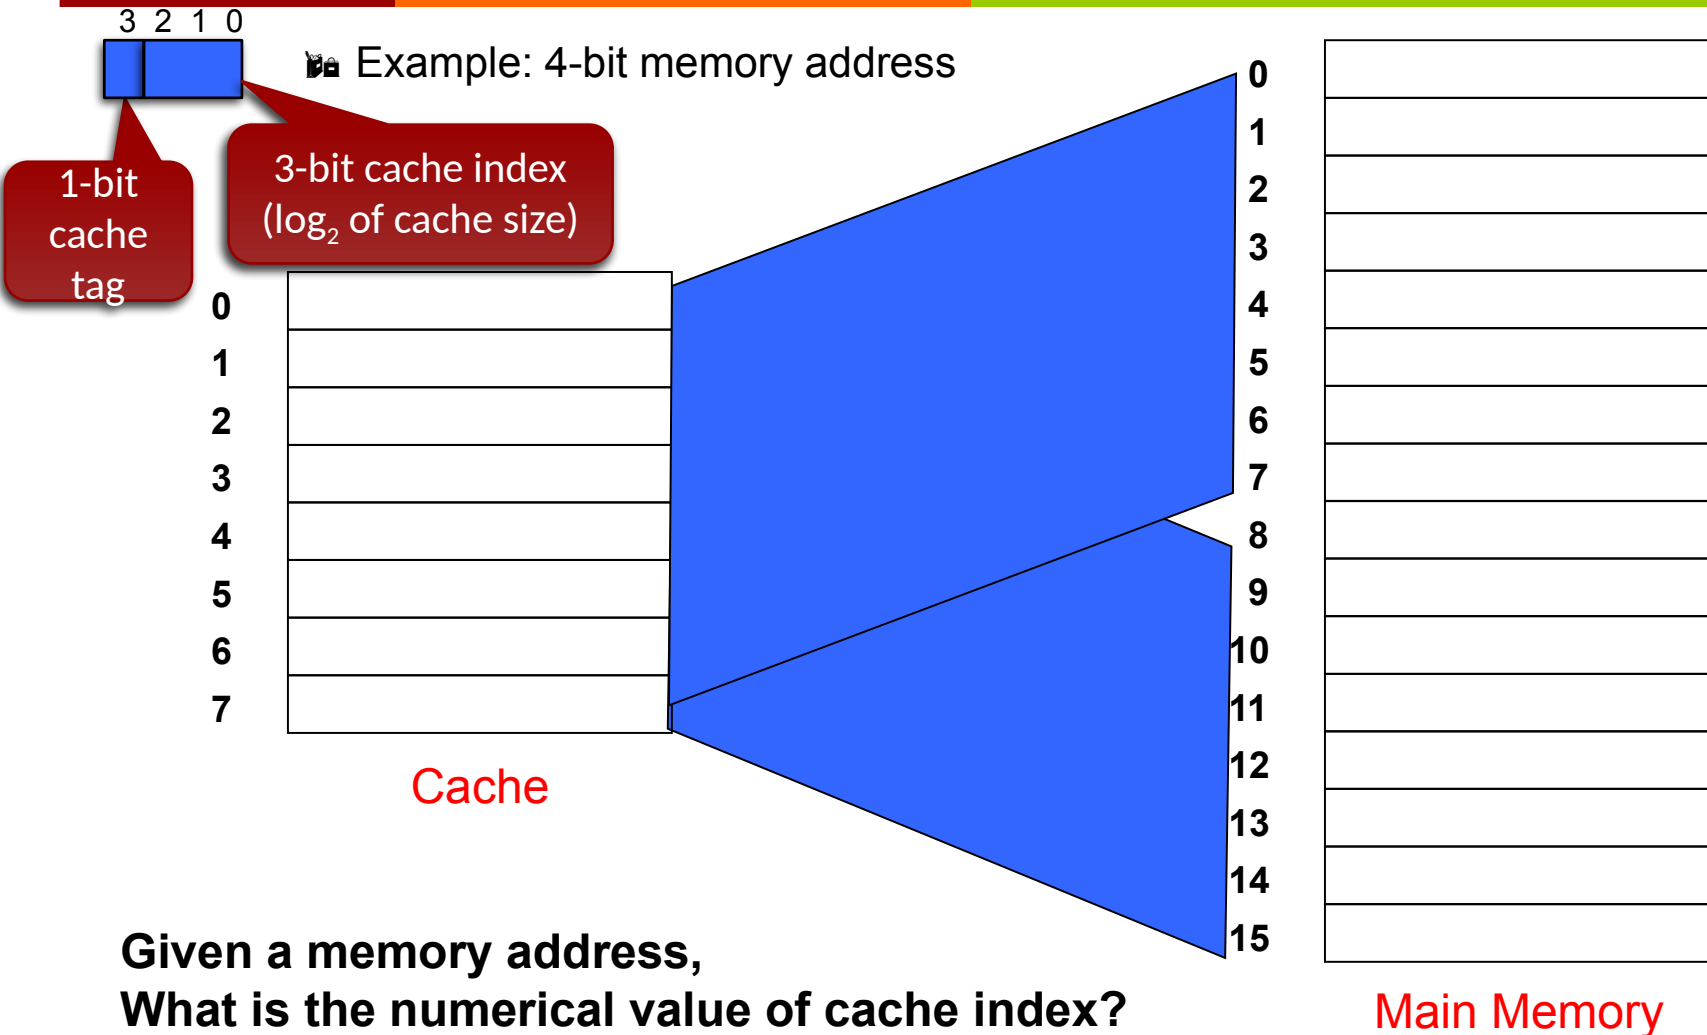
$$EMAT_1 = 10ns + 0.10 * 100ns$$

$$EMAT_1 = 20ns$$

Cache Organizations

- Direct mapped
- Fully associative
- Set associative

Direct mapped cache




Types of misses

- Compulsory (first time a block is brought in)
 - Capacity (cache is full)
 - Conflict (vying for space in a full set in the cache)
- Compulsory > Capacity > Conflict
- Another way to think about Conflict misses: Conflict misses are misses that would not occur if the cache were fully associative with LRU replacement. Fully associative caches can't have conflict misses by definition.

Example

Memory references: 0, 1, 2, 3, 1, 3, 0, 8, 0, 9, 10



These were
compulsory
misses.


This is the first
time the
address is
begin brought
into the cache.

| | |
|---|-----------|
| 0 | mem loc 0 |
| 1 | mem loc 1 |
| 2 | mem loc 2 |
| 3 | mem loc 3 |
| 4 | empty |
| 5 | empty |
| 6 | empty |
| 7 | empty |

| | |
|----|--|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |
| 6 | |
| 7 | |
| 8 | |
| 9 | |
| 10 | |
| 11 | |
| 12 | |
| 13 | |
| 14 | |
| 15 | |

Example

Memory references: 0, 1, 2, 3, 1, 3, 0, 8, 0, 9, 10



This is also a **compulsory** miss.


This is the first time address 8 has been brought into the cache.

| | |
|---|-------------|
| 0 | mem loc 0 8 |
| 1 | mem loc 1 |
| 2 | mem loc 2 |
| 3 | mem loc 3 |
| 4 | empty |
| 5 | empty |
| 6 | empty |
| 7 | empty |

| | |
|----|--|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |
| 6 | |
| 7 | |
| 8 | |
| 9 | |
| 10 | |
| 11 | |
| 12 | |
| 13 | |
| 14 | |
| 15 | |

Example

Memory references: 0, 1, 2, 3, 1, 3, 0, 8, 0, 9, 10



This is now a **conflict** miss.

There are two memory addresses mapped to the same cache cell.

| | |
|---|--|
| 0 | mem loc 0 8 0 |
| 1 | mem loc 1 |
| 2 | mem loc 2 |
| 3 | mem loc 3 |
| 4 | empty |
| 5 | empty |
| 6 | empty |
| 7 | empty |

| | |
|----|--|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |
| 6 | |
| 7 | |
| 8 | |
| 9 | |
| 10 | |
| 11 | |
| 12 | |
| 13 | |
| 14 | |
| 15 | |

Example

Memory references: 0, 1, 2, 3, 1, 3, 0, 8, 0, 9, 10

misses

hits

miss

| | | | |
|---|-----|---|---|
| 3 | 2 | 1 | 0 |
| 1 | 001 | | |

🛒 4-bit memory address

3-bit index

9 MOD 8 is 1


```
index = memory_address
        mod cache_size
```

| | |
|---|--------------------------|
| 0 | mem loc 0 8/0 |
| 1 | mem loc 1 |
| 2 | mem loc 2 |
| 3 | mem loc 3 |
| 4 | empty |
| 5 | empty |
| 6 | empty |
| 7 | empty |

| | |
|----|--|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |
| 6 | |
| 7 | |
| 8 | |
| 9 | |
| 10 | |
| 11 | |
| 12 | |
| 13 | |
| 14 | |
| 15 | |

How do we disambiguate data?

We use the rest of the memory address as the tag to label the data in the cache



4-bit memory address

3-bit index

This remaining bit determines if cache data is for memory address 0 or 8, etc.

| | tag | data |
|---|-----|-------------|
| 0 | 1 | mem loc 0 8 |
| 1 | 0 | mem loc 1 |
| 2 | 0 | mem loc 2 |
| 3 | 0 | mem loc 3 |
| 4 | | empty |
| 5 | | empty |
| 6 | | empty |
| 7 | | empty |


| | |
|----|--|
| 0 | |
| 1 | |
| 2 | |
| 3 | |
| 4 | |
| 5 | |
| 6 | |
| 7 | |
| 8 | |
| 9 | |
| 10 | |
| 11 | |
| 12 | |
| 13 | |
| 14 | |
| 15 | |

How do we know data is valid?


- At power-up, a cache may contain garbage!
- Tags help disambiguate, not validate

| | valid | tag | data |
|---|-------|-----|-------|
| 0 | 1 | 1 | loc 8 |
| 1 | 1 | 0 | loc 1 |
| 2 | 1 | 0 | loc 2 |
| 3 | 1 | 0 | loc 3 |
| 4 | 0 | X | empty |
| 5 | 0 | X | empty |
| 6 | 0 | X | empty |
| 7 | 0 | X | empty |

Fields in a Direct Mapped Cache




What type of locality affects direct mapped cache performance?

- A. Spatial locality
- B. Temporal locality
- C. Virtual locality
-  D. Spatial and temporal locality

What does the cache entry contain?

You have a 64 entry direct-mapped cache for a memory with 16-bit addresses and 16-bit words (like LC-3).

- A. 1 bit valid flag, 6 bit tag, 10 bit data
-  B. 1 bit valid flag, 10 bit tag, 16 bit data
- C. 1 bit valid flag, 10 bit tag, 6 bit data
- D. 2 bit valid flag, 12 bit tag, 16 bit data

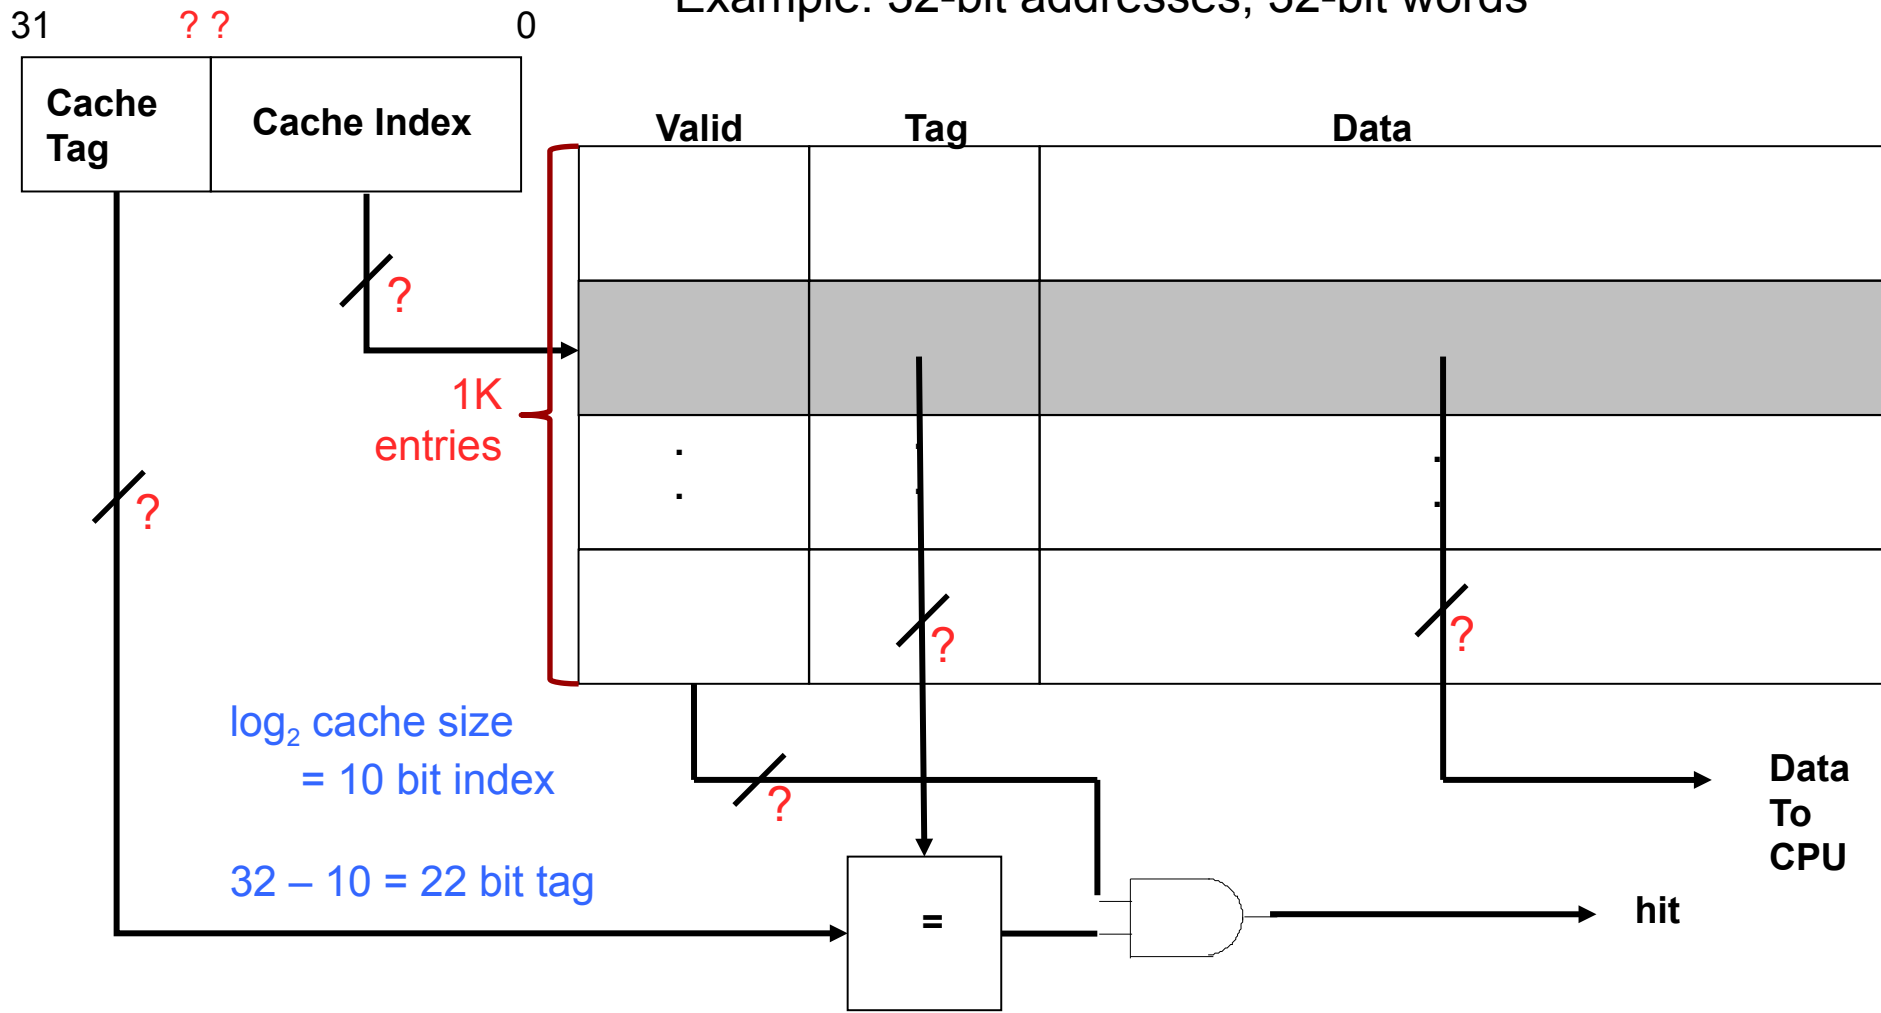
One bit is required for the valid flag

The cache index is going to be 6 because $\log_2(64)=6$, so tag will be $16-6=$ **10 bits**

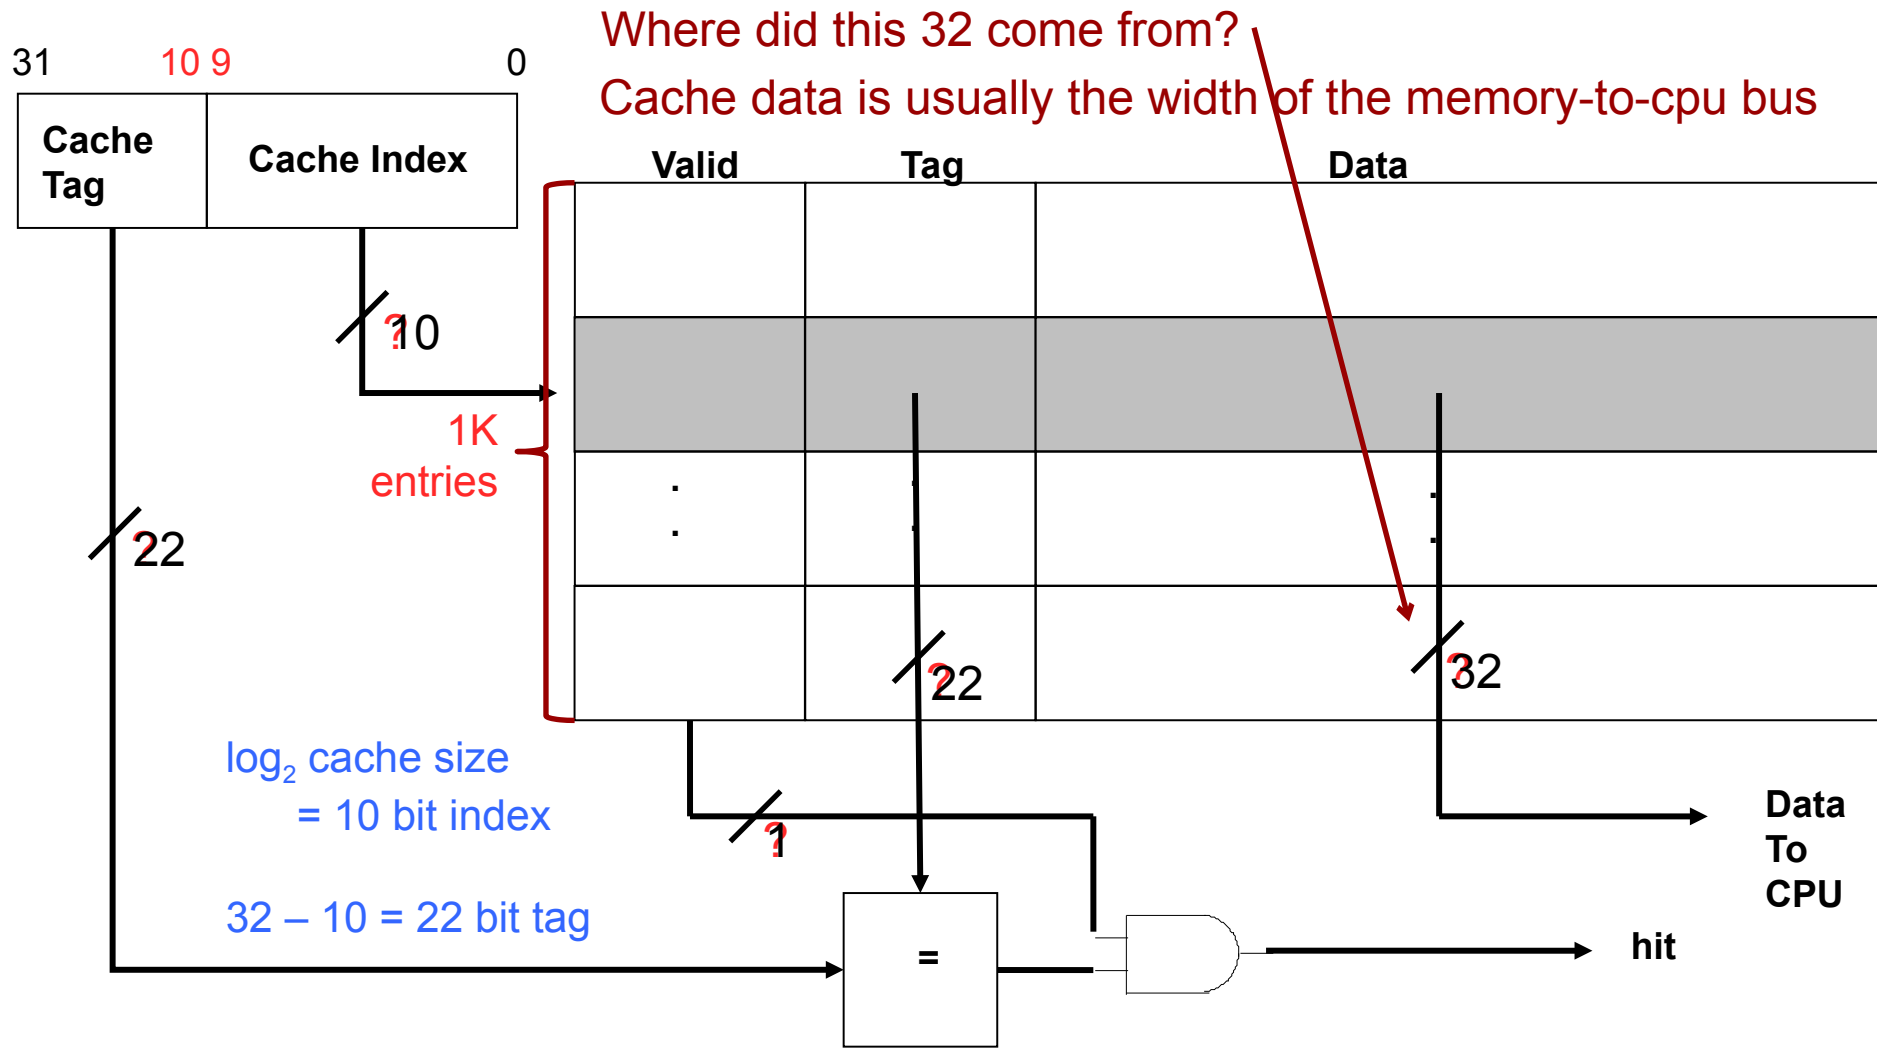
Memory word is 16 bits, so data has to be a **16 bits** (or a multiple)

Hardware

Example: 32-bit addresses, 32-bit words



Hardware



Interpreting memory addresses

Memory address

| Cache Tag | Cache Index |
|-----------|-------------|
|-----------|-------------|

- $\text{index} = \text{memory addr} \bmod \text{cache size}$
 - $\text{MemAddr} = 8 \text{ } \textcircled{P} \text{ index} = 0$
 - $\text{MemAddr} = 0 \text{ } \textcircled{P} \text{ index} = 0$
- $\text{number of tag bits} = \text{memory addr size} - \text{cache index size}$

Why not the other way around?

\textcircled{P} use the low bits for cache tag?

Index first, tag last?

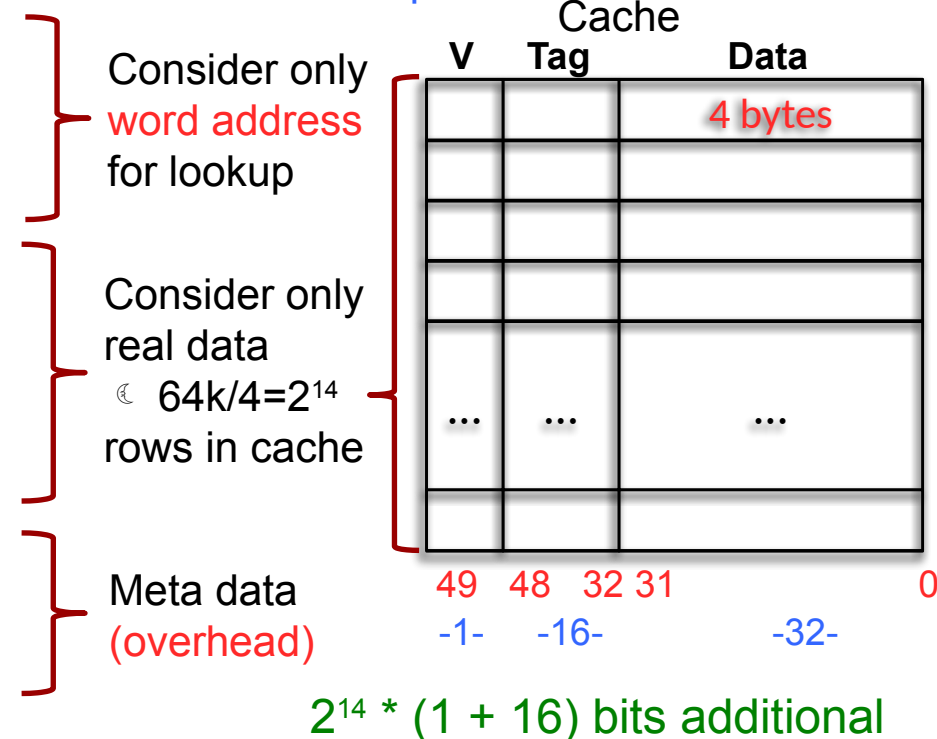

| | valid | tag | data | | valid | tag | data | | valid | tag | data | | valid | tag | data |
|---|-------|-----|-------|---|-------|-----|---------|---|-------|-----|---------|---|-------|-----|---------|
| 0 | 1 | 0 | loc 0 | 0 | 1 | 1 | loc 0 1 | 0 | 1 | 1 | loc 0 1 | 0 | 1 | 1 | loc 0 1 |
| 1 | 0 | X | empty | 1 | 0 | X | empty | 1 | 1 | 0 | loc 2 | 1 | 0 | 1 | loc 2 3 |
| 2 | 0 | X | empty | 2 | 0 | X | empty | 2 | 0 | X | empty | 2 | 0 | X | empty |
| 3 | 0 | X | empty | 3 | 0 | X | empty | 3 | 0 | X | empty | 3 | 0 | X | empty |
| 4 | 0 | X | empty | 4 | 0 | X | empty | 4 | 0 | X | empty | 4 | 0 | X | empty |
| 5 | 0 | X | empty | 5 | 0 | X | empty | 5 | 0 | X | empty | 5 | 0 | X | empty |
| 6 | 0 | X | empty | 6 | 0 | X | empty | 6 | 0 | X | empty | 6 | 0 | X | empty |
| 7 | 0 | X | empty | 7 | 0 | X | empty | 7 | 0 | X | empty | 7 | 0 | X | empty |

Access location 0 Access location 1 Access location 2 Access location 3

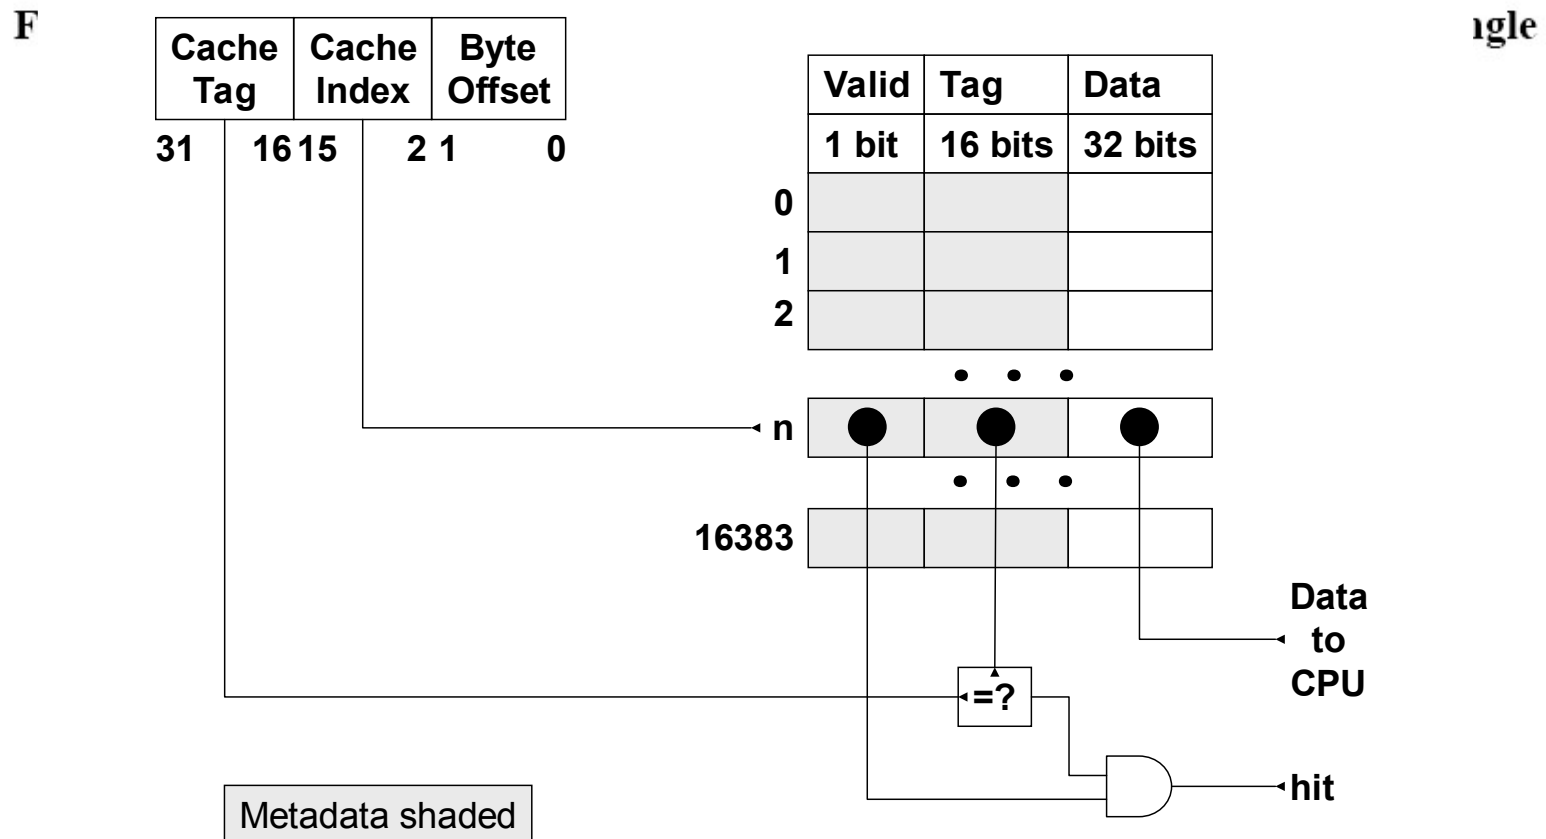
Cache occupancy if we switch index and tag is
BAD!! ⑥ **loss of spatial locality**

Example

- Let us consider the design of a direct-mapped cache for a realistic memory system.
- Assume that the CPU generates a 32-bit **byte-addressable** memory address.
- Each memory word contains **4 bytes**.
- A memory access **brings** in a **full word** into the cache.
- The **direct-mapped** cache is **64K bytes** in size (this is the amount of data that can be stored in the cache), with each cache entry containing **one word** of data.
- Compute the **additional storage space** needed for the valid bits and the tag fields of the cache.



Memory address interpretation when single cache block contains multiple bytes




A direct-mapped cache




- A. Has a many-to-one mapping between a memory location and a cache location
- B. Allows a memory location to be en-cached wherever there is space in the cache
- C. Is so-called because there is a directory associated with the contents of the cache
- D. Is usually much smaller than any other type of cache organization


Today's number is 69,727

In a direct-mapped cache with a t -bit tag

- A. There is a 1-bit tag comparator for each cache line
- B. There is a t -bit tag comparator for each cache line
- C. There is a 1-bit tag comparator for the entire cache
-  D. There is a t -bit tag comparator for the entire cache
- E. What is a comparator?

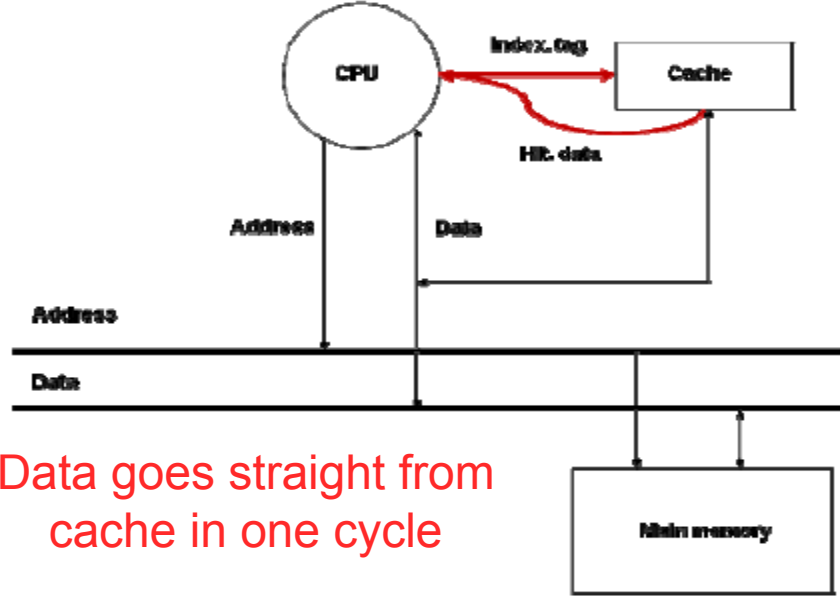
Pipelined processor with caches





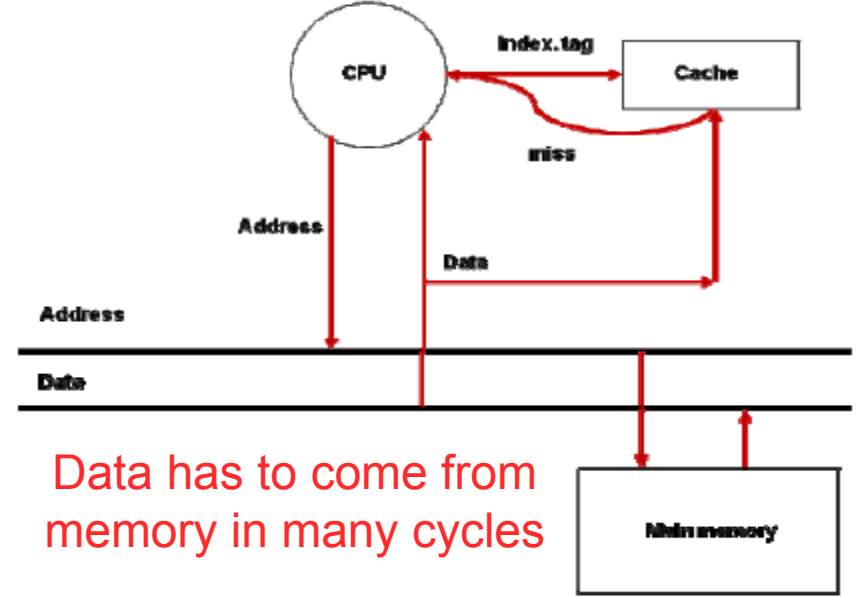
Data goes straight from
cache in one cycle

(a) Read Hit



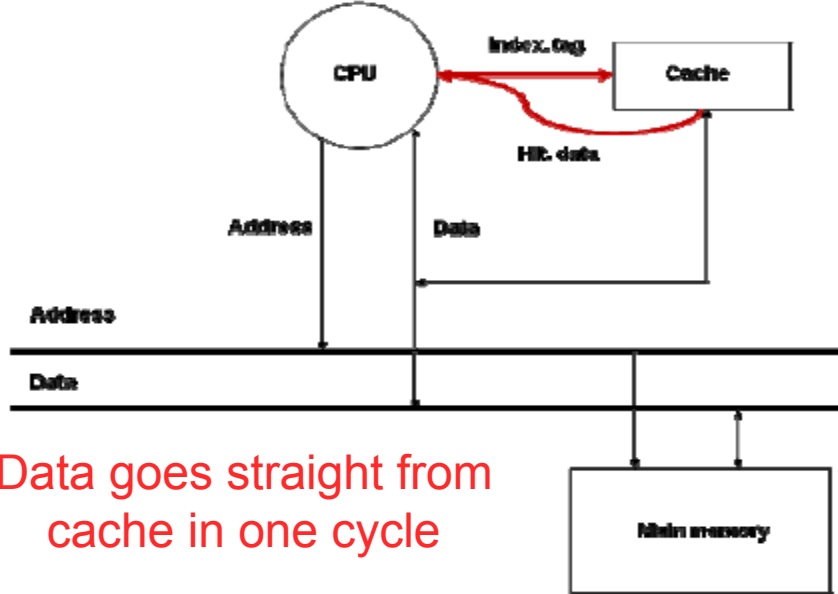
Data goes straight from
cache in one cycle

(a) Read Hit




Data has to come from
memory in many cycles

(b) Read miss



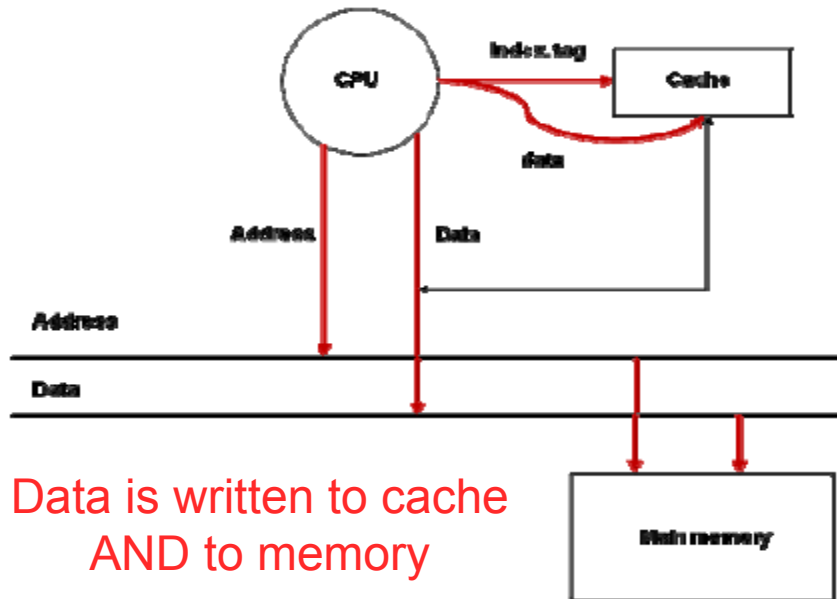
Data goes straight from cache in one cycle

(a) Read Hit



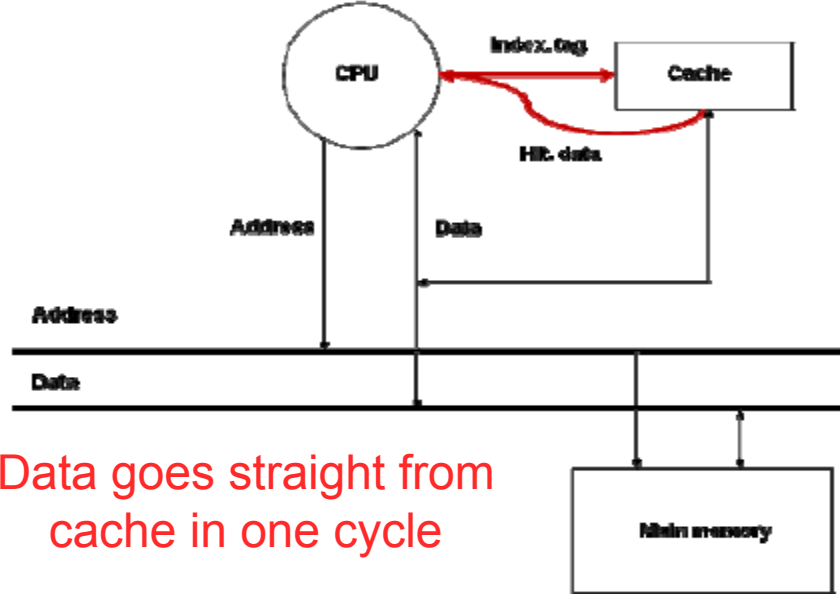
Data has to come from memory in many cycles

(b) Read miss



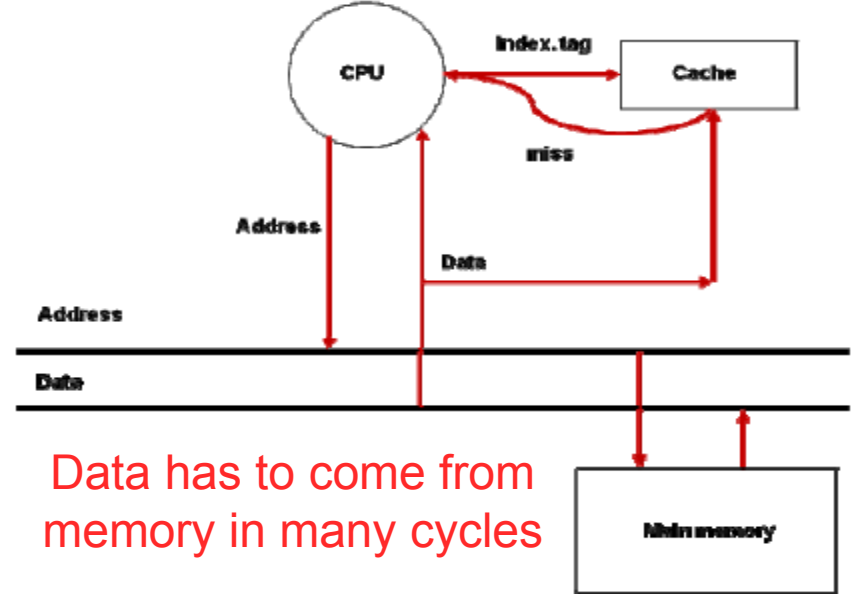
Data is written to cache AND to memory

(c) Write-through




Data goes straight from cache in one cycle

(a) Read Hit



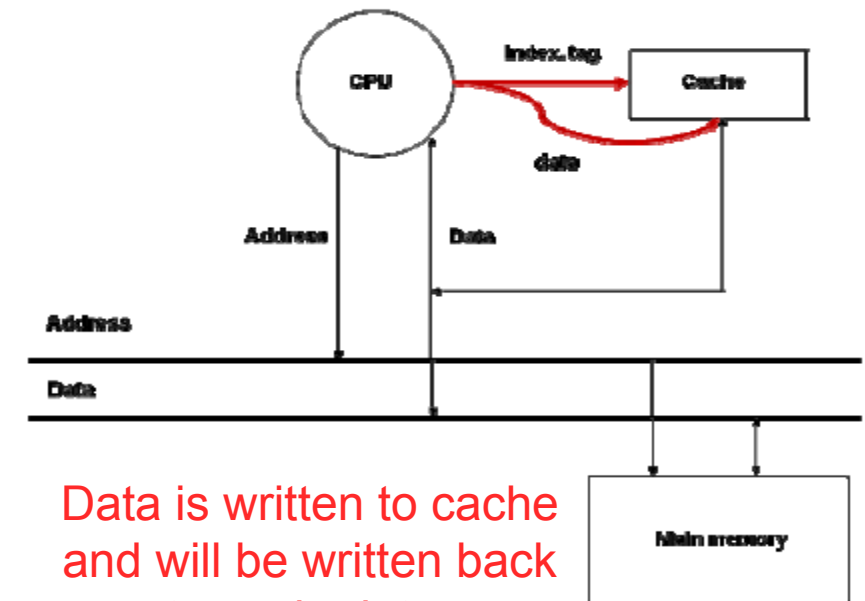
Data has to come from memory in many cycles

(b) Read miss



Data is written to cache AND to memory

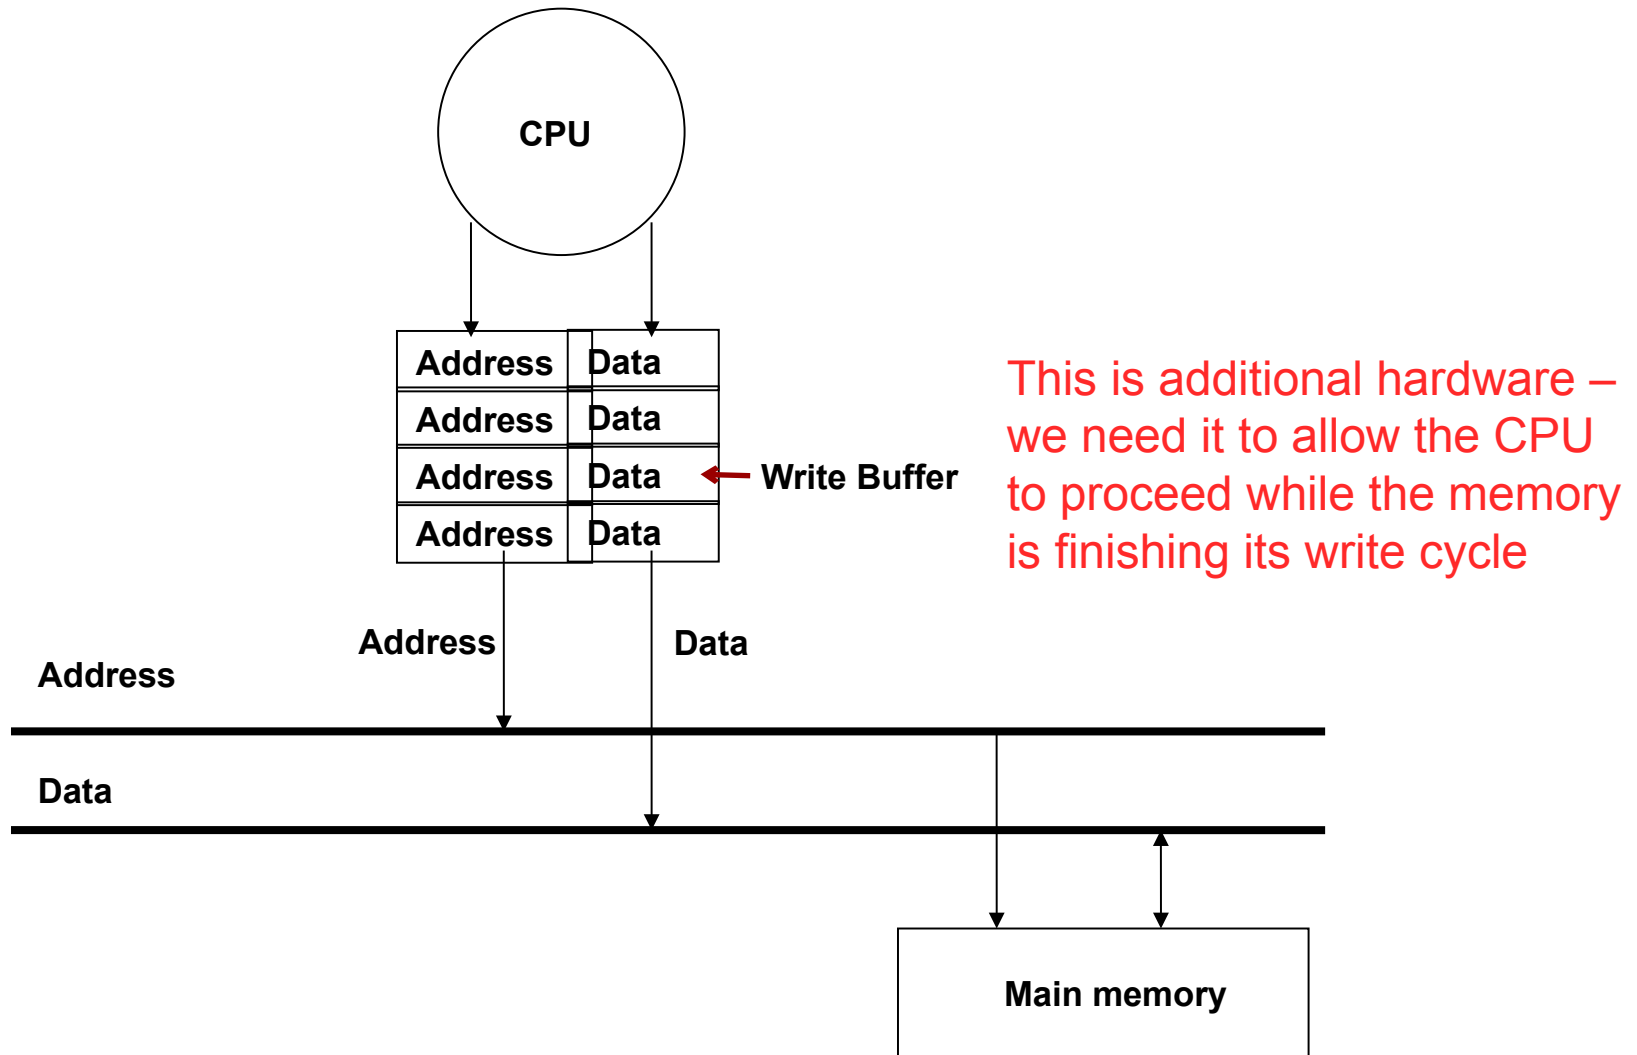
(c) Write-through



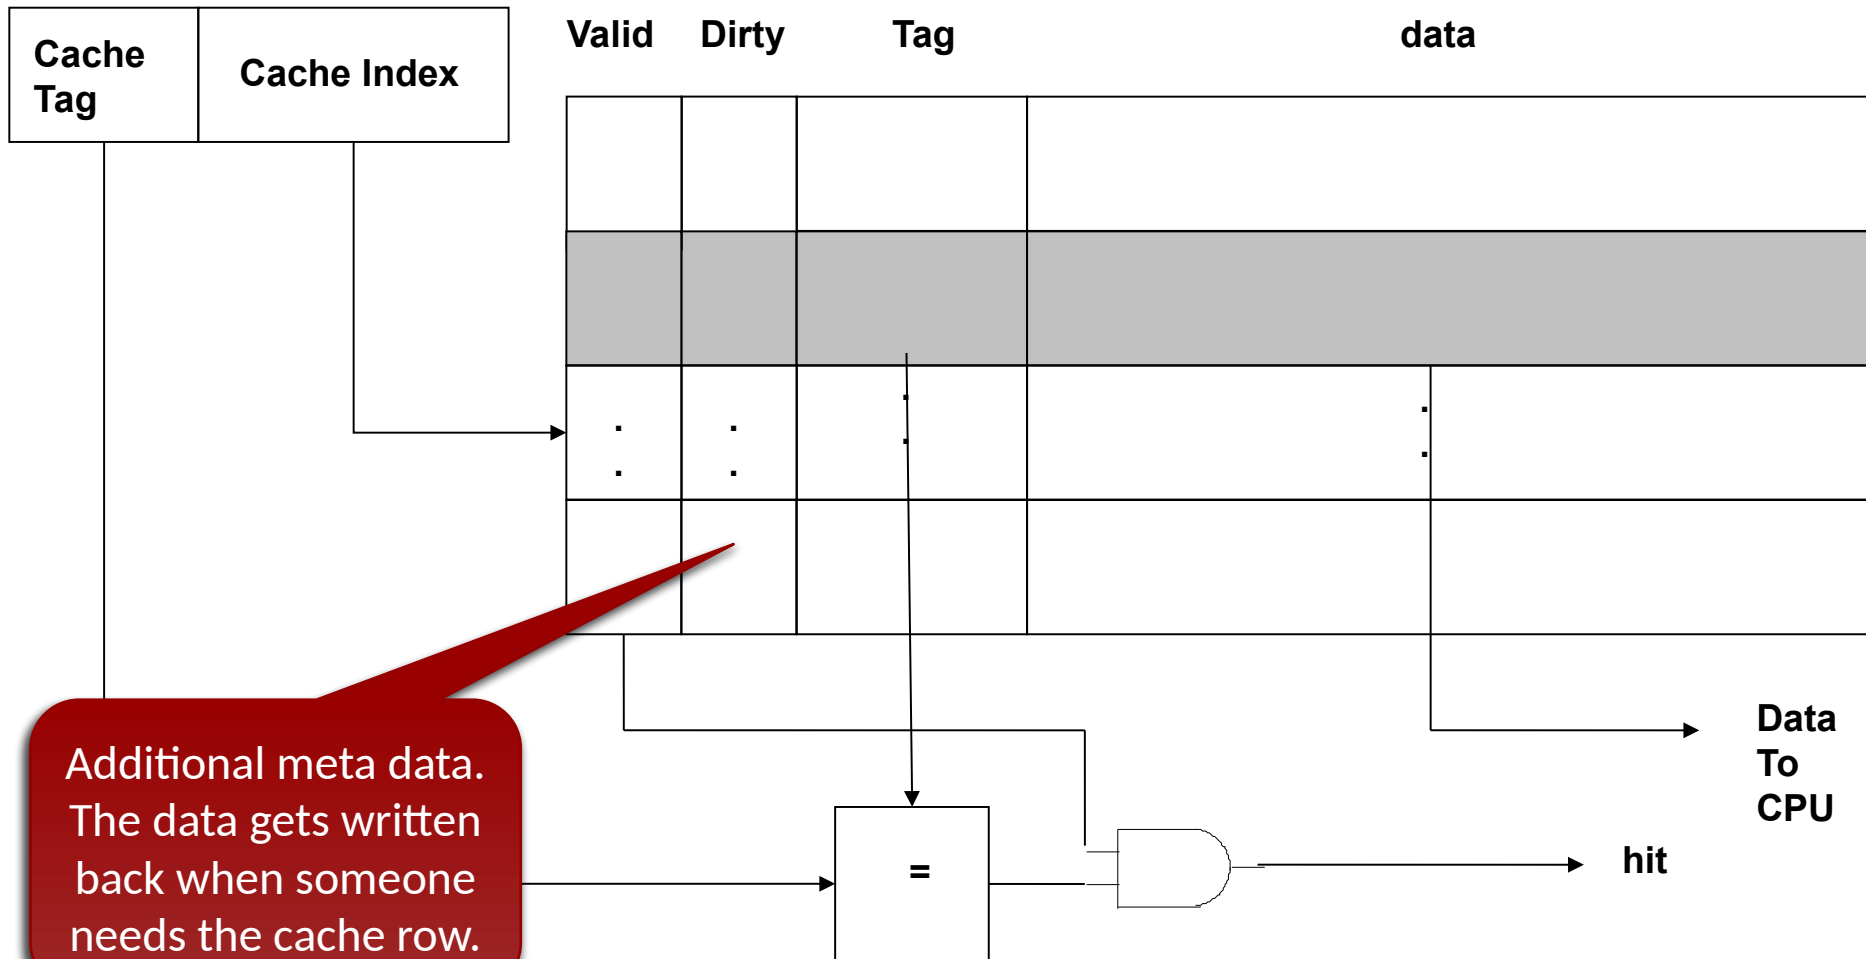
Data is written to cache and will be written back to cache later

(d) Write-back

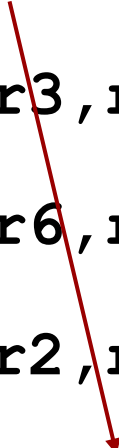
Write-through operation



Write back cache



Read miss stalls



```
I1: ld  r1,a;      r1 <- memory at a
I2: add r3,r4,r5; r3 <- r4 + r5
I3: and r6,r7,r8; r6 <- r7 & r8
I4: add r2,r4,r5; r2 <- r4 + r5
I5: add r2,r1,r2; r2 <- r1 + r2
```

- We can treat a read-miss in MEM in a similar fashion as we did previously with registers and the busy bits
- MEM can reset the busy bit for r1 when it sees the read complete for I1 (instead of waiting for WB as we did before)

Execution time with caches

~~Execution time = $N * CPI_{Avg} * \text{cycle time}$~~

$$CPI_{eff} = CPI_{Avg} + \text{Memory-stalls}_{Avg}$$

$$\text{Execution time} = N * CPI_{eff} * \text{cycle time}$$

$$\text{Execution time} = N * (CPI_{Avg} + \text{M-stalls}_{Avg}) * \text{cycle time}$$

$$\text{Memory-stalls}_{Avg} = \text{misses per instruction}_{Avg} * \text{miss-penalty}_{Avg}$$

$$\text{Total memory stalls} = N * \text{Memory-stalls}_{Avg}$$

The effective CPI is...

Average CPI = 1.5

Average cache miss per instruction = 3%

Miss penalty = 20

- A. 1.8
- B. 2.1
- C. 21.5
- D. 7.5
- E. No clue

$$\text{CPI}_{\text{eff}} = 1.5 + (3\% \text{ of } 20) = 1.5 + 0.6 = 2.1 \text{ CPI}$$

Example

- Consider a pipelined processor that has an average CPI of 1.8 without accounting for memory stalls.
 - I-Cache has a hit rate of 95%
 - D-Cache has a hit rate of 98%.
- Assume that memory reference instructions account for 30% of all the instructions executed.
 - 80% are loads
 - 20% are stores
- On average
 - read-miss penalty is 20 cycles and the
 - write-miss penalty is 5 cycles.

Compute the effective CPI of the processor accounting for the memory stalls.

Solution

Cost of instruction misses:

$$\begin{aligned} &= \text{I-cache miss rate} * \text{read miss penalty} \\ &= (1 - 0.95) * 20 \\ &= 1 \text{ cycle per instruction} \end{aligned}$$

Cost of data read misses:

$$\begin{aligned} &= \% \text{ memory reference instructions} \\ &\quad * \text{fraction that are loads} * \text{D-cache miss rate} \\ &\quad * \text{read miss penalty} \\ &= 0.3 * 0.8 * (1 - 0.98) * 20 \\ &= 0.096 \text{ cycles per instruction} \end{aligned}$$

Cost of data write misses:

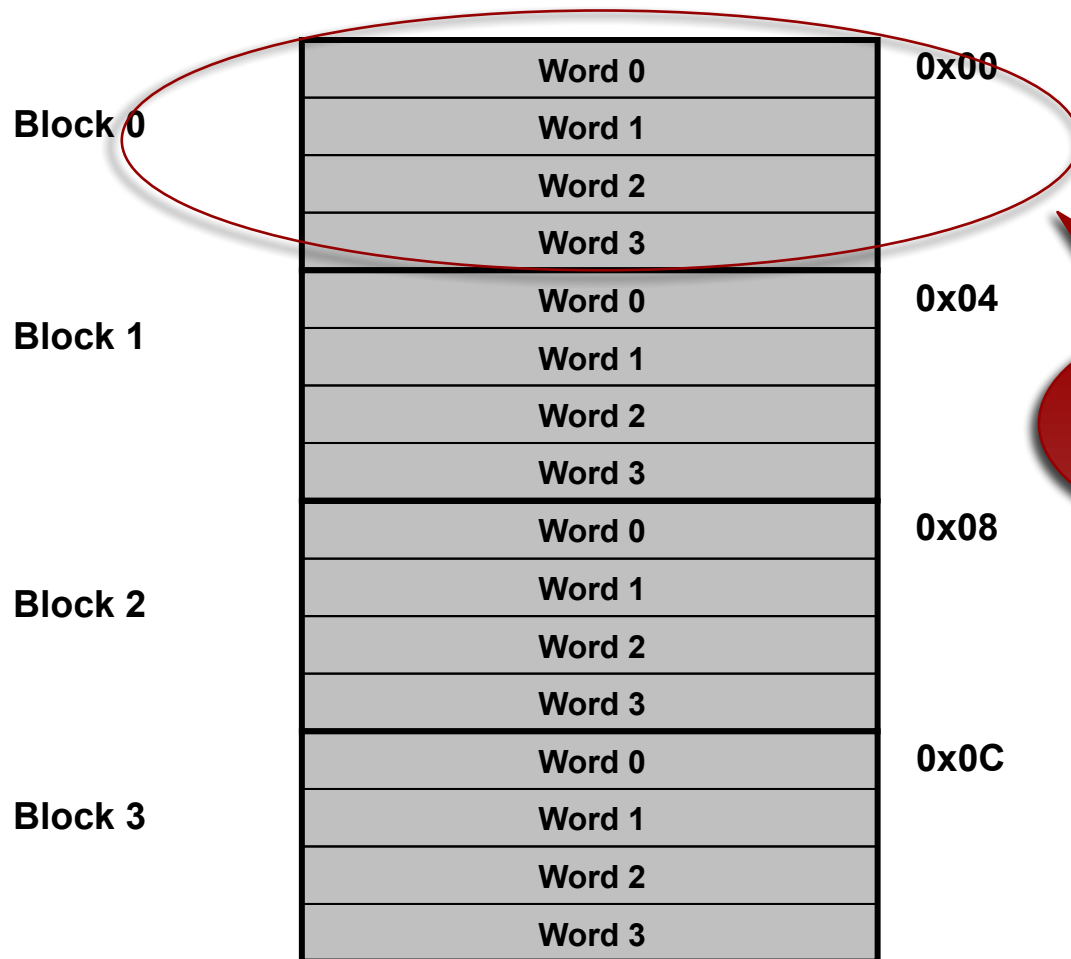
$$\begin{aligned} &= \% \text{ memory reference instructions} \\ &\quad * \text{fraction that are stores} * \text{D-cache miss rate} \\ &\quad * \text{write miss penalty} \\ &= 0.3 * 0.2 * (1 - 0.98) * 5 \\ &= 0.006 \text{ cycles per instruction} \end{aligned}$$

$$\begin{aligned} \text{CPI}_{\text{eff}} &= \text{CPI}_{\text{avg}} + \text{cost of I-cache misses} + \text{cost of D-cache misses} \\ &= 1.8 + 1 + (0.096 + 0.006) = 2.902 \end{aligned}$$

How to improve cache efficiency

- Further exploit spatial locality
 - Bring more from memory into cache at a time
- Better organization
 - Exploit working set concept

Exploiting Spatial Locality




On miss bring in
a whole block!

Interpreting memory address

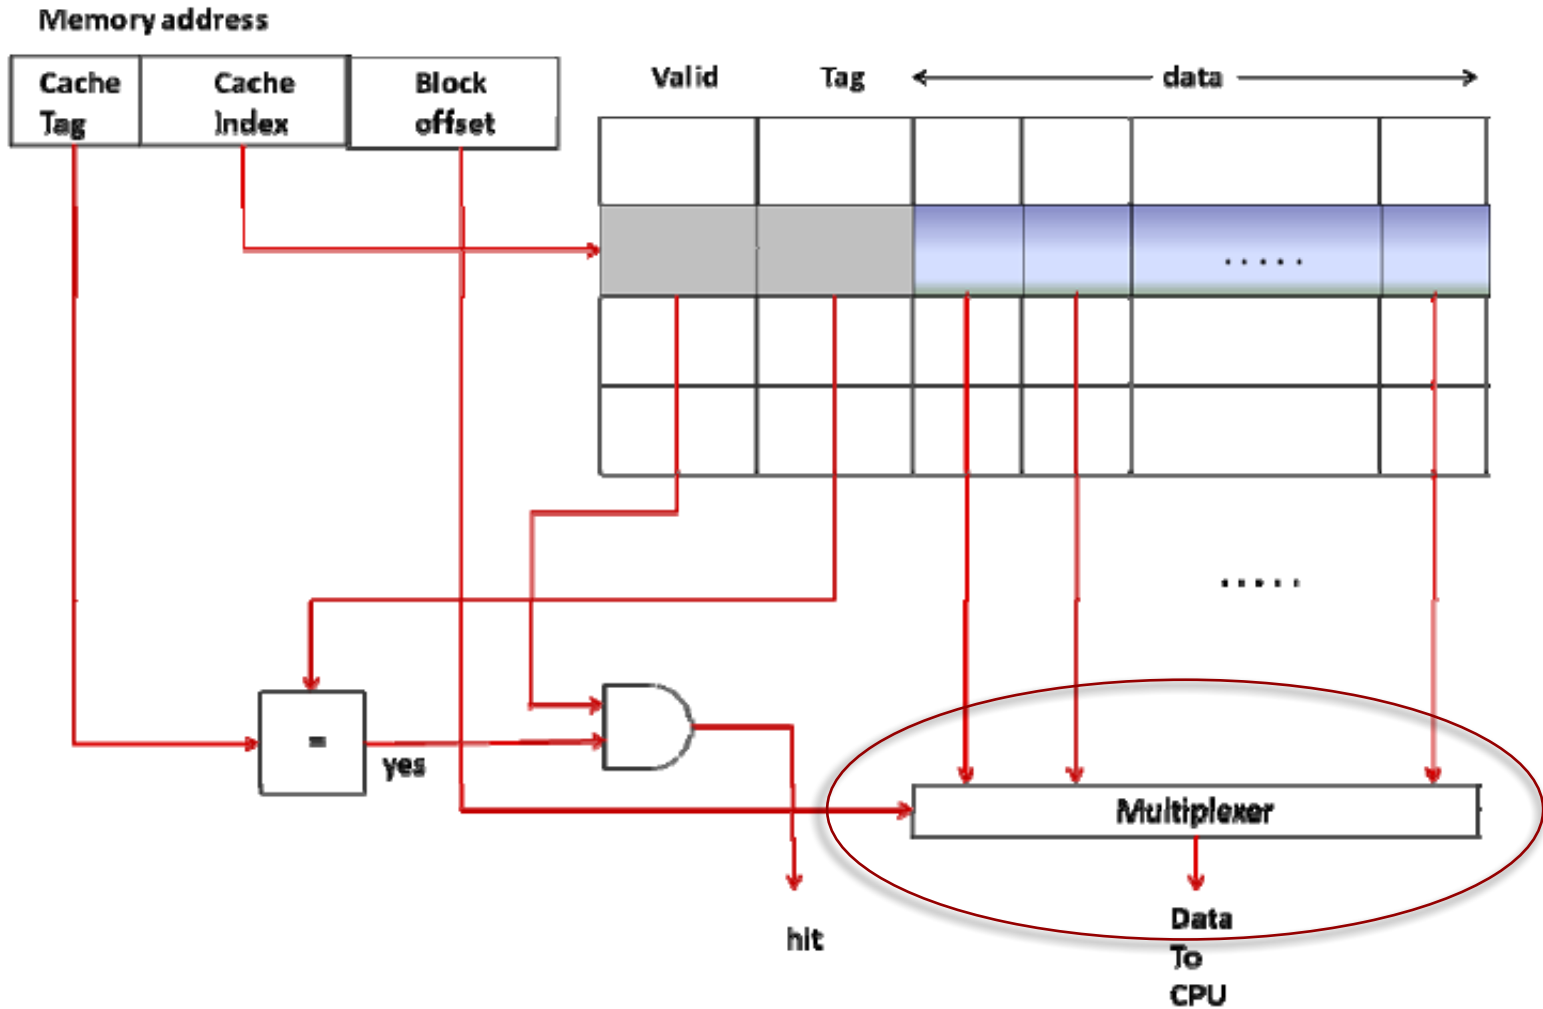
S = Size of cache; B = Block size; L = lines in cache

| Cache Tag | Cache Index | Block Offset |
|-----------|-------------|--------------|
|-----------|-------------|--------------|

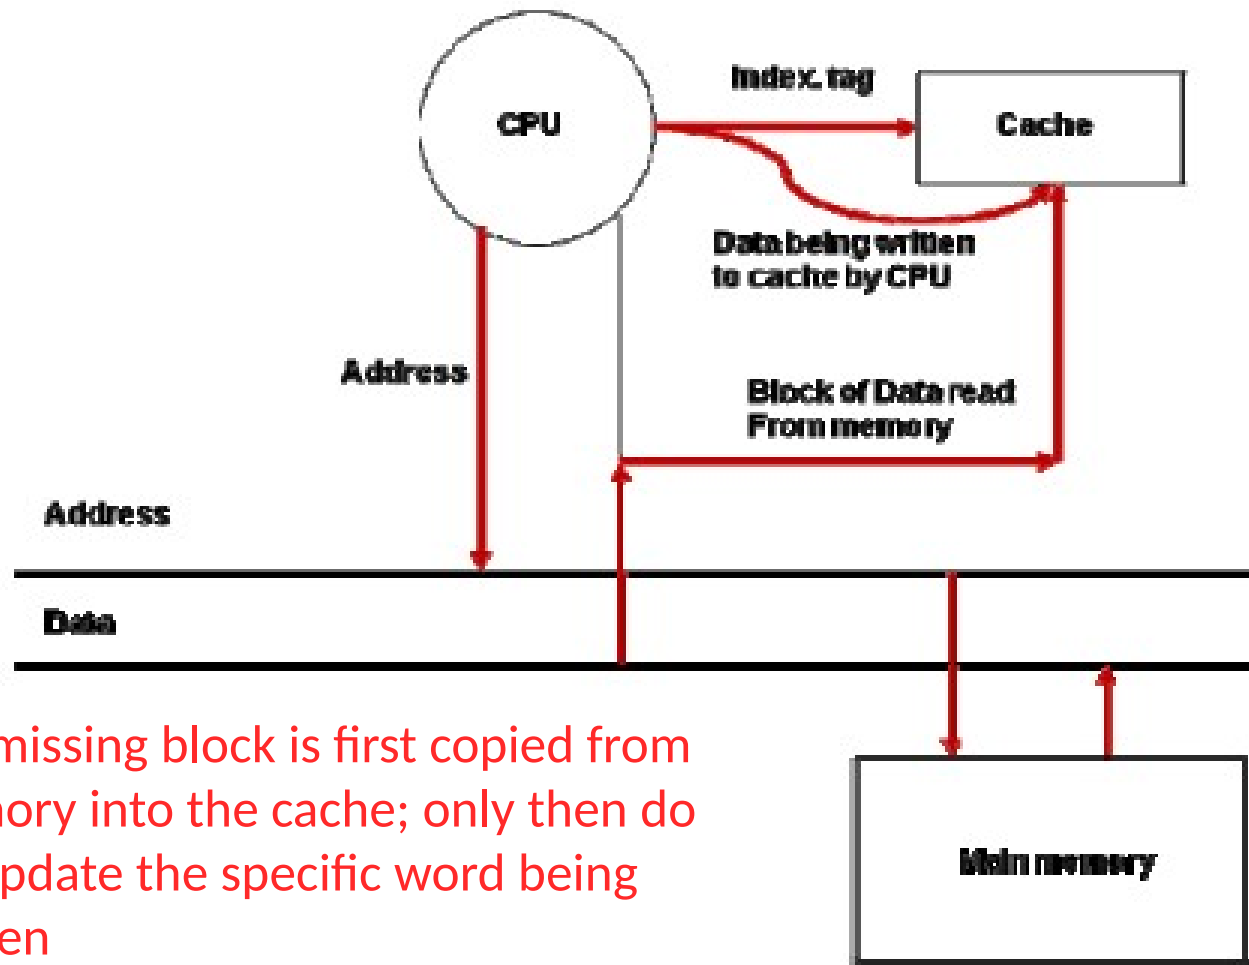


This is really just
the word
number within
the block

Multi-word cache organization




Write miss with multiple word cache



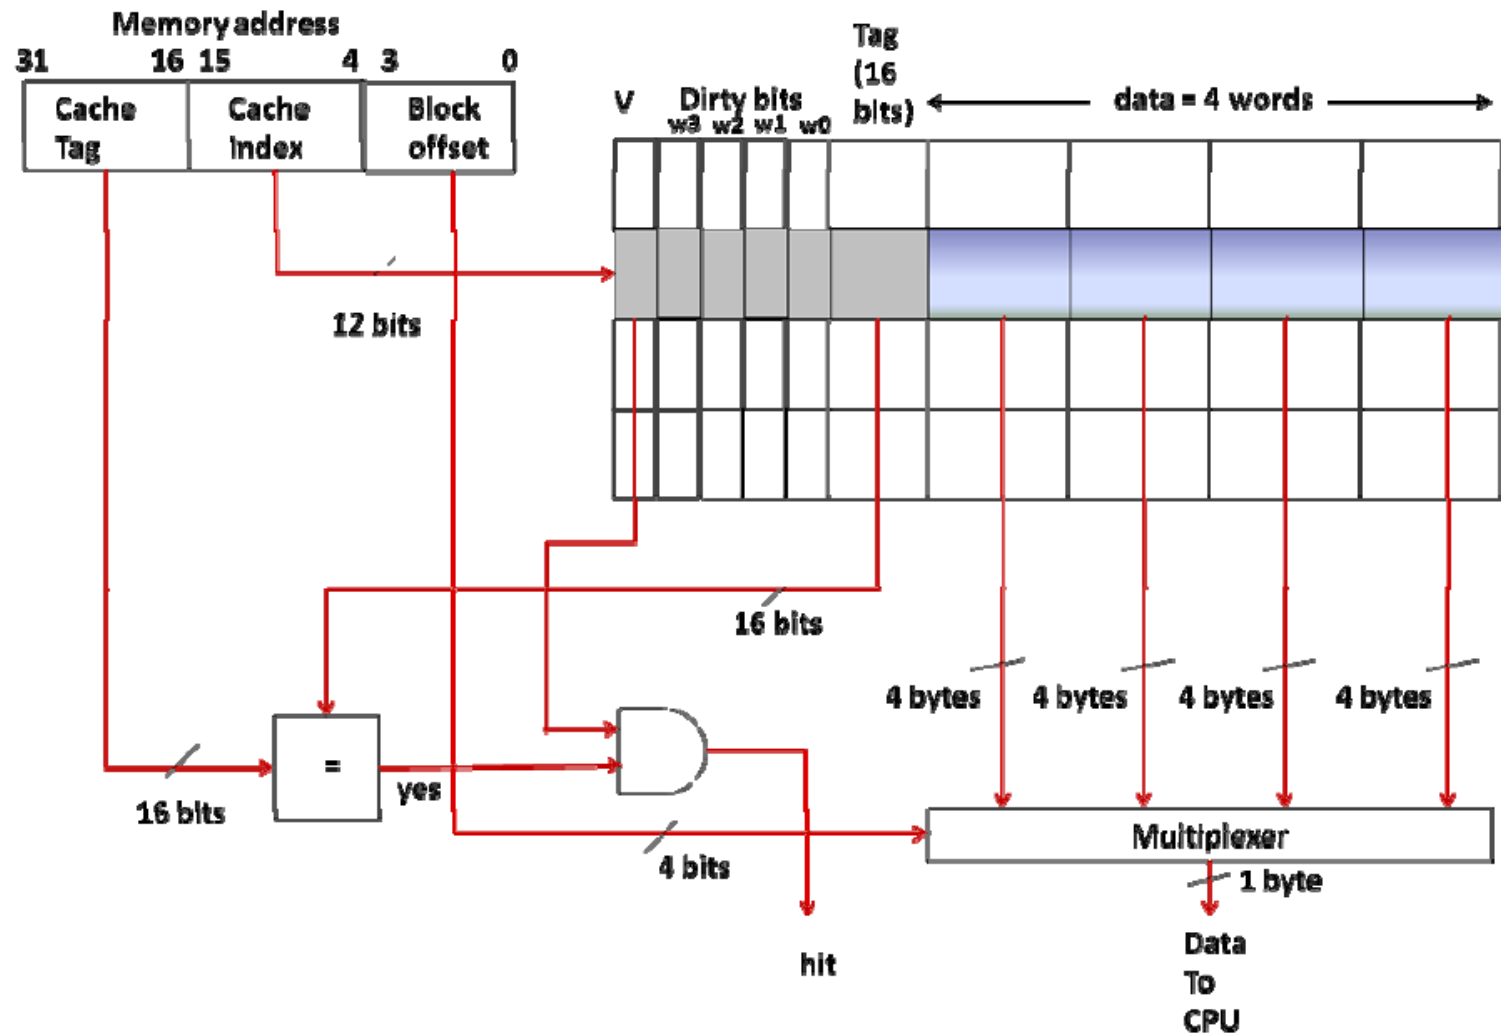
Multiple word cache example

➤ Direct-mapped cache

- 32-bit byte-addressable memory address
- Each memory word contains 4 bytes
- Block size = 4 words (16 bytes)
- A memory access brings in a block
- 64K byte cache.
- Write back cache with a dirty bit per word



Implementing a multi word cache




Increased block size?

- Exploits more spatial locality
- Reduces miss rate



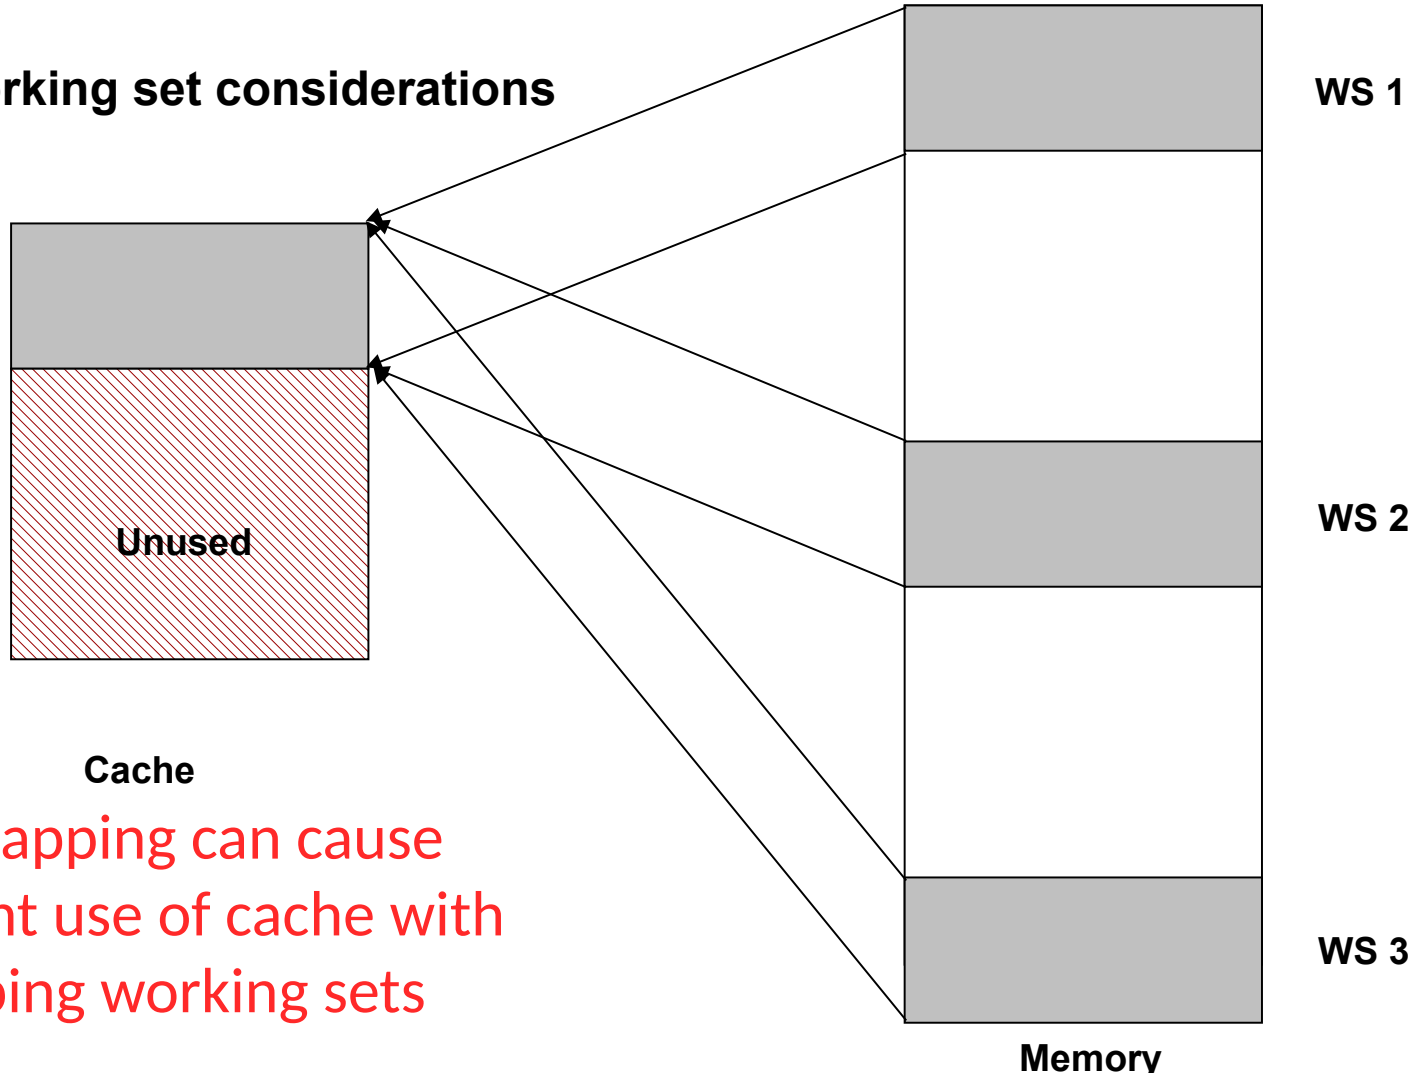
Increased block size?

- There is a point where things get worse
- When the working set changes, larger blocks have to be fetched
- Memory can only transfer so fast and it can become the bottleneck



Working set considerations

Working set considerations



Direct mapping can cause inefficient use of cache with overlapping working sets

What would be best?

- Allow any memory block to be brought into any cache block
- This is similar to be able to bring in a virtual page into any available physical page frame

Ⓟ Fully associative mapping

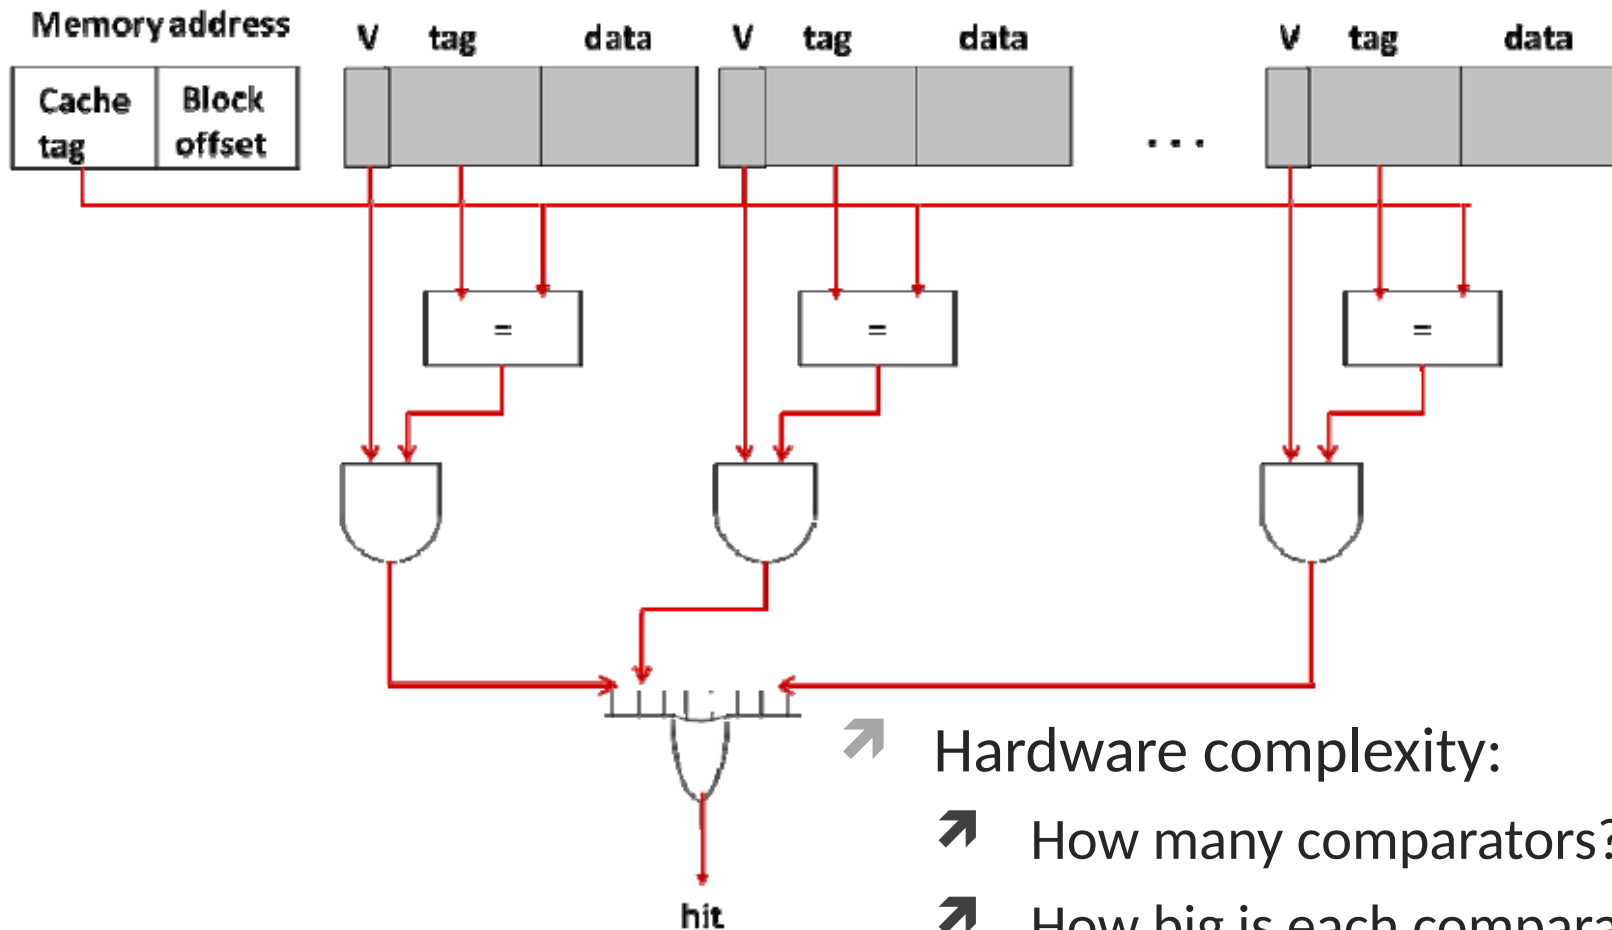
Address interpretation in FA cache

| Cache Tag | Index |
|-----------|-------|
|-----------|-------|

- No splitting memory addresses into “index” and “tag”
- **It all becomes tag!**

| Cache Tag |
|-----------|
|-----------|

Fully associative cache circuitry



➤ Hardware complexity:



- How many comparators?
- How big is each comparator?

Cache organizations


- Fully associative cache ⑨
 - Too much hardware complexity
 - Most flexible
- Direct mapped cache ⑨
 - Least hardware complexity
 - Least flexible
- Can we do better? Is there a compromise?
- Yes! It's called a set-associative cache
- Turns out direct-mapped and fully-associative caches are degenerate cases of a set-associative cache!

Generalization?

Direct Mapped




Fully Associative



Equivalent to n
parallel caches



Set Associative



p parallel caches

| | V | Tag | data |
|----|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |
| 4 | | | |
| 5 | | | |
| 6 | | | |
| 7 | | | |
| 8 | | | |
| 9 | | | |
| 10 | | | |
| 11 | | | |
| 12 | | | |
| 13 | | | |
| 14 | | | |
| 15 | | | |

| | V | Tag | data |
|---|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |
| 4 | | | |
| 5 | | | |
| 6 | | | |
| 7 | | | |

(b) 2-way set associative

Memory block in one of two places

(a) Direct-mapped cache

Memory block in exactly one place

(c) 4-way set associative

Memory block in one of four places

| | V | Tag | data |
|---|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |

| | V | Tag | data |
|---|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |

| | V | Tag | data |
|---|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |

| | V | Tag | data |
|---|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |

Synonyms

- Unfortunate but...
 - Four different ways of saying the same thing
 - Cache **line**
 - Cache **block**
 - Cache **entry**
 - Cache **element**
 - All mean the same thing...the basic unit of data transferred into/out of the cache at a time
 - The textbook has a couple of typos in which it implies *cache set* is also a synonym; that's wrong based on current usage as well as the text on 9-14 and 9-34!
- A cache **set** is a “row” in the cache. The number of blocks per set is determined by the type of the cache
 - Direct mapped: **n** sets, **1** element
 - P-way set associative: **n/p** sets, **p** elements
 - Fully associative: **1** set, **n** elements

16 sets

| | V | Tag | data |
|----|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |
| 4 | | | |
| 5 | | | |
| 6 | | | |
| 7 | | | |
| 8 | | | |
| 9 | | | |
| 10 | | | |
| 11 | | | |
| 12 | | | |
| 13 | | | |
| 14 | | | |
| 15 | | | |

| | V | Tag | data |
|---|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |
| 4 | | | |
| 5 | | | |
| 6 | | | |
| 7 | | | |

| | V | Tag | data |
|---|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |
| 4 | | | |
| 5 | | | |
| 6 | | | |
| 7 | | | |

8 sets of 2 elements

(b) 2-way set associative

(a) Direct-mapped cache

| | V | Tag | data |
|---|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |

| | V | Tag | data |
|---|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |

| | V | Tag | data |
|---|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |

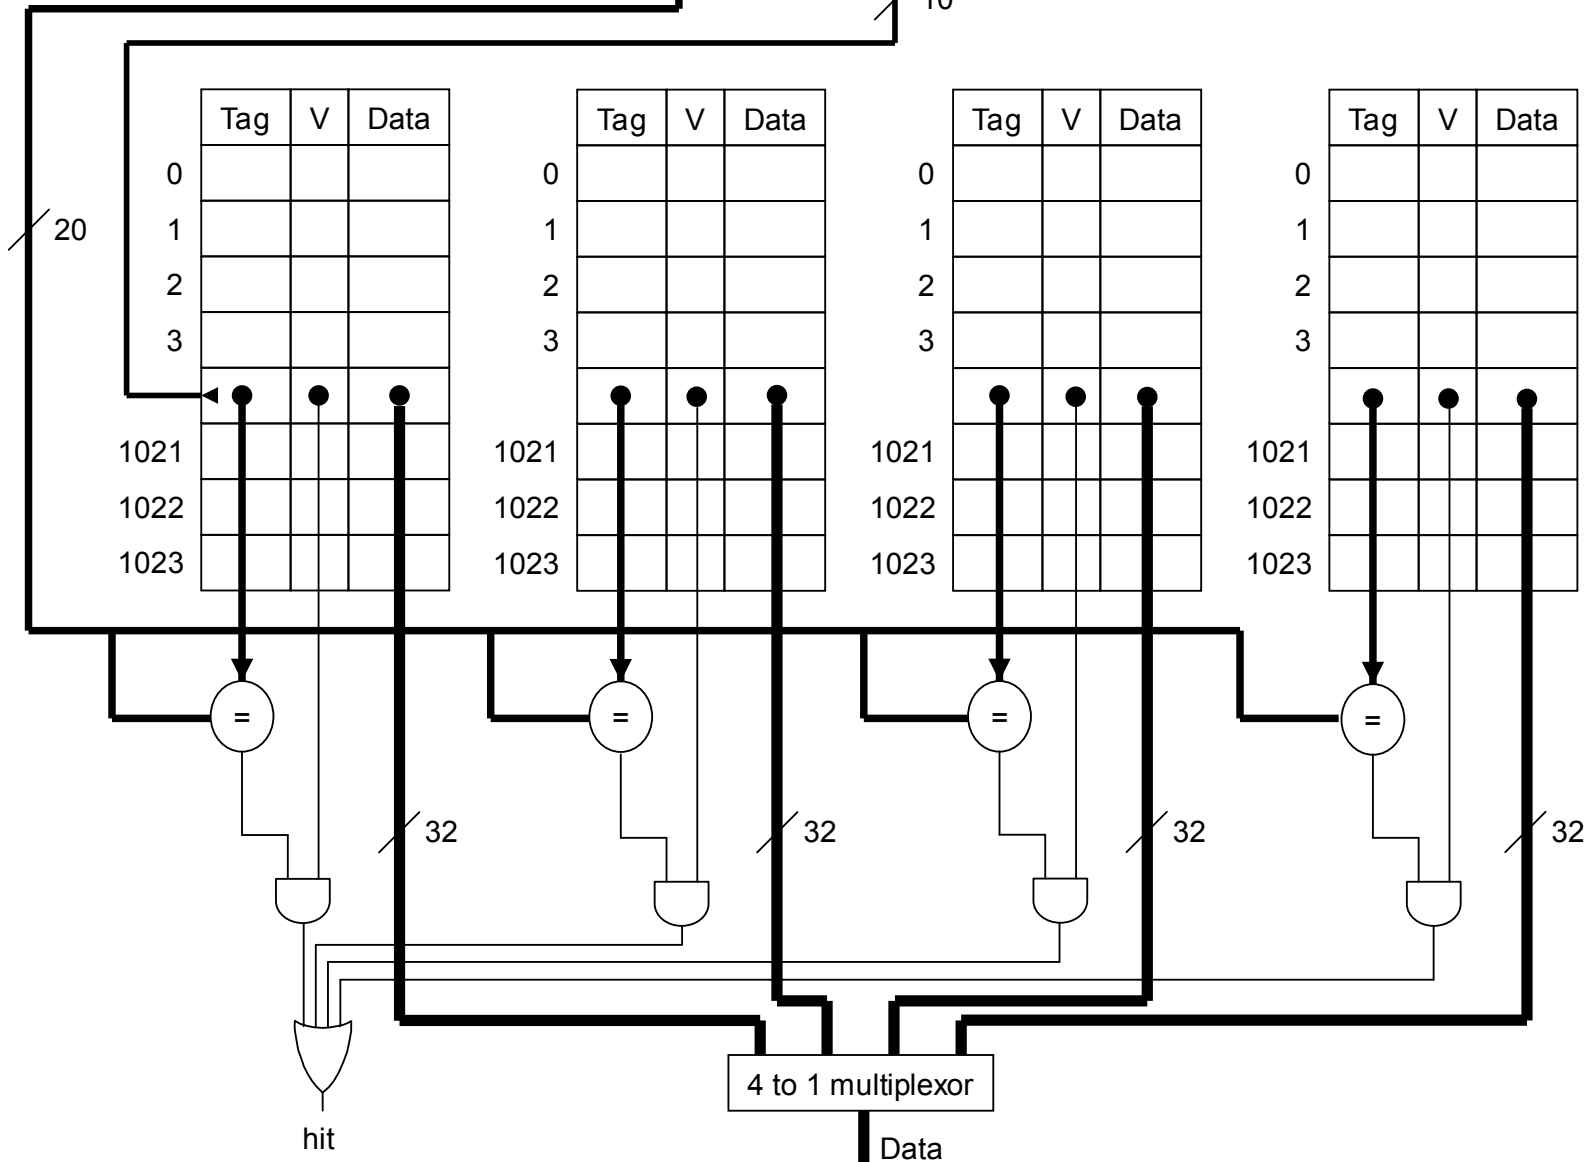
| | V | Tag | data |
|---|---|-----|------|
| 0 | | | |
| 1 | | | |
| 2 | | | |
| 3 | | | |

4 sets of 4 elements

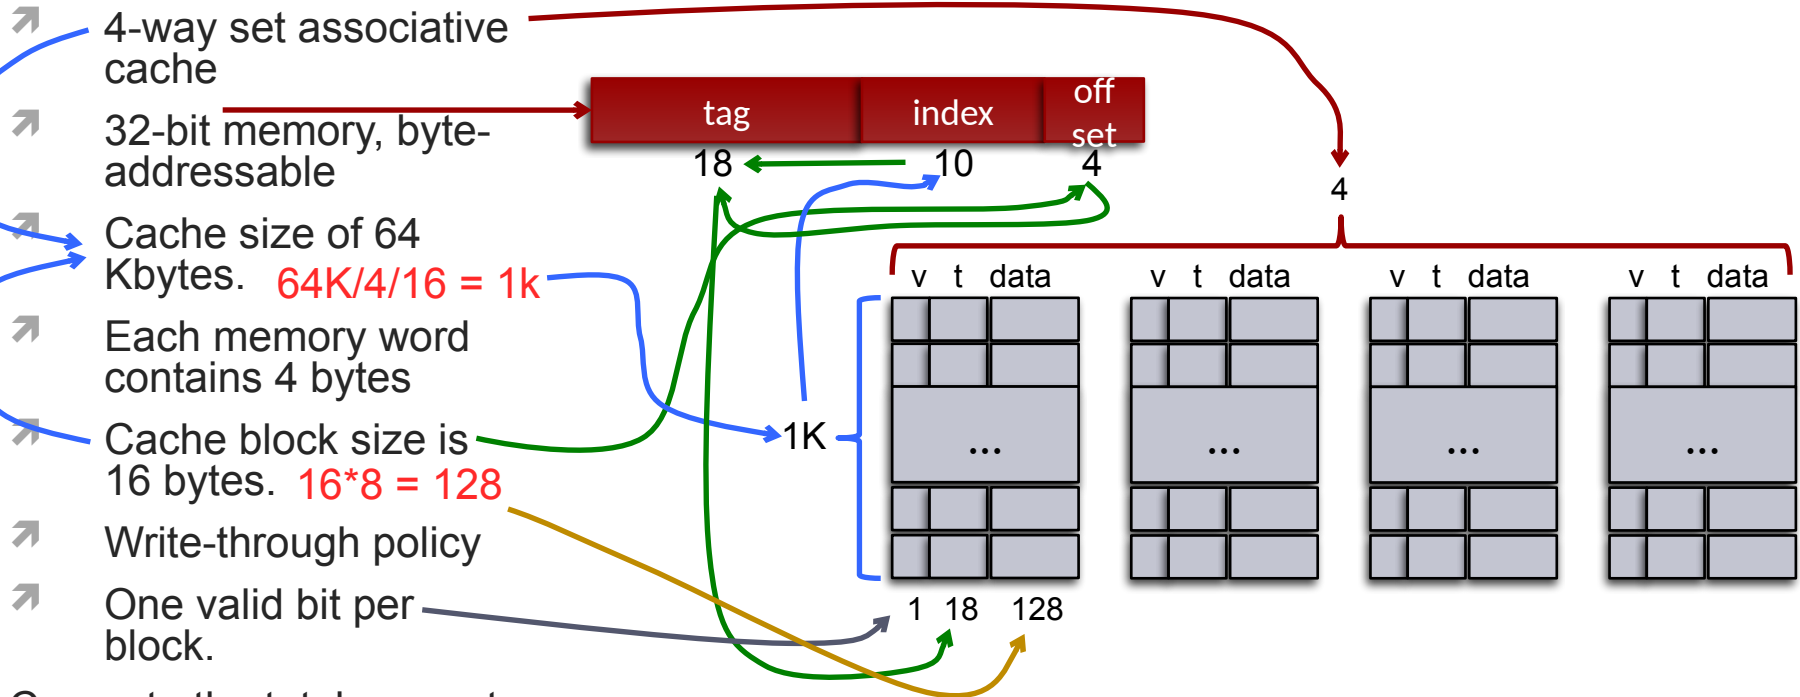
(c) 4-way set associative



10



Example: 4-way set associative cache



Compute the total amount of storage for implementing the cache (i.e. actual data plus the meta-data)

$$(1+18+128) * 1K * 4 / 8$$

$$= 75,264$$


bytes

Hardware complexity?

4 18-bit comparators

In a fully associative cache, ...

...with 64K bytes of data, 64 bytes / block and a t -bit tag

- A. There are four t -bit tag comparators
- B. There are 64 t -bit tag comparators
-  C. There are 1k t -bit tag comparators
- D. There is one t -bit tag comparator for the entire cache

Today's
number is
11,111

FA caches have one comparator for each block/line/entry in the cache. That means $64K/64 = 1K$ is the number of comparators.

In a 4-way set associative cache, ...

...with 64K bytes of data, 64 bytes / block, with a t -bit tag



- A. There are four t -bit tag comparators
- B. There are 64 t -bit tag comparators
- C. There are 1k t -bit tag comparators
- D. There is one t -bit tag comparator for the entire cache

What about cache replacement policy?

| | C0 | | | C1 | | | LRU |
|---|----|-----|------|----|-----|------|-----|
| | V | Tag | data | V | Tag | data | |
| 0 | | | | | | | |
| 1 | | | | | | | |
| 2 | | | | | | | |
| 3 | | | | | | | |
| 4 | | | | | | | |
| 5 | | | | | | | |
| 6 | | | | | | | |
| 7 | | | | | | | |

- What kind of cache is this? **2-way set associative**
- How many LRU bits do we need? **Just one.**
- What happens on every memory access
Set LRU to 0/1 if we read from/store in C0/C1
- So what do we do with a 4-way set associative cache?

LRU replacement in a 4-way cache

| | C0 | | | C1 | | | C2 | | | C3 | | | LRU | | |
|---|----|-----|------|----|-----|------|----|-----|------|----|-----|------|----------------------|--|--|
| | V | Tag | data | V | Tag | data | V | Tag | data | V | Tag | data | | | |
| 0 | | | | | | | | | | | | | c1 -> c3 -> c0 -> c2 | | |
| 1 | | | | | | | | | | | | | c0 -> c2 -> c1 -> c3 | | |
| 2 | | | | | | | | | | | | | c2 -> c3 -> c0 -> c1 | | |
| 3 | | | | | | | | | | | | | c3 -> c2 -> c1 -> c0 | | |

- Do we need a state machine for each cache line?
- Using as many state machines as the number of rows in the cache is a lot of hardware
- Each state machine @ 4! States @ 5 bit state register

What happens on a context switch?


➤ TLB? Flush user portion

➤ Cache? **Nothing!** We're using physical addresses




Note: If the OS brings in a page from disk directly to a physical page frame and bypasses the cache, **it must flush** any cache locations for the previous contents. **This isn't a context switch issue**, it's an issue that if I/O bypasses the cache, then any cache entries referencing the I/O buffer are definitely invalid.

On a process context switch

- A. The cache entries for the process being suspended are flushed.
- B. The cache and TLB entries for the process being suspended are flushed
-  C. The TLB entries for the process being suspended are flushed
- D. None of the above
- E. What does bathroom plumbing have to do with computer science?

Today's number is 11,111

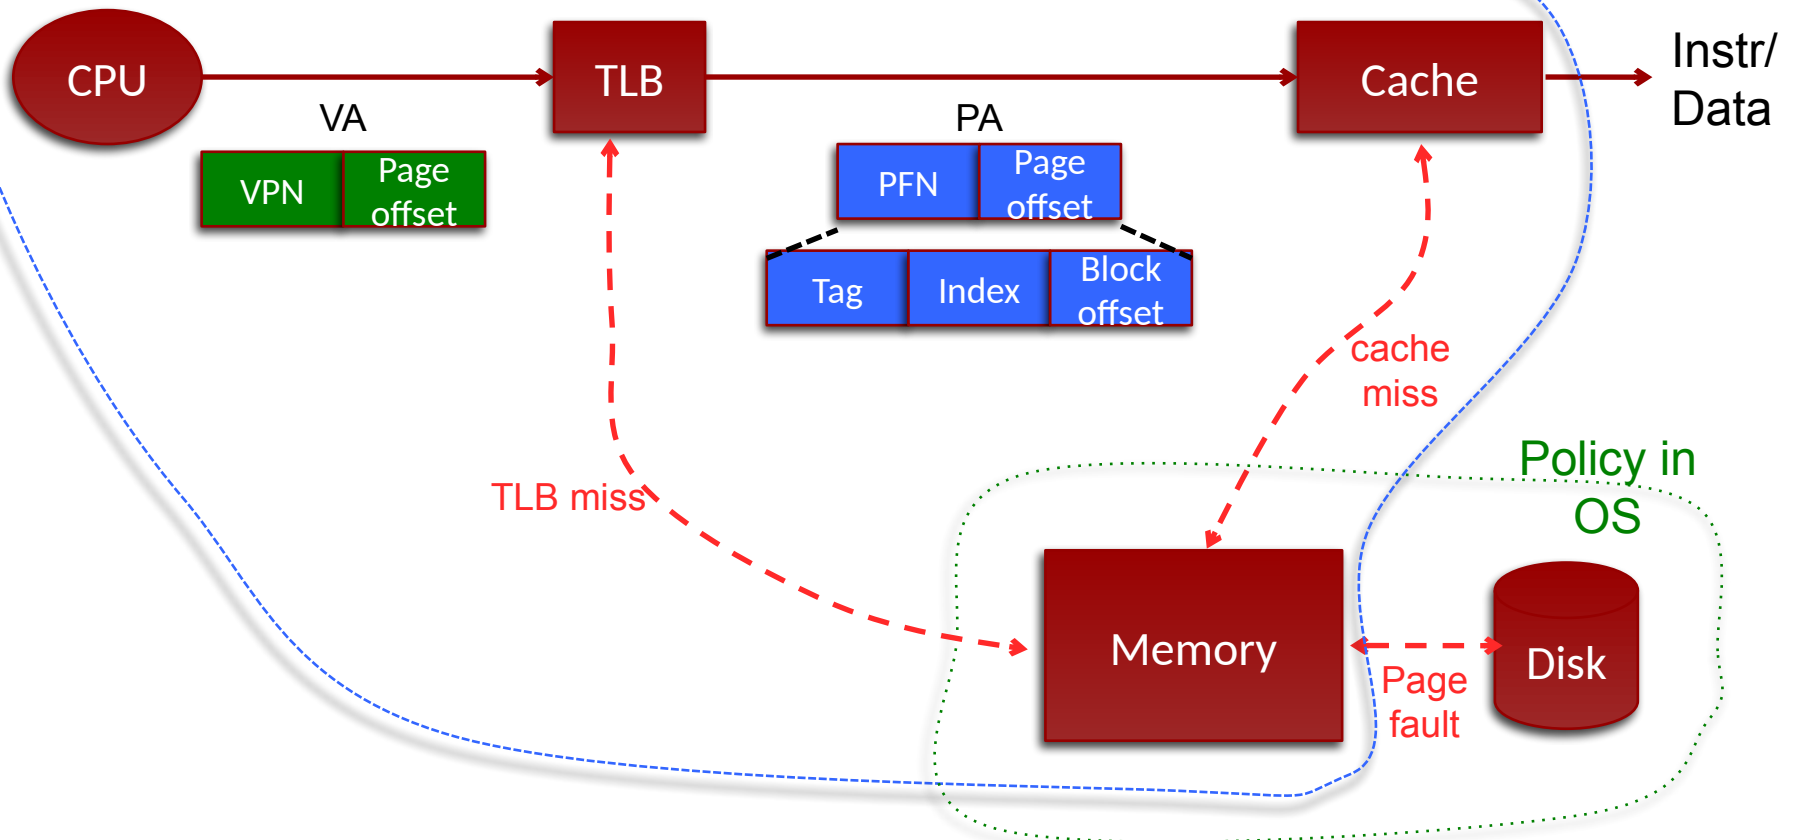
When a page is evicted from a page frame in memory

- A. The cache entries for the evicted page are flushed.
- B. The cache and TLB entries for the evicted page are flushed
-  C. The TLB entries for the evicted page are flushed
- D. None of the above
- E. Still fixated on bathroom plumbing, huh?

The caches we've studied so far hold physical addresses (address translation has already occurred before the cache gets involved), then no flush is needed on page replacement. **If the cache holds virtual addresses, then flushing cache entries from the page is required.**

Putting cache and VM together

Policy in
Hardware



Adding some speed

Policy in
Hardware

TLB miss
PA (VPN)

Cache

Instr/
Data

CPU

VA

TLB

PA

VPN

Page
offset

PFN

Page
offset

Tag

Index

Block
offset

We can have the
TLB try cache first

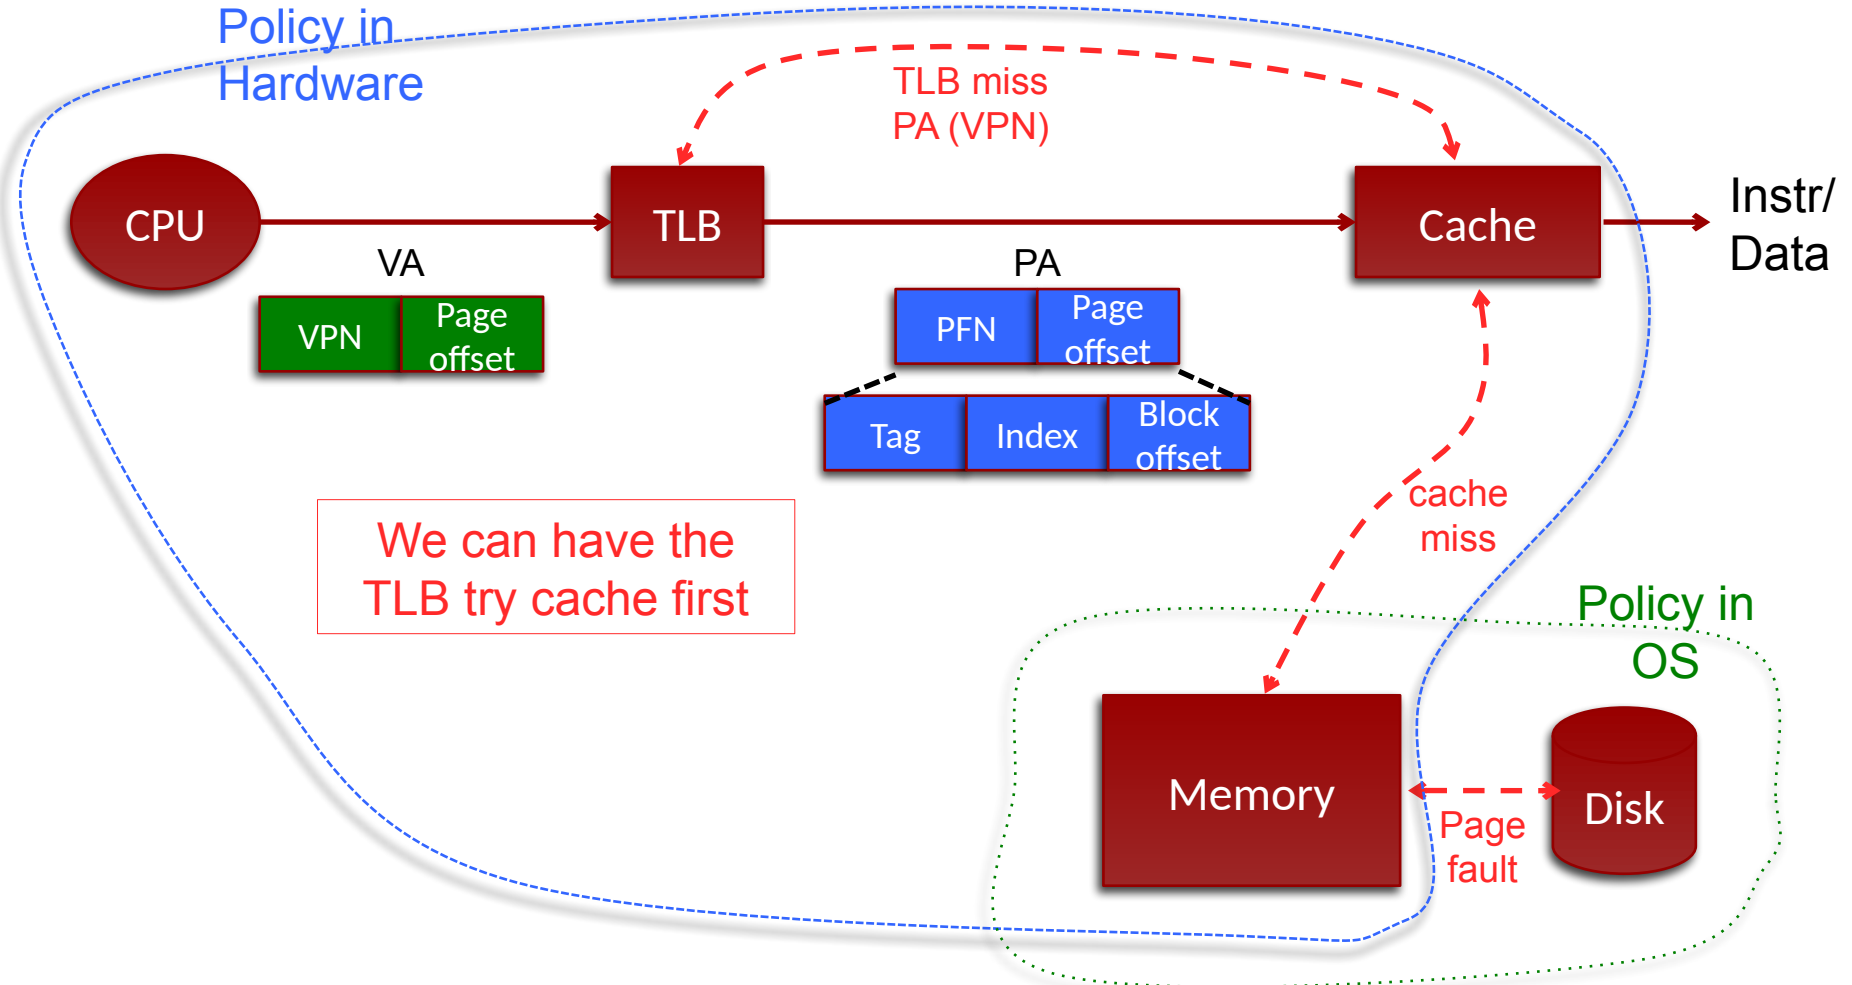
cache
miss

Policy in
OS


Memory

Page
fault

Disk

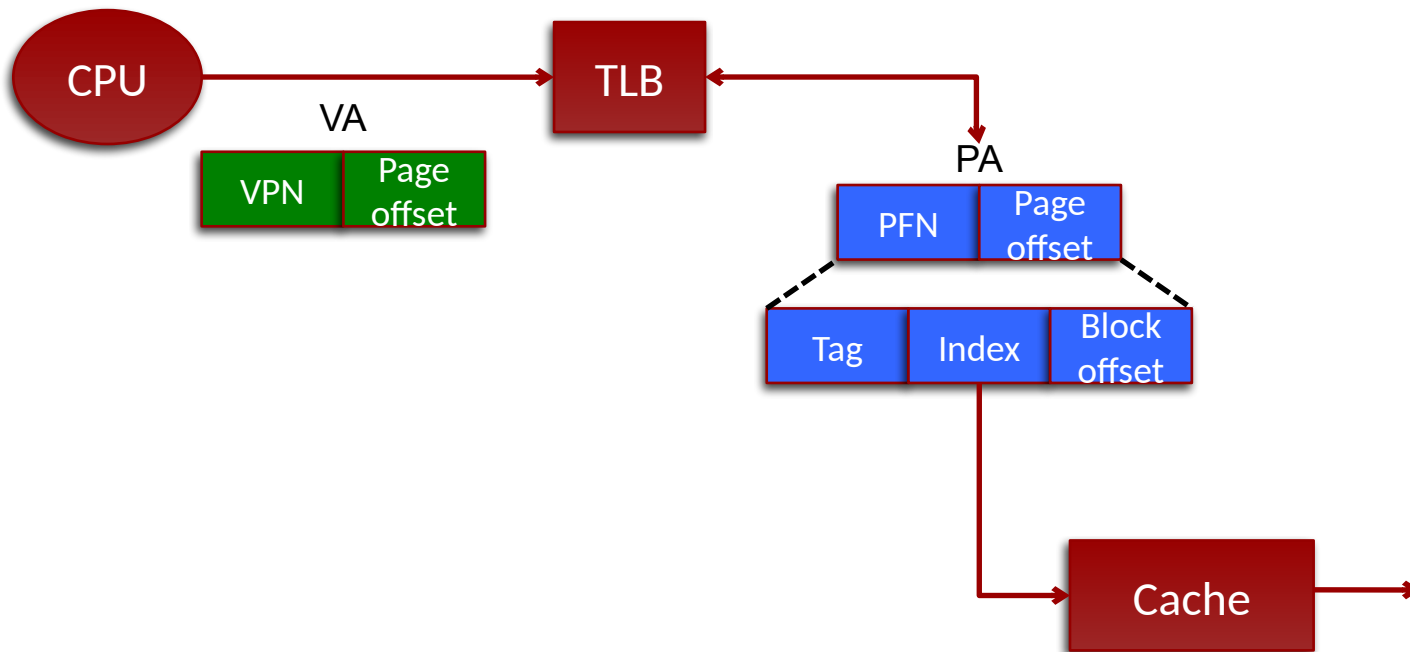


Putting cache and TLB together

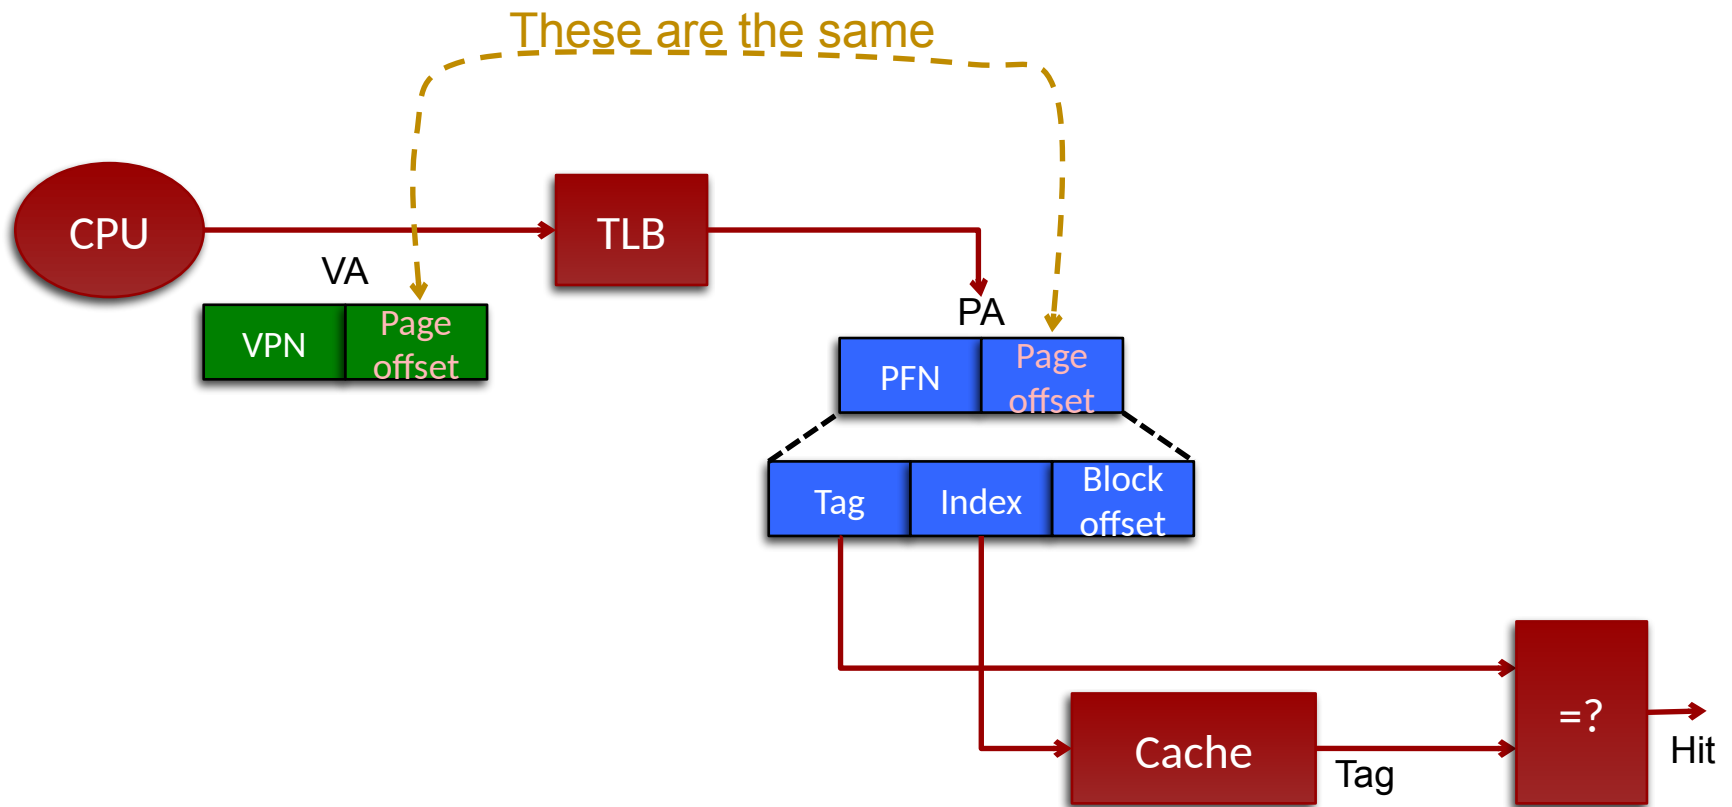


- TLB access is in the critical path of cache access Ⓢ increased clock cycle time
- Can we remove TLB from the critical path?

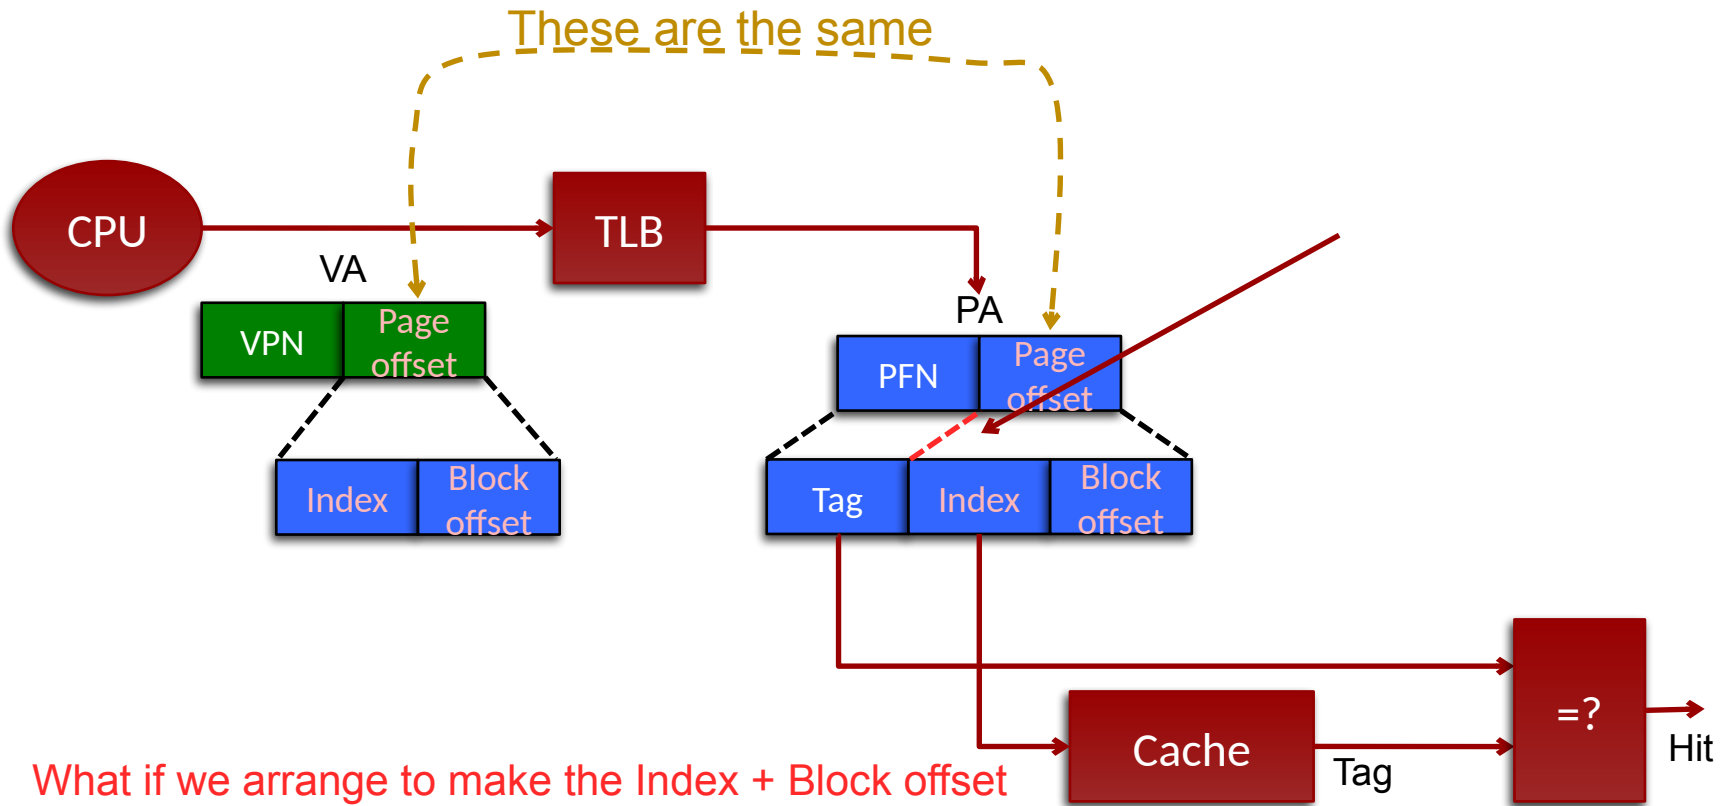
Recall how TLB and Cache work together



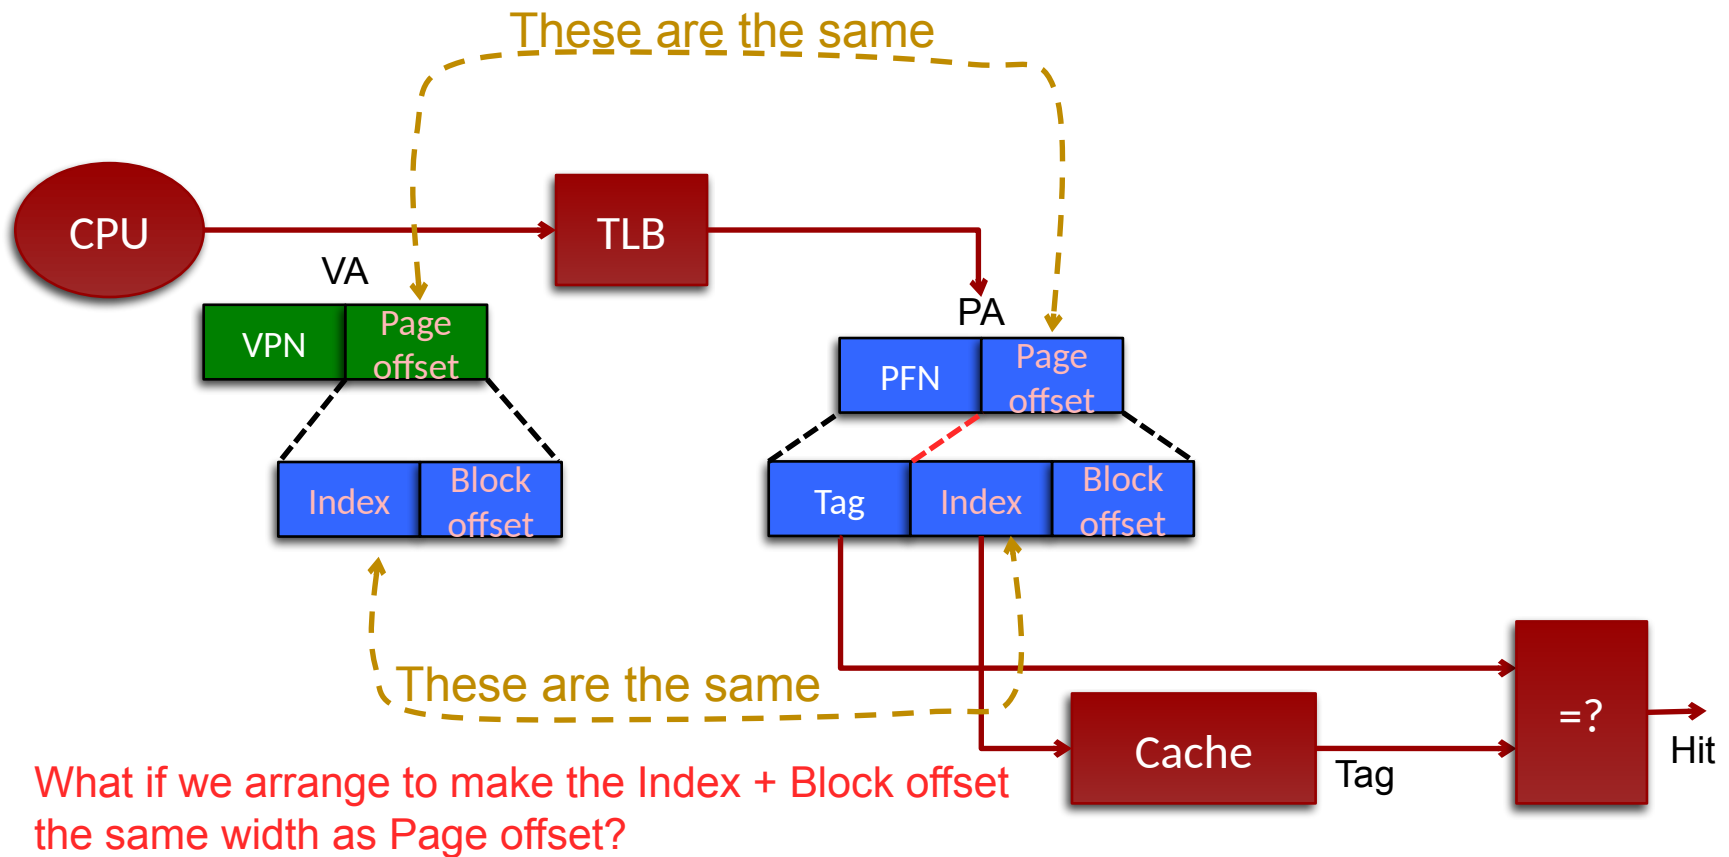
Recall



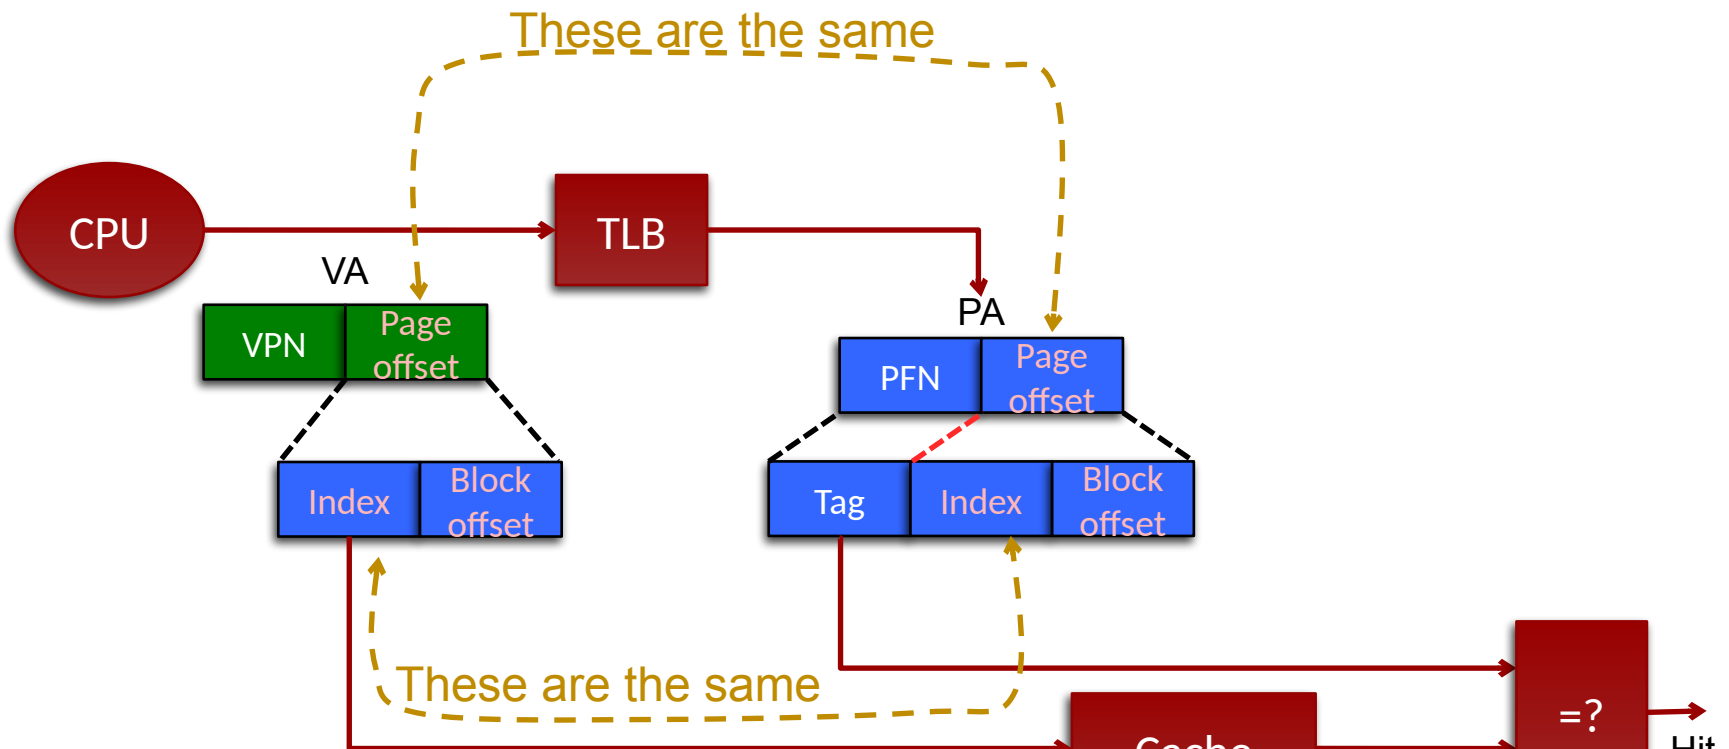
Make the cache index the same size as the page offset?



Now Index is the same for VA and PA!



Cache and TLB can start in parallel!



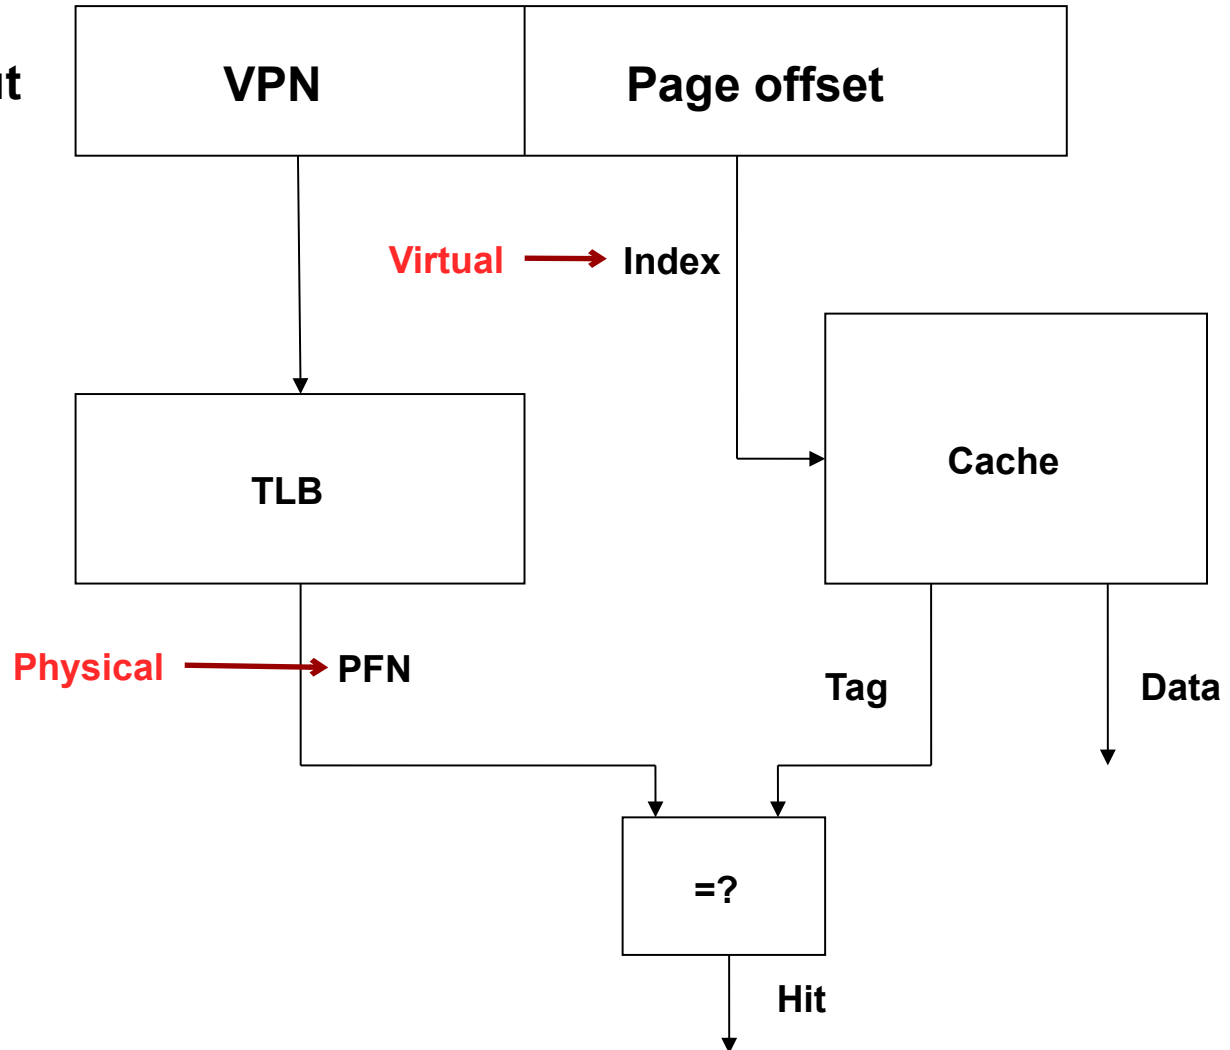
What if we arrange to make the Index + Block offset the same width as Page offset?

Then we can get the cache index from the VA to start the cache read without waiting for the TLB!

Virtually indexed physically tagged cache

How to get TLB out
of the critical path

TLB & cache
access in **parallel**

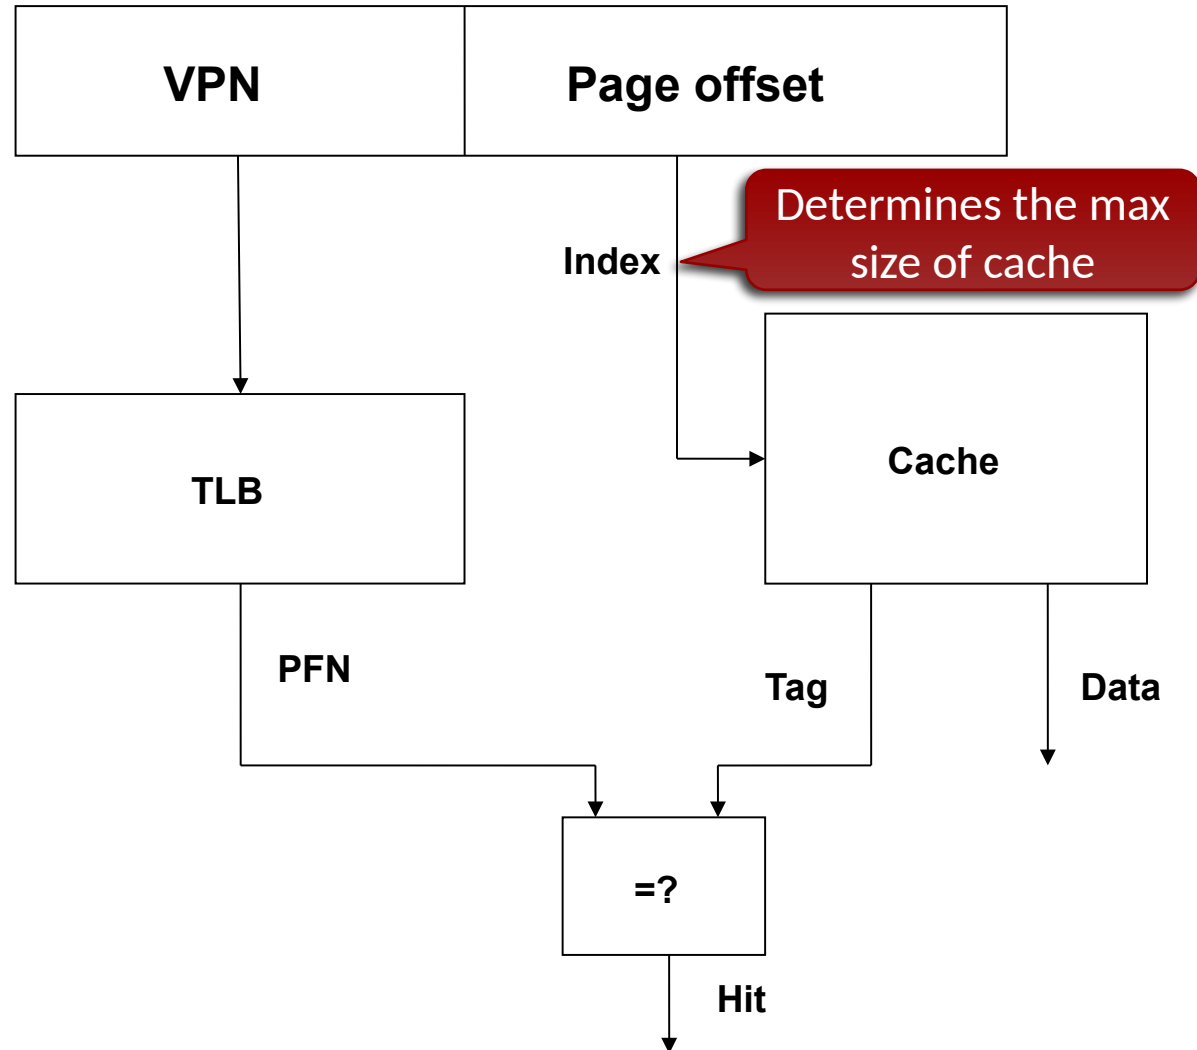


In a virtually indexed physically tagged cache




- A. The TLB and cache are accessed in **parallel**
- B. The TLB and cache are accessed **sequentially**
- C. The TLB and cache are accessed at **random**
- D. None of the above

Is there a limitation?



Page coloring example

OS guarantees some bits of VPN will remain unchanged

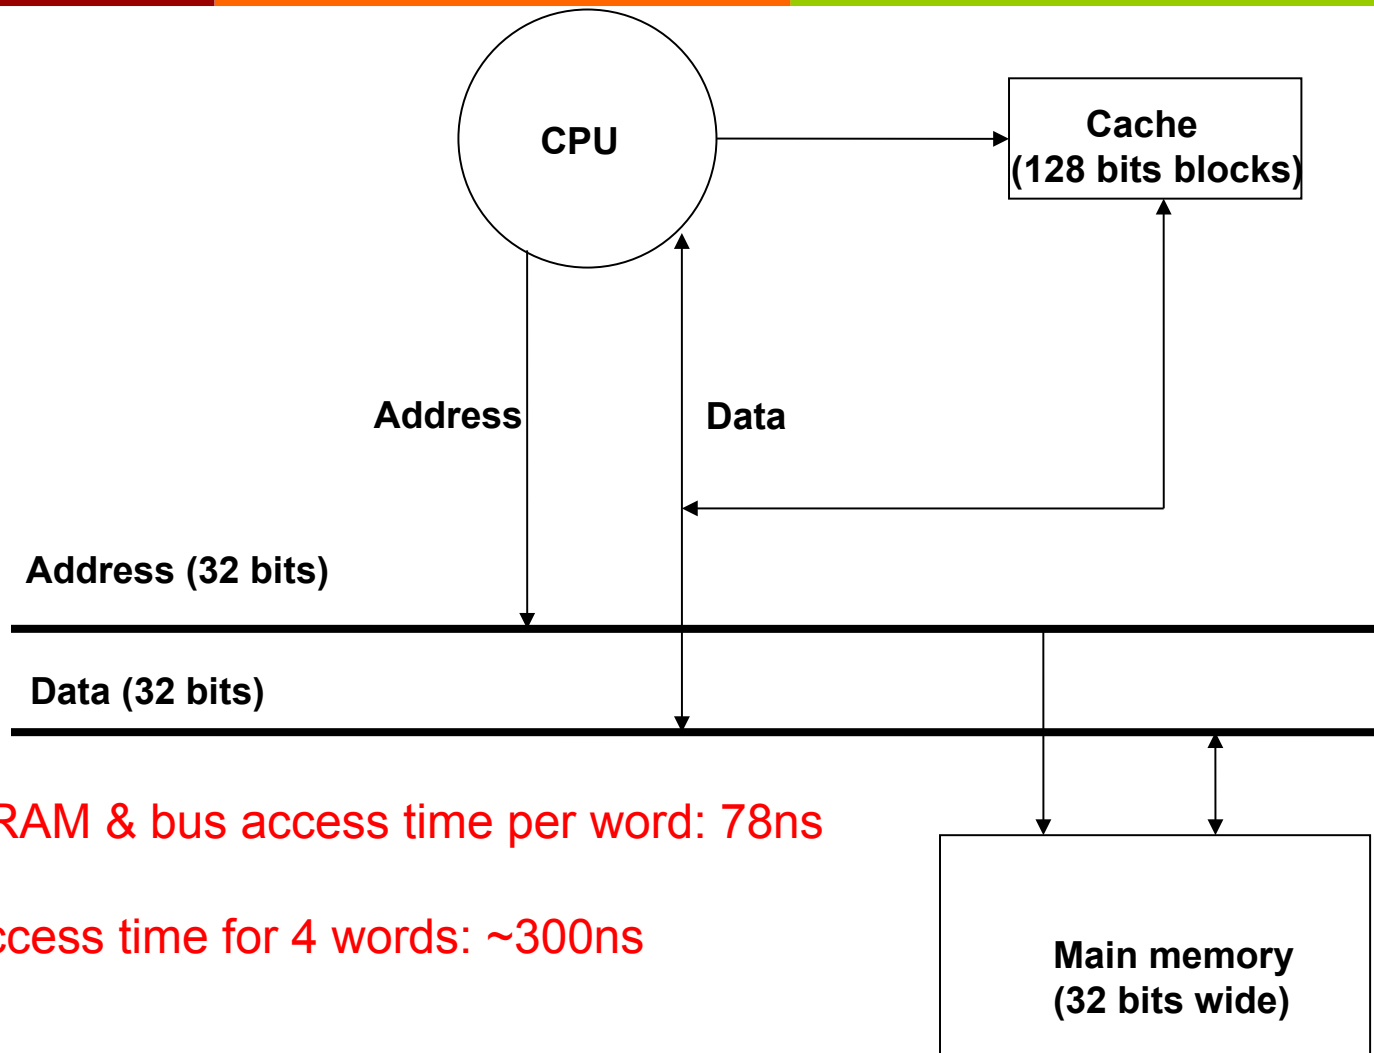


Work out Example 10 on your own and check the solution

Page coloring impacts

- Require the low bits of the VPN and PFN to be identical (4 low order bits in this example)
 - This allows the use of these bits as part of either the virtual or physical address, just like we use the offset
- What does that impact?
 - A virtual page can only occupy a subset of page frames in which the low bits of the VPN match the low bits of the PFN.
 - We refer to this as **page coloring** which means the page replacement algorithm must keep track of the “color” of the VPNs and PFNs (namely those low bits) to ensure virtual pages are only loaded into like-colored page frames.
 - It could make it harder to fit a working set into physical memory, but it's often workable because processes tend to use contiguous pages so the VPNs of the pages are spread evenly among the colors
 - In this example using the low 4 bits of the VPN/PFN, we get 16 different “colors” of page/page frame. So a page with a VPN ending in 0010_2 can only be placed in a page frame with a number that ends in 0010_2

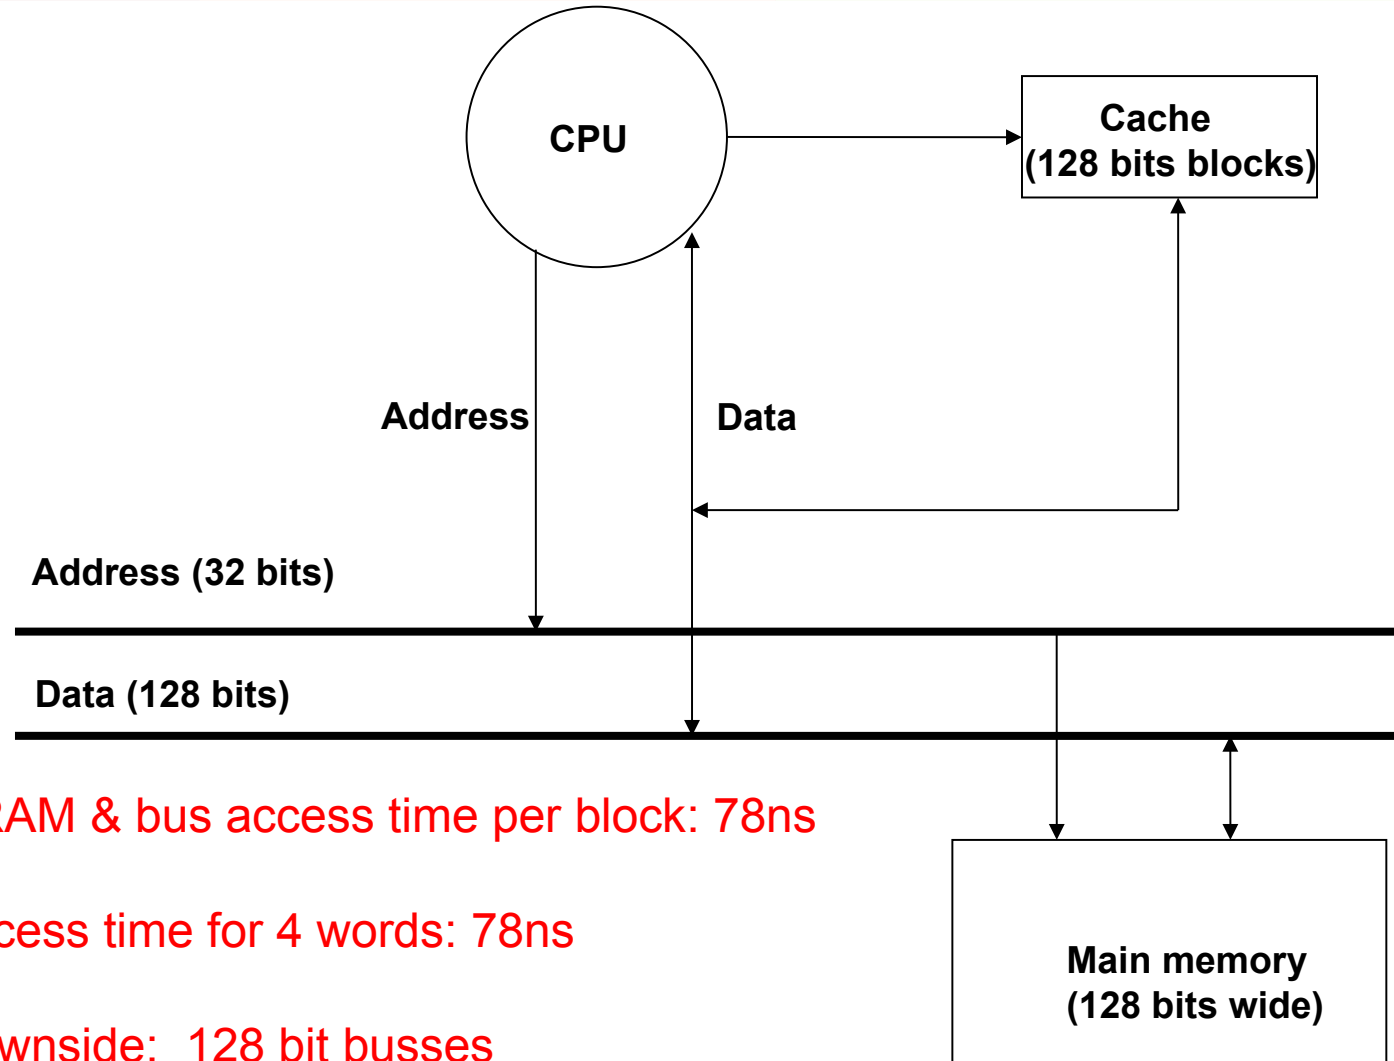
Simple memory system



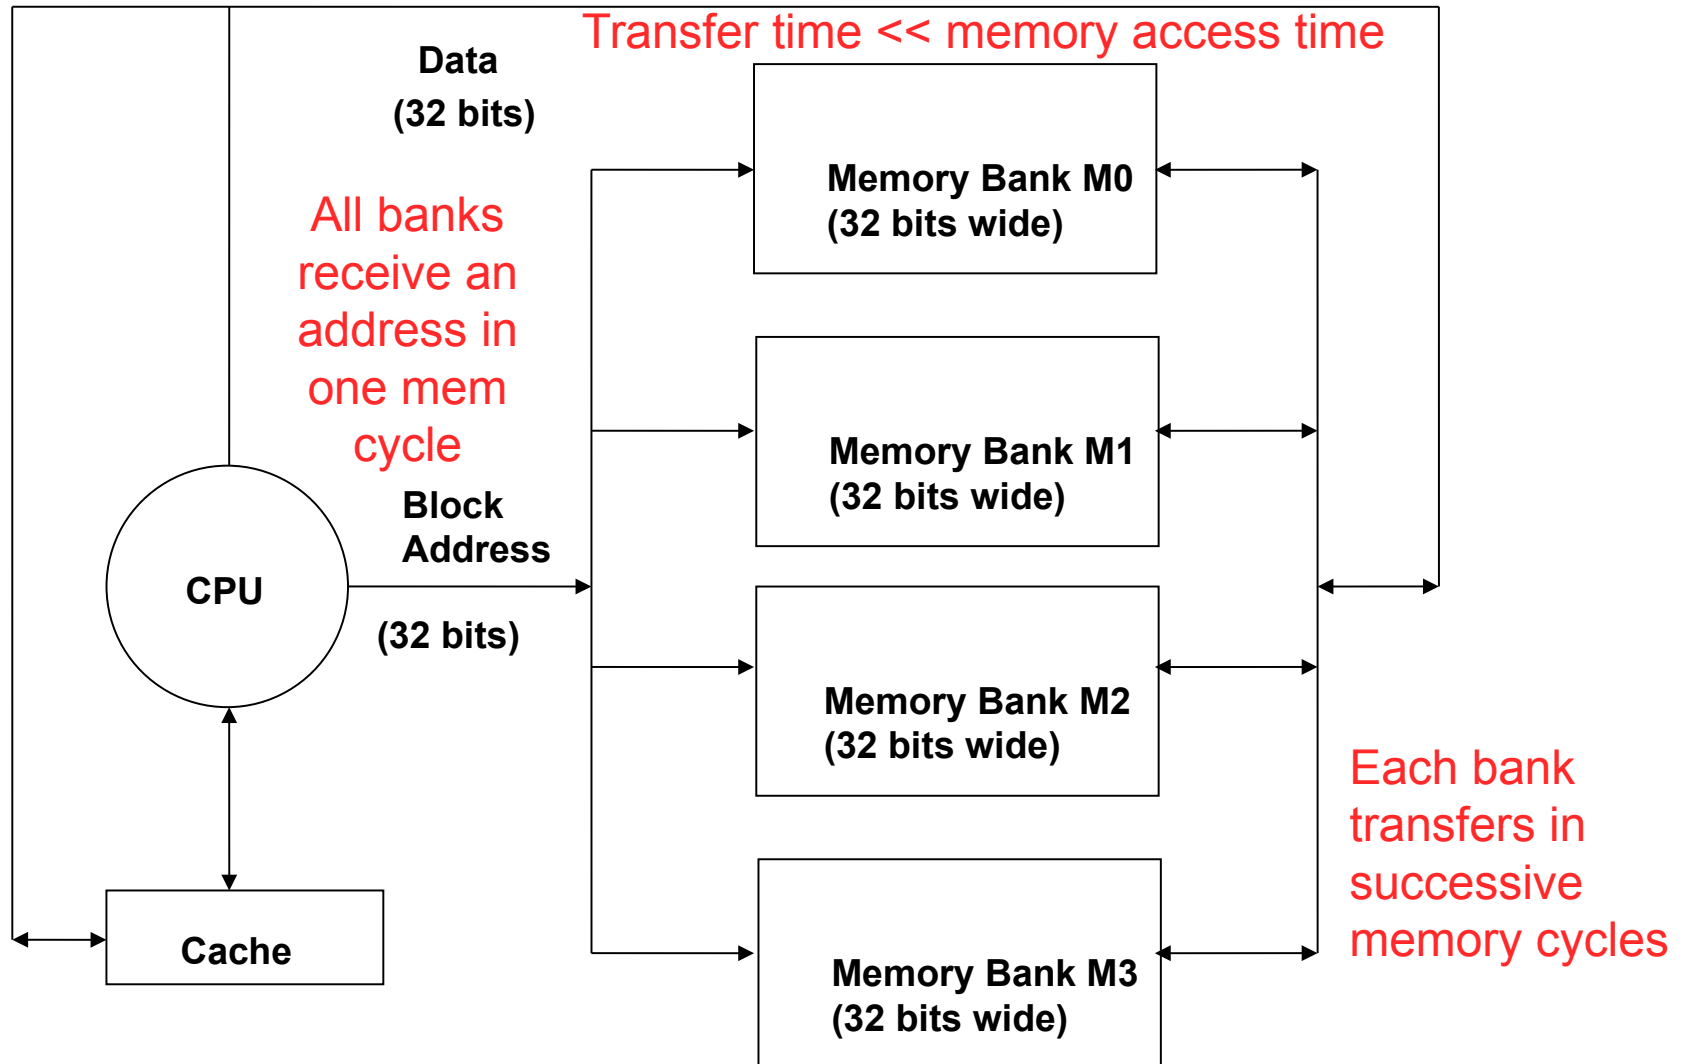
DRAM & bus access time per word: 78ns

Access time for 4 words: ~300ns

Memory system matching cache blocksize



Interleaved memory



Example


- 4-way interleaved memory
- DRAM access time: 70 cycles.
- Memory bus cycle time: 5 cycles
- Compute the block transfer time for a block size of 4 words. Assume all 4 words are first retrieved from the DRAM before the data transfer to the CPU.

Example solution sketch


➤ Start memory
fetches at 0, 5, 10,
15ns

➤ Bank 4 memory
fetch finishes at
85ns

➤ Data transfers to
CPU/cache occur
at 70, 75, 80, 85ns

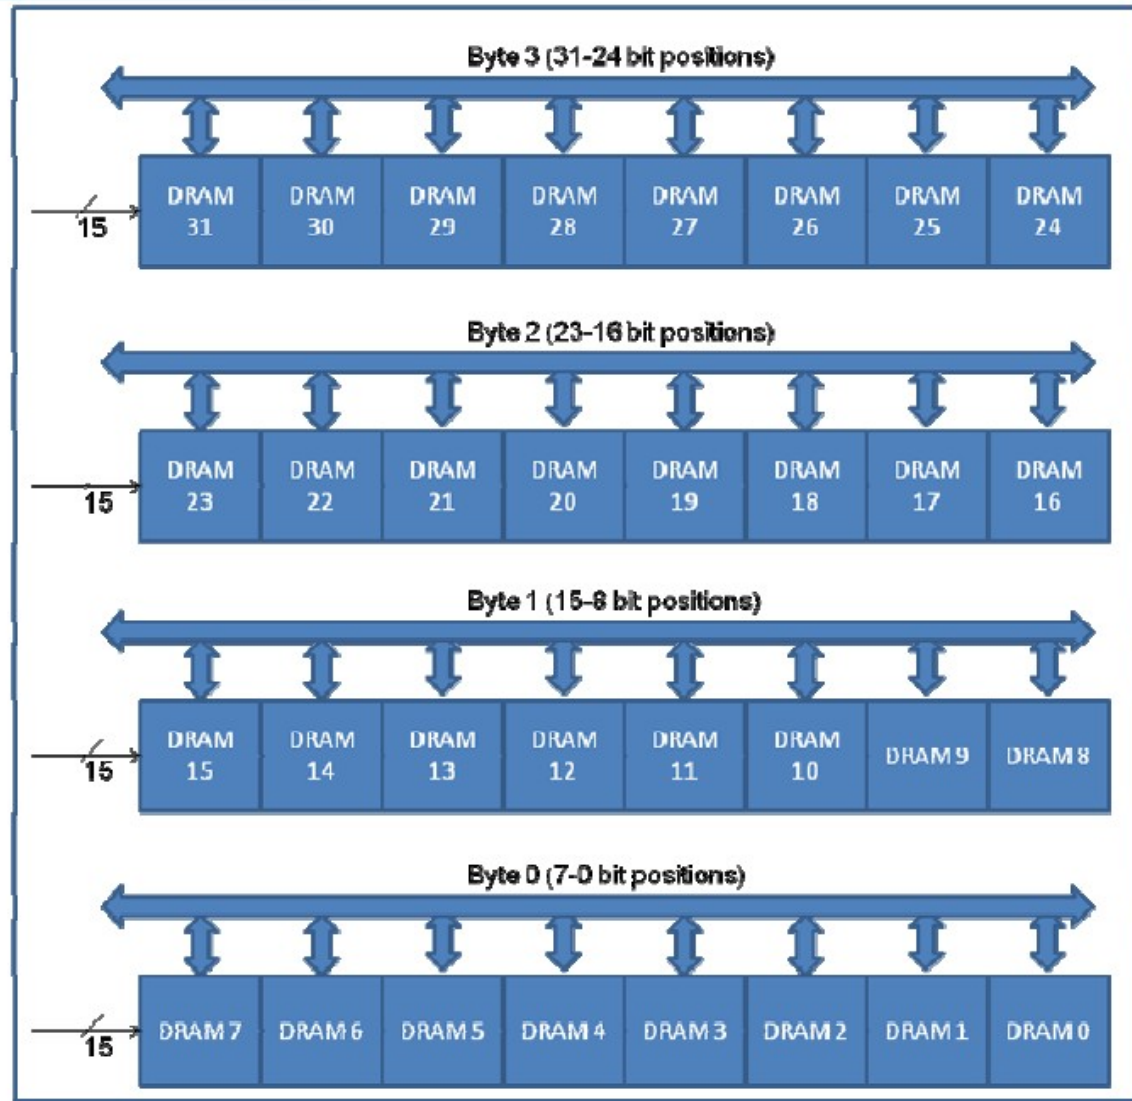



Interleaved memory is useful because

- A. It increases the cache block size
- B. It reduces the cache block size
-  C. It allows concurrent fetches of data blocks from memory without requiring wider busses
- D. It decreases the size of the cache index



15-bit row/column address from Memory controller



Two-cycle addressing using the same 15 addressing pins:

- Present row address and RAS'
- Present column address and CAS'

Figure 9.42: A 4 GByte memory system using a 1 Gbit DRAM chip

Relative sizes & latencies, ca. 2009

| Type of Memory | Typical Size | Approximate latency in CPU clock cycles to read one word of 4 bytes |
|--------------------------|----------------------------|--|
| CPU registers | 8 to 32 | Usually immediate access (0-1 clock cycles) |
| L1 Cache | 32 (Kilobyte) KB to 128 KB | 3 clock cycles |
| L2 Cache | 128 KB to 4 Megabyte (MB) | 10 clock cycles |
| Main (Physical) Memory | 256 MB to 4 Gigabyte (GB) | 100 clock cycles |
| Virtual Memory (on disk) | 1 GB to 1 Terabyte (TB) | 10^6 to 10^7 clock cycles (not accounting for the software overhead of handling page faults) |

Relative sizes & latencies, ca. 2022

| Type of Memory | Typical Size | Approximate latency in CPU clock cycles to read one word of 4 bytes |
|--------------------------|----------------|---|
| CPU registers | 8 to 32 | Usually immediate access (0-1 clock cycles) |
| L1 Cache | 32 KB to 4MB | 1-3 clock cycles |
| L2 Cache | 32 KB to 32 MB | ~10 clock cycles |
| L3 Cache | 1 MB to 768 MB | ~30 clock cycles |
| Main (Physical) Memory | 256 MB to 8 TB | ~100 clock cycles |
| Virtual Memory (on disk) | 1 GB to 8 TB? | 10^6 to 10^7 clock cycles (not including the software overhead of handling page faults) |

Summary of Chapter 9

terminology

| Category | Vocabulary | Details |
|-------------------------------------|-----------------------------|---|
| Principle of locality (Section 9.2) | Spatial | Access to contiguous memory locations |
| | Temporal | Reuse of memory locations already accessed |
| Cache organization | Direct-mapped | One-to-one mapping (Section 9.6) |
| | Fully associative | One-to-any mapping (Section 9.12.1) |
| | Set associative | One-to-many mapping (Section 9.12.2) |
| Cache reading/writing (Section 9.8) | Read hit/Write hit | Memory location being accessed by the CPU is present in the cache |
| | Read miss/Write miss | Memory location being accessed by the CPU is not present in the cache |
| Cache write policy (Section 9.8) | Write through | CPU writes to cache and memory |
| | Write back | CPU only writes to cache; memory updated on replacement |
| Cache parameters | Total cache size (S) | Total data size of cache in bytes |
| | Block Size (B) | Size of contiguous data in one data block |
| | Degree of associativity (p) | Number of homes a given memory block can reside in a cache |
| | Number of cache lines (L) | S/pB |
| | Cache access time | Time in CPU clock cycles to check hit/miss in cache |
| Memory address interpretation | Unit of CPU access | Size of data exchange between CPU and cache |
| | Unit of memory transfer | Size of data exchange between cache and memory |
| | Miss penalty | Time in CPU clock cycles to handle a cache miss |
| | Index (n) | $\log_2 L$ bits, used to look up a particular cache line |
| | Block offset (b) | $\log_2 B$ bits, used to select a specific byte within a block |
| | Tag (t) | a – (n+b) bits, where a is number of bits in memory address; used for matching with tag stored in the cache |

Summary of Chapter 9

terminology

| Category | Vocabulary | Details |
|------------------------------------|--|---|
| Cache entry/cache block/cache line | Valid bit | Signifies data block is valid |
| | Dirty bits | For write-back, signifies if the data block is more up to date than memory |
| | Tag | Used for tag matching with memory address for hit/miss |
| Performance metrics | Data | Actual data block |
| | Hit rate (h) | Percentage of CPU accesses served from the cache |
| | Miss rate (m) | $1 - h$ |
| | Avg. Memory stall | $\text{Misses-per-instruction}_{\text{Avg}} * \text{miss-penalty}_{\text{Avg}}$ |
| | Effective memory access time (EMAT _i) at level i | $\text{EMAT}_i = T_i + m_i * \text{EMAT}_{i+1}$ |
| Types of misses | Effective CPI | $\text{CPI}_{\text{Avg}} + \text{Memory-stalls}_{\text{Avg}}$ |
| | Compulsory miss | Memory location accessed for the first time by CPU |
| | Conflict miss | Miss incurred due to limited associativity even though the cache is not full |
| | Capacity miss | Miss incurred when the cache is full |
| Replacement policy | FIFO | First in first out |
| | LRU | Least recently used |
| Memory technologies | SRAM | Static RAM with each bit realized using a flip flop |
| | DRAM | Dynamic RAM with each bit realized using a capacitive charge |
| Main memory | DRAM access time | DRAM read access time |
| | DRAM cycle time | DRAM read and refresh time |
| | Bus cycle time | Data transfer time between CPU and memory |
| | Simulated interleaving using DRAM | Using page mode bits of DRAM |